

A traits system for the uniform expression of run-time and compile-time polymorphism in Fortran

Konstantinos Kifonidis, Ondrej Certik, Derick Carnazzola

October 21, 2024

Abstract

Based on conclusions drawn from a survey of modern languages, a traits system for Fortran is developed that is fully backwards compatible with the present Fortran language, and allows for the uniform management of source code dependencies on both user-defined and language-intrinsic types via run-time and compile-time polymorphism. The feature set that is described here is small enough to facilitate a first prototype implementation in an open source compiler (like LFortran, LLVM Flang, or GNU Fortran), but at the same time comprehensive enough to already endow Fortran with polymorphism capabilities that equal those of modern programming languages like Swift, Rust, Go, or Carbon. The discussed extensions allow for modern-day, traits-based, object-oriented programming, and powerful, easy to use, fully type-checked, generics. The latter support seamlessly both the procedural, functional, and object-oriented programming styles, and they largely get by *without* requiring manual instantiations by the user. The presented design can also be naturally extended in the future to support rank-genericity for arrays, and compile-time polymorphic union/sum types. The new capabilities are expected to transform the way both Fortran applications and libraries will be written in the future. Decoupled software plugin architectures with enormously improved source code flexibility, reusability, maintainability, and reliability will become possible, without any need for run-time type inspections, and without any loss in computational performance.

Chapter 1

Introduction

Polymorphism was discovered in the 1960ies by Kristen Nygaard and Ole-Johan Dahl during their development of Simula 67, the world's first object-oriented (OO) language [3]. Their work introduced into programming what is nowadays known as “virtual method table” (i.e. function-pointer) based run-time polymorphism, which is both the first focus of this document, and the decisive feature of all OO languages. Several other forms of polymorphism are known today, the most important of them being parametric polymorphism [1], also known as “generics”, which is the second focus of this document, and which has historically developed disjointly from run-time polymorphism since it makes use of compile-time mechanisms.

1.1 The purpose of polymorphism

But what is the purpose of polymorphism in a programming language? What is polymorphism actually good for? One of the more comprehensive answers to this question was given by Robert C. Martin in numerous books (e.g. [11]), as well as in the following quotation from his blog [10]:

“There really is only one benefit to polymorphism; but it's a big one. It is the inversion of source code and run time dependencies.

In most software systems when one function calls another, the runtime dependency and the source code dependency point in the same direction. The calling module depends on the called module. However, when polymorphism is injected between the two there is an inversion of the source code dependency. The calling module still depends on the called module at run time. However, the source code of the calling module does not depend upon the source code of the called module. Rather both modules depend upon a polymorphic interface.

This inversion allows the called module to act like a plugin. Indeed, this is how all plugins work.”

Notice the absence of the words “code reuse” in these statements. The purpose of polymorphism, according to Martin, is the “inversion” (i.e. replacement, or management) of rigid source code dependencies by means of particular abstractions, i.e. polymorphic interfaces (or protocols/traits, as they are also known today). The possibility to reuse code is then merely the logical consequence of such proper dependency management in a polymorphism-based software plugin architecture.

1.2 Source code dependencies in statically typed languages

Which then are the source code dependencies that polymorphism helps us manage? It has been customary to make the following distinction when answering this question:

- Firstly, most larger programs that are written in statically typed languages (like Fortran) have dependencies on *user-defined* procedures and data types. If the programmer employs encapsulation of both a program's procedures and its data, i.e. its state, both these dependencies can actually be viewed as dependencies on user-defined abstract data types. These are the dependencies that Martin is concretely referring to in the above quotation, and it is these dependencies on (volatile) user-defined data and implementations that are particularly troublesome, because they lead to rigid coupling between the various different *parts* of an application. Their results are recompilation cascades, the non-reusability of higher-level source code units, the impossibility to comprehend a large application incrementally, and fragility of such an application as a whole.
- Secondly, every program, that is written in a statically typed language, also has dependencies on abstract data types that are provided by the language itself. Fortran's integer, real, etc. intrinsic types are examples of *language-intrinsic* abstract data types. While hard-wired dependencies on such intrinsic types do not couple different parts of a program (because the implementations of these types are supplied by the language), they nevertheless make a program's source code rigid with respect to the data that it can be used on.

The most widely used approaches to manage dependencies on language-intrinsic types have so far been through generics, while dependency management of user-defined (abstract data) types has so far been the task of OO programming and OO design patterns. Martin [11] has, for instance, defined object-orientation as follows:

"OO is the ability, through the use of polymorphism, to gain absolute control over every source code dependency in [a software] system. It allows the architect to create a plugin architecture, in which modules that contain high-level policies are independent of modules that contain low-level details. The low-level details are relegated to plugin modules that can be deployed and developed independently from the modules that contain high-level policies."

1.3 Modern developments

Notice how Martin's modern definition of object-orientation, that emphasizes source code decoupling, is the antithesis to the usually taught "OO" approaches of one class rigidly inheriting implementation code from another. Notice also how his definition does not require some specific type of polymorphism for the task of dependency management, as long as (according to Martin's first quotation) the mechanism is based on polymorphic interfaces.

Martin's statements on the purpose of both polymorphism and OO simply reflect the two crucial developments that have taken place in these fields over the last decades. Namely, the realizations that

- run-time polymorphism should be freed from the conflicting concept of implementation inheritance (to which it was originally bound given its Simula 67 heritage), and be formulated exclusively in terms of conformance to polymorphic interfaces, i.e. function signatures, or purely procedural abstractions, and that

- compile-time polymorphism should be formulated in exactly the same way as well.

These two developments, taken together, have recently opened up the possibility to treat polymorphism, and hence the dependency management of both user-defined and language-intrinsic types, uniformly in a programming language. As a consequence, it has become possible to use the potentially more efficient (but also less flexible) mechanism of compile-time polymorphism also for certain tasks that have traditionally been reserved for run-time polymorphism (in the form of OO programming), and to mix and match the two polymorphism types inside a single application to better satisfy a user's needs for both flexibility and efficiency.

1.4 Historical background

The road towards these realizations was surprisingly long. Over the last five decades, a huge body of OO programming experience first had to demonstrate that the use of (both single and multiple) implementation inheritance breaks encapsulation in OO languages, and therefore results in extremely tightly coupled, rigid, fragile, and non-reusable code. This led to an entire specialized literature on OO design patterns [4, 6, 9], that aimed at avoiding such rigidity by replacing the use of implementation inheritance with the means to formulate run-time polymorphism that are discussed below. It also led to the apprehension that implementation inheritance (but *not* run-time polymorphism) should be abandoned [17]. In modern languages, implementation inheritance is either kept solely for backwards compatibility reasons (e.g. in the Swift and Carbon languages), or it is foregone altogether (e.g. in Rust, and Go).

The first statically typed mainstream programming language that offered a proper separation of run-time polymorphism from implementation inheritance was Objective-C. It introduced “protocols” (i.e. polymorphic interfaces) in the year 1990 [2]. Protocols in Objective-C consist of pure function signatures, that lack implementation code. Objective-C provided a mechanism to implement multiple such protocols by a class, and to thus make classes conform to protocols. This can be viewed as a restricted form of multiple inheritance, namely inheritance of object *specification*, which is also known as *subtyping*. Only a few years later, in 1995, the Java language hugely popularized these ideas using the terms “interfaces” and “interface inheritance” [2]. Today, nearly all modern languages support polymorphic interfaces/protocols, and the basic mechanism of multiple interface inheritance that was introduced to express run-time polymorphism in Objective-C, often in even improved, more flexible, manifestations. The only negative exceptions in this respect being modern Fortran, and C++, which both still stick to the obsolescent Simula 67 paradigm.

A similarly lengthy learning process, as that outlined for run-time polymorphism, took also place in the field of compile-time/parametric polymorphism. Early attempts, notably templates in C++, to render function arguments and class parameters polymorphic, did not impose any constraints on such arguments and parameters, that could be checked by C++ compilers. With the known results on compilation times and cryptic compiler error messages [5]. Surprisingly, Java, the language that truly popularized polymorphic interfaces in OO programming, did not provide an interface based mechanism to constrain its generics. Within the pool of mainstream programming languages, this latter realization was only made with the advent of Rust [12].

Rust came with a traits (i.e. polymorphic interfaces) system with which it is possible for the user to uniformly and transparently express both generics (i.e. compile-time) and run-time polymorphism in the same application, and to relatively easily switch between the two, where possible. Rust's traits are an improved form of protocols/interfaces in that the user can implement them for a type without having these implementations be coupled to the type's actual definition. Thus, ex-

isting types can be made to retroactively implement new traits, and hence be used in new settings (with some minor restrictions on user ownership of either the traits or the types).

Rust's main idea was quickly absorbed by almost all other mainstream modern languages, most notably Swift, Go, and Carbon, with the difference that these latter languages tend to leave the choice between static and dynamic binding of procedures to the compiler, or language implementation, rather than the programmer. C++ is in the process of adopting generics constraints for its "templates" under the term "strong concepts", but without implementing the greater idea to uniformly express *all* the polymorphism in the language through them. An implementation of this latter idea must today be viewed as a prerequisite in order to call a language design "modern". The purpose of this document is to describe extensions to Fortran, that aim to provide the Fortran language with such modern capabilities.

Chapter 2

Case study: Calculating the average value of a numeric array

To illustrate the advanced features and capabilities of some of the available modern programming languages with respect to polymorphism, and hence dependency management, we will make use here of a case study: the simple test problem of calculating the average value of a set of numbers stored inside a one-dimensional array. In the remainder of this chapter, we will first provide an account and some straightforward monomorphic (i.e. rigidly coupled) functional implementation of this test problem, followed by a functional implementation that makes use of both run-time and compile-time polymorphism to manage rigid source code dependencies. In the survey of programming languages that is presented in Chapter 3, we will then recode this standard test problem in an encapsulated fashion, to highlight how the source code dependencies in this problem can be managed in different languages even in more complex situations, that require OO techniques.

2.1 Monomorphic functional implementation

We have chosen Go here as a language to illustrate the basic ideas. Go is easily understood, even by beginners, and is therefore well suited for this purpose (another good choice would have been the Swift language). The code in the following Listing 2.1 should be self explanatory for anyone who is even only remotely familiar with the syntax of C family languages. So, we'll make only a few remarks regarding syntax.

- While mostly following a C like syntax, variable declarations in Go are essentially imitating Pascal syntax, where a variable's name precedes the declaration of the type.
- Go has two assignment operators. The usual `=` operator, as it is known from other languages, and the separate operator `:=` that is used for combined declaration and initialization of a variable.
- Go has array slices that most closely resemble those of Python's Numpy (which exclude the upper bound of an array slice).

Our basic algorithm for calculating the average value of an array of integer elements employs two different implementations for averaging. The first makes use of a "simple" summation of all the array's elements, in ascending order of their array index. While the second sums in a "pair-wise" manner, dividing the array in half to carry out the summations recursively, and switching

to the “simple” method once subdivision is no longer possible. In both cases, the resulting sum is then divided by the array’s number of elements, to obtain the desired average.

Listing 2.1: Monomorphic functional version of the array averaging example in Go.

```
1 package main
2
3 import "fmt"
4
5 func simple_sum(x []int32) int32 {
6     var s int32
7     s = int32(0)
8     for i := 0; i < len(x); i++ {
9         s += x[i]
10    }
11    return s
12 }
13
14 func pairwise_sum(x []int32) int32 {
15     if len(x) <= 2 {
16         return simple_sum(x)
17     } else {
18         m := len(x) / 2
19         return pairwise_sum(x[:m+1]) + pairwise_sum(x[m+1:])
20     }
21 }
22
23 func simple_average(x []int32) int32 {
24     return simple_sum(x) / int32(len(x))
25 }
26
27 func pairwise_average(x []int32) int32 {
28     return pairwise_sum(x) / int32(len(x))
29 }
30
31 // .....
32 // main program
33 // .....
34
35 func main() {
36
37     xi := []int32{1,2,3,4,5}
38
39     var key int32
40
41     fmt.Println("Simple_sum_average:_1")
42     fmt.Println("Pairwise_sum_average:_2")
43     fmt.Print("Choose_an_averaging_method:_")
44     fmt.Scan(&key)
```



```

45
46  switch key {
47  case 1:
48      fmt.Println(simple_average(xi))
49  case 2:
50      fmt.Println(pairwise_average(xi))
51  default:
52      fmt.Println("Case_not_implemented!")
53  }
54
55 }

```

An inspection of Listing 2.1 will readily reveal that this code has three levels of rigid (i.e. hard-wired) dependencies. Namely,

1. function `pairwise_sum` depending on function `simple_sum`'s implementation,
2. functions `simple_average` and `pairwise_average` depending on functions' `simple_sum`, and `pairwise_sum` implementation, respectively, and
3. the entire program depending rigidly on the `int32` data type in order to declare both the arrays that it is operating on, and the results of its summation and averaging operations.

The first two items are dependencies on user-defined implementations, while the third is a typical case of rigid dependency on a language-intrinsic type, which renders the present code incapable of being applied to arrays of any other data type than `int32`s. Given that we are dealing with three levels of dependencies, three levels of polymorphism will accordingly be required to remove all these dependencies.

2.2 Polymorphic functional implementation

Listing 2.2 gives an implementation of our test problem, that employs Go's generics and functional features in order to eliminate the last two of the rigid dependencies that were listed in Sect. 2.1. The code makes use of Go's generics to admit arrays of both the `int32` and `float64` types as arguments to all functions, and to express the return values of the latter. It also makes use of the run-time polymorphism inherent in Go's functional features, namely closures and variables of higher-order functions, to replace the two previous versions of function `average` (that depended on specific implementations), by a single polymorphic version. Only the rigid dependency of function `pairwise_sum` on function `simple_sum` has not been removed, in order to keep the code more readable. In the OO code versions, that will be presented in Chapter 3, even this dependency is eliminated.

A few remarks are in order for a better understanding of Listing 2.2's code:

- In Go, generic type parameters to a function, like the parameter `T` here, are provided in a separate parameter list, that is enclosed in brackets `[]`.
- Generic type parameters have a constraint that follows their declared name. Go exclusively uses interfaces as such constraints (like the interface `INumeric` in the following code).

- Interfaces consist of either explicit function signatures, or *type sets*, like `int32 | float64` in the present example. The latter actually signify a set of function signatures, too, namely the signatures of the intersecting set of all the operations and intrinsic functions for which the listed types provide implementations.
- The code makes use of type conversions to the generic type `T`, where required. For instance, `T(0)` converts the constant `0` to the corresponding zero constant of type `T`.
- The code instantiates closures and stores these by value in two variables named `avi` and `avf` for later use (Fortran and C programmers should note that `avi` and `avf` are *not* function pointers!).

Listing 2.2: Polymorphic functional version of the array averaging example in Go.

```

1 package main
2
3 import "fmt"
4
5 type INumeric interface{ int32 | float64 }
6
7 func simple_sum[T INumeric](x []T) T {
8     var s T
9     s = T(0)
10    for i := 0; i < len(x); i++ {
11        s += x[i]
12    }
13    return s
14 }
15
16 func pairwise_sum[T INumeric](x []T) T {
17     if len(x) <= 2 {
18         return simple_sum(x)
19     }
20     m := len(x) / 2
21     return pairwise_sum(x[:m+1]) + pairwise_sum(x[m+1:])
22 }
23
24 func average[T INumeric](sum func([]T) T, x []T) T {
25     return sum(x) / T(len(x))
26 }
27
28 func main() {
29     xi := []int32{1, 2, 3, 4, 5}
30     xf := []float64{1, 2, 3, 4, 5}
31
32     var key int32
33     var avi func([]int32) int32
34     var avf func([]float64) float64
35

```

```

36 fmt.Println("Simple_sum_average:", 1)
37 fmt.Println("Pairwise_sum_average:", 2)
38 fmt.Print("Choose_an_averaging_method:_")
39 fmt.Scan(&key)
40
41 switch key {
42     case 1:
43         avi = func(x []int32) int32 {
44             return average(simple_sum[int32], x)
45         }
46         avf = func(x []float64) float64 {
47             return average(simple_sum[float64], x)
48         }
49     case 2:
50         avi = func(x []int32) int32 {
51             return average(pairwise_sum[int32], x)
52         }
53         avf = func(x []float64) float64 {
54             return average(pairwise_sum[float64], x)
55         }
56     default:
57         fmt.Println("Case_not_implemented!")
58         return
59 }
60
61 fmt.Println(avi(xi))
62 fmt.Println(avf(xf))
63 }

```

Notice how, in order to instantiate the closures `avi` and `avf` (see the switch statement), manual instantiations of the `simple_sum` and `pairwise_sum` generic functions are required – with the arguments `int32` or `float64` being substituted for the generic type parameter, `T`, of these functions.

The motivation to code the example as in Listing 2.2 is that once the two closures, `avi` and `avf`, have been properly instantiated, they may then be passed from the main program to any other client code that may need to make use of the particular averaging algorithm that was selected by the user. This latter client code would *not* have to be littered with switch statements itself, and it would *not* have to depend on any specific implementations. It would merely depend on the closures' interfaces. The same holds for the OO code versions that are discussed in the next chapter, with objects replacing the closures (both being merely slightly different realizations of the same idea).

Chapter 3

Survey of modern languages

In the present chapter, we give encapsulated (i.e. OO) code versions of the test problem in various modern languages. As in the functional code version that was presented in Sect. 2.2, we employ run-time polymorphism to manage the dependencies on user-defined implementations (in this case abstract data types), and generics in order to manage the dependencies on language-intrinsic types. This serves to illustrate how both run-time and compile-time polymorphism can be typically used for dependency management in an OO setting in these modern languages. The survey also aims to highlight the many commonalities but also some of the minor differences in the approaches to polymorphism that were taken in these different languages. As a final disclaimer, we do not advocate to code problems in an OO manner that can be easily coded in these languages in a functional way (as it is the case for this problem). However, in more complex cases, where many more nested functions would need to be used, and where state would have to be hidden, the OO programming style would be the more appropriate one. Hence, our test problem will stand in, in this chapter, for emulating such a more complex problem, that would benefit from an encapsulated coding style.

3.1 Go

Go has supported run-time polymorphism through (polymorphic) “interfaces” (and thus modern-day OO programming) since its inception. In Go, encapsulation is done by storing state in a `struct` and by binding procedures, that need to use that state, to this same `struct`. Thereby creating a user-defined abstract data type (or ADT) with methods. Go allows the programmer to implement multiple polymorphic interfaces for such a type (i.e. to use multiple interface inheritance), even though it offers no explicit language statement for this purpose.

Instead, a user-defined type is implicitly assumed to implement an interface whenever it provides implementations of all the interface’s function signatures. This (structural) way of implementing interfaces requires only an object reference of the type to be passed to its methods (by means of a separate parameter list, in front of a method’s actual name). It is otherwise decoupled from the type’s (i.e. the ADT’s `struct`) definition. Go, finally, makes it explicit in its syntax that interfaces (like `structs`) are types in their own right, and that hence polymorphic variables (i.e. objects) can be declared in terms of them.

Restrictions in Go are that language-intrinsic types cannot have user-provided methods, and that methods and interfaces cannot be directly implemented for user-defined types whose definitions are located in other packages. That is, the programmer has to write wrappers in the latter case.

Since version 1.18, Go also supports compile-time polymorphism through generics. Go's generics make use of "strong concepts", since they are bounded by constraints that are expressed through interfaces. Hence, the Go compiler will fully type-check generic code. In Go, structures, interfaces, and functions, but not methods, can all be given their own generic type parameters.

3.1.1 Encapsulated version coded in Go

Listing 3.1 gives an encapsulated version of the test problem coded in Go. The two different implementations of the sum function have been encapsulated in two different ADTs named `SimpleSum` and `PairwiseSum`, whereas a third ADT named `Averager` encapsulates the functionality that is required to perform the actual averaging. The latter two ADTs contain the lower-level objects `other` and `drv` of `ISum[T]` type as components, to which they delegate calls to these objects' sum methods. Notice, how the use of the polymorphic interface `ISum[T]` for the declarations of the objects `other` and `drv` enables either `SimpleSum` or `PairwiseSum` instances to be plugged into these objects' clients.

A second interface, named `IAverager`, is used to enable polymorphism for different averaging algorithms. Finally, there's a third interface, `INumeric`, that serves exactly the same purpose as in the functional polymorphic version that was given in Sect. 2.2, namely to make all function arguments and return values polymorphic, by admitting as input and output parameters both the `int32` and `float64` intrinsic types.

Hence, three polymorphic interfaces were required in this code, in order to eliminate the three levels of rigid dependencies that were listed in Sect. 2.1. Notice also that, exempting `INumeric`, all the interfaces and all the user-defined ADTs need to take in generic type parameters in this example. In Go, this is required in order to enable all the sum and average methods to use generic type parameters.

Listing 3.1: Encapsulated Go version of the array averaging example.

```
1 package main
2
3 import "fmt"
4
5 // .....
6 // Interfaces
7 // .....
8
9 type INumeric interface {
10     int32 | float64
11 }
12
13 type ISum[T INumeric] interface {
14     sum(x []T) T
15 }
16
17 type IAverager[T INumeric] interface {
18     average(x []T) T
19 }
20
21 // .....
22 // SimpleSum ADT
```

```

23 // .....
24
25 type SimpleSum[T INumeric] struct {
26 }
27
28 func (self SimpleSum[T]) sum(x []T) T {
29     var s T
30     s = T(0)
31     for i := 0; i < len(x); i++ {
32         s += x[i]
33     }
34     return s
35 }
36
37 // .....
38 // PairwiseSum ADT
39 // .....
40
41 type PairwiseSum[T INumeric] struct {
42     other ISum[T]
43 }
44
45 func (self PairwiseSum[T]) sum(x []T) T {
46     if len(x) <= 2 {
47         return self.other.sum(x)
48     } else {
49         m := len(x) / 2
50         return self.sum(x[:m+1]) + self.sum(x[m+1:])
51     }
52 }
53
54 // .....
55 // Averager ADT
56 // .....
57
58 type Averager[T INumeric] struct {
59     drv ISum[T]
60 }
61
62 func (self Averager[T]) average(x []T) T {
63     return self.drv.sum(x) / T(len(x))
64 }
65
66 // .....
67 // main program
68 // .....
69
70 func main() {

```

```

71
72     var avi IAverager[int32]
73     var avf IAverager[float64]
74
75     xi := []int32{1,2,3,4,5}
76     xf := []float64{1.,2.,3.,4.,5.}
77
78     var key int32
79
80     fmt.Println("Simple_sum_average:_1")
81     fmt.Println("Pairwise_sum_average:_2")
82     fmt.Print("Choose_an_averaging_method:_")
83     fmt.Scan(&key)
84
85     switch key {
86     case 1:
87         avi = Averager[int32]{ SimpleSum[int32]{} }
88         avf = Averager[float64]{ SimpleSum[float64]{} }
89     case 2:
90         avi = Averager[int32]{ PairwiseSum[int32]{ SimpleSum[int32]{} } }
91         avf = Averager[float64]{ PairwiseSum[float64]{ SimpleSum[float64]{} } }
92     default:
93         fmt.Println("Case_not_implemented!")
94         return
95     }
96
97     fmt.Println(avi.average(xi))
98     fmt.Println(avf.average(xf))
99 }

```

The main program makes use of Go’s built-in structure constructors, and chaining of their calls, in order to instantiate objects of the required ADTs. In particular, it instantiates run-time polymorphic Averager objects (depending on whether simple or pairwise sum averaging is to take place), and it does so for both the `int32` and `float64` types separately, in order to then use these objects on `int32` and `float64` data, respectively. That *two* such objects are required (one for each language-intrinsic data type) is connected to the aforementioned fact that in order to make methods use generic type parameters in Go, one has to parameterize interfaces, and instantiate these with different actual data types, as in func `main`’s first two code lines. A single (i.e. unparameterized) `IAverager` interface therefore doesn’t suffice, which is unfortunate from the user’s perspective, as some code duplication in client code cannot be avoided in this way.

3.2 Rust

Like Go, Rust supports both run-time and compile-time polymorphism through polymorphic interfaces, which Rust calls “traits”. Unlike Go, Rust has its programmers implement traits in a nominal manner, by using explicit `impl` code blocks to provide a trait’s method implementations. These same `impl` blocks can also be used to bind to a type so-called “associated functions” which aren’t methods, i.e. which don’t take in a `self` (passed-object dummy) argument. A typical exam-

ple for this are user-defined constructors for structs. See the functions that are named `new` in the following code Listing 3.2, and which are called using `::` syntax.

In contrast to Go, Rust allows the programmer to implement traits for both user-defined *and* language-intrinsic types, and to do so for types that are located in external libraries (called “crates” in Rust), as long as the traits themselves are defined in the programmer’s own crate. The reverse, namely implementing an external trait for a user-owned type, is also possible. Only the (edge) case of implementing an external trait for an external type is not allowed (this is called the “orphan rule” [7]). The latter case requires the use of wrappers.

Comparable to Go, Rust’s generics model allows for the generic parameterization of functions, traits, and user-defined types like structs. Rust does not explicitly forbid generic methods. However, if one defines such a method’s signature within a trait, then this will make the trait unusable for the declaration of any “trait objects” [8], i.e. for the employment of run-time polymorphism. Thus, the Rust programmer will in general (need to) parameterize traits and structs rather than any methods themselves. Rust generics are fully type-checked at compilation time, i.e. Rust supports “strong concepts”.

3.2.1 Encapsulated version coded in Rust

The encapsulated Rust version of our test problem that is given in the following Listing 3.2 is in its outline quite similar to the corresponding Go version. There are, however, a few minor differences, that are listed in the following notes.

- Rust uses angled brackets, `< >`, to indicate generic parameter lists.
- Generics constraints in Rust are typically enforced by specifying the required traits in `impl` blocks using `where` statements.
- Rust does not supply an equivalent to Go’s `int32 | float64` syntax and semantics. In order to enable numeric operations on generic types, Rust instead provides a `Num` trait via an external `num` crate (i.e. library). However, the use of this external dependency makes available neither the `AddAssign` (`+=`) operator, nor casts to generic types within generic routines (as they are provided in Go).
- Since the employed algorithm relies on these features, the following code uses a homemade `INumeric` trait, that derives from Rust’s `Num` and `AddAssign` traits, and extends them by the requirement to also implement an associated function, `new`, that provides the needed type casting (notice that `Self` stands in here for the implementing type). The conformance of the `i32` and `f64` intrinsic types to the `INumeric` trait is then acknowledged via `impl` blocks, that also implement any of the still outstanding functionality of this trait for these two types.
- Where necessary, use of the `Copy` trait is also made, to work around Rust’s default move semantics.
- In order to help make all of the source code dependencies explicit, our Rust version employs `modules`, and `use` statements to import the required functionality.
- Rust’s default structure constructors suffer from the same flaw as Fortran’s. That is, they are unable to initialize from an external scope, structure components that are declared being private to their module. As in Fortran, use of user-defined constructors must be made instead (cf. the `new` functions that are defined in separate `impl` blocks for the ADTs `PairwiseSum` and `Averager`).

- To declare run-time polymorphic variables one has to put so-called “trait objects” into “Boxes”, i.e. to declare smart pointers of them, for dynamic instantiation and heap memory allocation (this is the Rust equivalent to using allocatable polymorphic objects in Fortran).

Listing 3.2: Encapsulated Rust version of the array averaging example.

```

1 pub mod interfaces {
2
3     use num::Num;
4     use std::ops::AddAssign;
5
6     // .....
7     // Interfaces
8     // .....
9
10    pub trait INumeric: Num + AddAssign {
11        fn new(input: usize) -> Self;
12    }
13
14    pub trait ISum<T> {
15        fn sum(&self, x: &[T]) -> T;
16    }
17
18    pub trait IAverager<T> {
19        fn average(&self, x: &[T]) -> T;
20    }
21 }
22
23 pub mod intrinsics {
24
25     use crate::interfaces::INumeric;
26
27     impl INumeric for i32 {
28         fn new(input: usize) -> i32 {
29             return input as i32
30         }
31     }
32
33     impl INumeric for f64 {
34         fn new(input: usize) -> f64 {
35             return input as f64
36         }
37     }
38
39 }
40
41 pub mod simple_library {
42

```

```

43 use crate::interfaces::{INumeric,ISum};
44
45 // .....
46 // SimpleSum ADT
47 // .....
48
49 pub struct SimpleSum;
50
51 impl<T> ISum<T> for SimpleSum where T: INumeric + Copy {
52     fn sum(&self, x: &[T]) -> T {
53         let mut s: T;
54         s = T::new(0);
55         for i in 0 .. x.len() {
56             s += x[i];
57         }
58         return s
59     }
60 }
61
62 }
63
64 pub mod pairwise_library {
65
66     use crate::interfaces::{INumeric,ISum};
67
68     // .....
69     // PairwiseSum ADT
70     // .....
71
72     pub struct PairwiseSum<T> {
73         other: Box<dyn ISum<T>>,
74     }
75
76     impl<T> PairwiseSum<T> where T: INumeric {
77         pub fn new(other: Box<dyn ISum<T>>) -> PairwiseSum<T> {
78             PairwiseSum{
79                 other: other,
80             }
81         }
82     }
83
84     impl<T> ISum<T> for PairwiseSum<T> where T: INumeric {
85         fn sum(&self, x: &[T]) -> T {
86             if x.len() <= 2 {
87                 return self.other.sum(x);
88             } else {
89                 let m = x.len() / 2;
90                 return self.sum(&x[..m+1]) + self.sum(&x[m+1..]);

```

```

91     }
92 }
93 }
94
95 }
96
97 pub mod averager_library {
98
99     use crate::interfaces::{INumeric,ISum,IAverager};
100
101     // .....
102     // Averager ADT
103     // .....
104
105     pub struct Averager<T> {
106         drv: Box<dyn ISum<T>>,
107     }
108
109     impl<T> Averager<T> where T: INumeric {
110         pub fn new(drv: Box<dyn ISum<T>>) -> Averager<T> {
111             Averager{
112                 drv: drv,
113             }
114         }
115     }
116
117     impl<T> IAverager<T> for Averager<T> where T: INumeric {
118         fn average(&self, x: &[T]) -> T {
119             return self.drv.sum(&x) / T::new(x.len());
120         }
121     }
122 }
123
124
125 // .....
126 // main program
127 // .....
128
129 #[macro_use] extern crate text_io;
130
131 fn main() {
132     use crate::interfaces::IAverager;
133     use crate::simple_library::SimpleSum;
134     use crate::pairwise_library::PairwiseSum;
135     use crate::averager_library::Averager;
136
137     let avsi = Averager::new(Box::new(SimpleSum{}));
138     let avsf = Averager::new(Box::new(SimpleSum{}));

```

```

139
140 let avpi = Averager::new(Box::new(PairwiseSum::new(Box::new(SimpleSum{}))));
141 let avpf = Averager::new(Box::new(PairwiseSum::new(Box::new(SimpleSum{}))));
142
143 let mut avi: Box<dyn IAverager::<i32>> = Box::new(avsi);
144 let mut avf: Box<dyn IAverager::<f64>> = Box::new(avsf);
145
146 let xi : [i32;5] = [1,2,3,4,5];
147 let xf : [f64;5] = [1.,2.,3.,4.,5.];
148
149 let key: i32;
150
151 println!("Simple_sum_average:_1");
152 println!("Pairwise_sum_average:_2");
153 scan!("{}",key);
154
155 match key {
156     1 => {}
157     2 => { avi = Box::new(avpi);
158           avf = Box::new(avpf); }
159     _ => { println!("Case_not_implemented!");
160           return; }
161 }
162
163 println!("{}", avi.average(&xi));
164 println!("{}", avf.average(&xf));
165 }

```

The Rust program version is somewhat longer than the corresponding Go version, because of the need to implement the aforementioned custom type conversion functionality, to provide user-defined constructors in order to initialize (opaque) module-hosted abstract data types, and to import dependencies from modules into the main function (as it would be necessary in realistic situations). The main function's logic could also not be expressed as concisely as in the Go version, because the various trait objects required "boxing" for their instantiation, because Rust's default move semantics had to be worked around, and because Rust wouldn't allow the declaration of variables without simultaneous initialization. Otherwise the codes are pretty much identical¹.

3.3 Swift

Being a successor language to Objective-C, Swift differs slightly from the languages considered so far in that it opted to retain implementation inheritance for backwards compatibility to Objective-C, whereas both Go and Rust do not support implementation inheritance *by design*. Swift therefore

¹Notice that the present Rust version makes universal use of dynamic method dispatch via trait objects, in order to correspond most closely to all the other implementations that we provide in both the present chapter, and in Sect. 6.2.1. An alternative, more idiomatic, Rust version that is equivalent to the Fortran version which we give in Sect. 6.2.2, and that effects static dispatch of the various sum methods through the use of generics, can be found in the Code subdirectory that is accompanying this document.

supports “classical” classes, but it also allows one to bind methods to structures (which, in contrast to classes, are value types in Swift).

Like Go and Rust, Swift (furthermore) supports a traits system in order to implement both run-time and compile-time polymorphism through polymorphic interfaces, that are called “protocols” in Swift. If the Swift programmer chooses to ignore implementation inheritance and classes, he can therefore very much program with structures and protocols in Swift as he would with structures and interfaces/traits in Go and Rust, respectively.

Given Swift’s backwards compatible design, implementation of a protocol (i.e. interface inheritance) is usually done as in classical OO languages, i.e. within a structure’s or a class’s definition. A colon (:) followed by one or more interface names must be supplied for this purpose after the structure’s or class’s own name. However, a very powerful facility for types to implement protocols retroactively is also provided, through so-called extensions, that work even if the types’ source code is inaccessible (because one is, e.g., working with a library in binary form). This same facility also allows for protocols to be implemented by language-intrinsic types. For instance, the following little program, given by Listing 3.3, prints out “I am 4.9”.

Listing 3.3: Swift example of implementing a protocol for an intrinsic data type.

```
1 protocol Printable {  
2     func output()  
3 }  
4 extension Float64: Printable {  
5     func output() {  
6         print("I_am_\(self)")  
7     }  
8 }  
9 var x: Float64 = 4.9  
10 x.output()
```

Swift generics support “strong concepts”, and are thus fully type-checked at compile time, and their capabilities are on par with those of Go and Rust. In one aspect they are even superior, namely in that Swift allows for parameterized *methods*, instead of parameterized protocols. This has some interesting, positive implications for the Swift programmer, that will be discussed in detail below.

3.3.1 Encapsulated version coded in Swift

Listing 3.4 gives an example of how the encapsulated version of the array averaging test problem can be programmed in Swift. See the following remarks in order to understand this code:

- Swift uses angled brackets to indicate generic parameter lists.
- Generic type constraints are formulated by supplying a protocol name after a type parameter (separated by a colon).
- Similar to Rust, Swift provides a Numeric protocol that is defined in its standard library, for use as a generics constraint for numeric types.
- Swift’s own Numeric protocol does *not* include the division operation! Hence, the following code defines and uses a custom INumeric protocol, that extends the former, to require any adopting types to implement also the division operator.

- Swift's Int32 and Float64 types already implement, individually, both the functionality of Swift's own Numeric protocol, and the division operation. It therefore suffices to simply acknowledge, through empty extension statements, that they also conform to the combined functionality of the new INumeric protocol.
- The Swift version makes use of language built-in, default, structure constructors (called "initializers").
- Array slices are not arrays themselves. Hence, an explicit conversion using an Array() constructor is required in such cases.
- By default, function and method calls in Swift make use of keyword arguments.
- Type conversion into a generic type, T, is effected in Swift by calling an appropriate initializer of this type. The call syntax is a bit peculiar, i.e. Go's T(0) is written as T(exactly:0)! (making use of both the initializer's mandatory keyword, exactly, and the ! operator to unwrap the returned optional type).

Listing 3.4: Encapsulated Swift version of the array averaging example.

```

1 // .....
2 // Interfaces
3 // .....
4
5 protocol INumeric: Numeric {
6     static func / (lhs: Self, rhs: Self) -> Self
7 }
8
9 protocol ISum {
10     func sum<T: INumeric>(x: [T]) -> T
11 }
12
13 protocol IAverager {
14     func average<T: INumeric>(x: [T]) -> T
15 }
16
17 // .....
18 // Intrinsics
19 // .....
20
21 extension Int32: INumeric {}
22 extension Float64: INumeric {}
23
24 // .....
25 // SimpleSum ADT
26 // .....
27
28 struct SimpleSum: ISum {
29
30     func sum<T: INumeric>(x: [T]) -> T {

```

```

31     var s: T
32     s = T(exactly:0)!
33     for i in 0 ..< x.count {
34         s += x[i]
35     }
36     return s
37 }
38 }
39
40 // .....
41 // PairwiseSum ADT
42 // .....
43
44 struct PairwiseSum: ISum {
45     var other: ISum
46
47     func sum<T: INumeric>(x: [T]) -> T {
48         if ( x.count <= 2 ) {
49             return other.sum(x: x)
50         } else {
51             let m = x.count / 2
52             return sum(x: Array(x[..<m])) + sum(x: Array(x[m...]))
53         }
54     }
55 }
56 }
57
58 // .....
59 // Averager ADT
60 // .....
61
62 struct Averager: IAverager {
63     var drv: ISum
64
65     func average<T: INumeric>(x: [T]) -> T {
66         return drv.sum(x: x) / T(exactly: x.count)!
67     }
68 }
69
70 // .....
71 // main function
72 // .....
73
74 func main() {
75
76     let avs = Averager(drv: SimpleSum())
77     let avp = Averager(drv: PairwiseSum(other: SimpleSum()))
78

```

```

79     var av : IAverage = avs
80
81     let xi: [Int32] = [1,2,3,4,5]
82     let xf: [Float64] = [1.0,2.0,3.0,4.0,5.0]
83
84     var key: Int32?
85
86     print("Simple_sum_average:_1")
87     print("Pairwise_sum_average:_2")
88     print("Choose_an_averaging_method:_")
89     key = Int32(readLine()!)
90
91     switch key {
92     case 1:
93         // simple sum case
94         av = avs
95     case 2:
96         // pairwise sum case
97         av = avp
98     default:
99         print("Case_not_implemented!")
100        return
101    }
102
103    print( av.average(x: xi) )
104    print( av.average(x: xf) )
105 }
106
107 // execute main function
108 main()

```

Even a casual glance at the Swift version will show that the Swift code is the easiest to read and understand among the generic object-oriented implementations that were presented in this chapter. This is largely the result of Swift supporting generic methods, and hence not requiring the programmer to parameterize and instantiate any generic interfaces (protocols), in contrast to both Go and Rust. The consequences are

- that method genericity for an ADT's objects can be expressed using only a single, as opposed to multiple protocols,
- that therefore merely a *single* object instance of that same protocol is required, in order to be able to operate on many different language-intrinsic data types, and
- that this also (largely) *obviates the need for manual instantiations of generics in Swift* (because generic functions/methods are easier to instantiate automatically by the compiler, as it can almost always infer the required types by checking the regular arguments that are passed to a function/method)!

As an example, consider the object `av` in the above Swift code that contains the functionality for array averaging. This object supports two different levels of polymorphism: Firstly, given that

it is an instance of the `IAverager` protocol, it can be polymorphically assigned different averaging algorithms (see the `switch` statement). Secondly, because it contains an `average` method that is generic, it can be used on data of different intrinsic types, like `Int32` and `Float64` here.

Notice that the `main` function in the Swift code needs to declare merely a *single* such object variable of `IAverager` type, to make use of all these capabilities. This is a direct consequence of there being only a single (i.e. unparameterized) version of the `IAverager` protocol, and of parameterizing the protocol's method signatures by generic types rather than the protocol itself.

Contrast this with Go's and Rust's model, where not only separate objects of `IAverager` type are required for *every* different intrinsic data type that the programmer wishes to use these objects with. But where also *manual* instantiations of corresponding versions of the generically parameterized `IAverager` interface/trait are required from the programmer, for declaring these objects.

Swift's generics model gets rid of all of that complexity, and therefore vastly simplifies client code. We consider this a very significant advantage of the generics approach that is taken in Swift vs. that of Go and Rust.

3.4 Conclusions

The use of run-time polymorphism by means of (polymorphic) interfaces is rather similar in all the languages considered here. The most significant differences (that were not concretely explored here) appear to be that Go has stricter limitations on retroactively implementing interfaces for existing types than the other languages. Whereas Rust (with some minor restrictions), and Swift allow the implementation of an interface by some type to be accomplished independently from the type's definition site. Rust and Swift thereby overcome Haverdaen et al.'s critique [5] of Java regarding this point. In fact, it is *interface inheritance* which makes the uniform polymorphic treatment of both intrinsic and user-defined types possible in the first place in Rust and Swift, that Haverdaen et al. seem to also (rightly) demand. All the considered languages are also very similar in that they support fully type checked generics via the mechanism of interfaces. In the following, we will thus only summarize the most significant differences in these languages' generics features.

3.4.1 Go

Go's basic model to implement generics allows structures, interfaces, and ordinary functions, but not methods, to be given their own generic type parameters. The lack of true generic methods makes some duplication of instantiation code in clients unavoidable. Nevertheless, generic Go code is quite easy to read and to understand. What sets Go apart from the other languages is its built-in, easy to use support for conversion to generic types, and especially its brilliant new notion to interpret type sets as interfaces, along with its syntax to support this notion. This enables the Go programmer to easily tailor constraints on generic types to his specific needs, which is what makes the use of generics in Go pleasant. We consider these latter particular features of Go as "must haves" for Fortran.

3.4.2 Rust

Rust's basic model for generics is similar to Go's in that it allows for parameterization of structures, interfaces, and ordinary functions, but not necessarily methods. Hence, what has been said above for Go in this respect holds also for Rust. Rust has, unfortunately, some quirks which render its use for the management of all types of dependencies through polymorphism somewhat

sub-optimal when compared to the other languages considered here. The language is unpleasant to use, because of its “borrow checker”, its employment of move semantics by default, its *excessive* obsession with type safety, and its overall C++-like philosophy to copiously rely on external dependencies, even for the most basic tasks. The Rust version of our test case is therefore marred by a few dependencies on external libraries, which is quite contrarian to the purpose of programming in a polymorphic fashion, namely to avoid rigid dependencies. Even with the functionality provided by these external dependencies, type conversion to generic types within generic routines wasn’t possible. A necessary capability for numerical work that is, for instance, built into Go, but has to be achieved in Rust by some gymnastics on the programmer’s side. The points we like most about the language are its idea to decouple trait implementations from a struct’s definition through explicit `impl` blocks, and the complete control over the use of dynamic vs. static method dispatch (via trait objects and generics, respectively) that Rust affords the programmer. These are particular features of Rust that, in our opinion, Fortran should borrow in some form.

3.4.3 Swift

Swift’s basic model of implementing generics by allowing parameterized structures, functions, and methods (but not parameterized interfaces) is both the easiest to read, and the easiest to use from a programmer’s perspective. Swift’s generics design allows the Swift compiler to instantiate generics largely automatically, through type inference of the regular arguments that are passed to functions, methods, and (structure or class) constructors. In contrast to the other languages, in Swift, the user almost never has to bother with instantiating any generics.

If the Swift programmer knows how to write generic functions, his knowledge automatically translates into coding generic methods, since generic functions can be transformed into generic methods without requiring changes to their function signatures. This property is helpful for the refactoring of non-OO codes into corresponding OO versions.

We hence consider Swift’s generics to be the most attractive model to base Fortran’s basic generic capabilities on, provided that it can be implemented sufficiently easily. The fact that Swift is a language that does not put emphasis on numerics, and whose present standard library therefore does not provide a truly useful Numeric protocol (that supports all the usual numeric operations), is of absolutely *no* consequence for adopting Swift’s basic generics design for Fortran.

Fortran will necessarily do a better job in this respect, both by borrowing Go’s idea of formulating interfaces in terms of type sets, so that the user can easily implement his own type constraints. But also by making accessible to the user a set of language-built in interfaces that are truly useful for a variety of numeric operations, that are implemented by Fortran’s intrinsic types.

Chapter 4

Fortran extensions I: Traits and types

The present and the next chapter, describe a number of simple extensions to Fortran, that we consider essential in order to enable dependency management through polymorphism at a level of functionality that is on par with modern languages like Swift, Rust, Go, or Carbon. The present chapter adds general subtyping capabilities to Fortran, while the next chapter aims at providing specific support for generics. The extensions related to subtyping concern abstract interfaces, derived types, and the `class` specifier for variable declarations. They also encompass an extremely useful, new, Fortran feature: the `implements` statement.

4.1 Named abstract interfaces (traits)

The most important of all the following extensions is the capability to define named abstract interfaces, or traits (i.e. named collections of procedure signatures), in order to suitably constrain (and hence to type-check) the declarations of polymorphic variables. Named abstract interfaces are the crucial feature that is required in order to uniformly and properly express both run-time and compile-time polymorphism (i.e. modern-day OO and generic programming, respectively) in the language, and to thereby enable a uniform management of dependencies on both user-defined *and* language-intrinsic types.

4.1.1 Definition

Fortran already allows the programmer to define unnamed abstract interfaces. But in order to use these as traits, named versions of them are required. As in the following example, that defines three such named interfaces, `IAddable`, `ISubtractable`, and `IPrintable` that are intended as abstract blueprints for actual implementations of three type-bound procedures, named `add`, `sub`, and `output`, respectively:

```
1  abstract interface :: IAddable
2      function add(self,other) result(res)
3          deferred(self), intent(in) :: self, other
4          deferred(self)           :: res
5      end function add
6  end interface IAddable
7
8  abstract interface :: ISubtractable
9      function sub(self,other) result(res)
```

```

10      deferred(self), intent(in) :: self, other
11      deferred(self)              :: res
12  end function sub
13 end interface ISubtractable
14
15 abstract interface :: IPrintable
16   subroutine output(self)
17     deferred(self), intent(in) :: self
18   end subroutine output
19 end interface IPrintable

```

Notice, how in all these interfaces, the explicit passed-object dummy argument, `self`, but also all the other arguments that need to have the same type as `self`, are declared by means of the new deferred declaration specifier for associated types. Here, `deferred(self)` refers to the types that `self` can take on in actual implementations of these interfaces. This is discussed in more detail in Sect. 4.1.2. Since both named abstract interfaces, and the new deferred declaration specifier, are merely simple additions to Fortran, that aim to extend the present use cases of abstract interfaces in the language (which currently serve as bounds on the signatures of procedure pointers, and deferred, i.e. abstract, methods), these new features are fully backwards compatible with the present language.

4.1.2 Associated (deferred) types

Associated types (which are also available in Swift and Rust, but in contrast to Fortran are not required for the declaration of any explicit passed-object dummy arguments in these languages) are essentially aliases. They are employed within abstract interfaces/traits in lieu of types, whose actual value is not known at the time a trait is formulated, but will be known by a compiler once the programmer has actually implemented that trait for some concrete derived or intrinsic data type.

Consider, for instance, the interface `IPrintable` of Sect. 4.1.1, and suppose that a user provides, in his source code, an implementation of this interface/trait for some derived type (that we shall assume here to be extensible by type extension, and named `MyType`). Then the declaration

```

1  deferred(self), intent(in) :: self

```

of the passed-object dummy argument `self`, in `IPrintable`'s signature of subroutine `output`, would need to be matched by a declaration

```

1  class(MyType), intent(in) :: self

```

in the user's actual implementation of that subroutine. If the user would need to implement `IPrintable` also for Fortran's `real(real64)` type (see Sect. 4.2.4 for the details of how this can be accomplished), then that same declaration in `IPrintable` would instead need to be matched by a declaration

```

1  real(real64), intent(in) :: self

```

in `output`'s implementation. Thus, `deferred(self)` serves to stand in as an alias, in `IPrintable`, for both the `class(MyType)`, and `real(real64)` types, that object `self` takes on in subroutine `output`'s two different implementations. Since the concept of an associated type only makes sense within traits, the deferred declaration specifier for associated types can, accordingly, only be used within named abstract interfaces.

4.1.3 Extends attribute

Abstract interface definitions must allow the programmer to define new abstract interfaces that inherit procedure signatures from *multiple* simpler interfaces (multiple interface inheritance). In the following example, the interface `ICalculable` inherits the procedure signatures contained in both the interfaces `IAddable`, and `ISubtractable`, making `ICalculable` at the same time a *subtype* of both these simpler interfaces:

```
1  abstract interface, extends(IAddable,ISubtractable) :: ICalculable
2  end interface ICalculable
```

That is, objects that implement (or adopt) the `ICalculable` interface (i.e. conform to it), can also be used in settings that require conformance to either the `IAddable`, or `ISubtractable` interfaces.

The next example finally shows an interface, `IAdmissible`, that both inherits a procedure signature, namely that of function `add` which is contained in the `IAddable` interface, and adds a second function signature of its own (which is named `cast` here):

```
1  abstract interface, extends(IAddable) :: IAdmissible
2      function cast(self,i) result(res)
3          deferred(self), intent(in) :: self
4          integer,          intent(in) :: i
5          deferred(self)          :: res
6      end function cast
7  end interface IAdmissible
```

4.1.4 Generic overloading

Named abstract interfaces support “generic” overloading of both operators and procedures, with lists of (other) procedure names. This is accomplished by making use, within such an interface, of the generic statement that was introduced by Fortran 2018. The following example shows how to overload, within the `IAdmissible` interface of Sect. 4.1.3, Fortran’s operator(+) with the name of function `add`:

```
1  abstract interface, extends(IAddable) :: IAdmissible
2      function cast(self,i) result(res)
3          deferred(self), intent(in) :: self
4          integer,          intent(in) :: i
5          deferred(self)          :: res
6      end function cast
7      generic :: operator(+) => add
8  end interface IAdmissible
```

The effect of this is that any two objects, of some given type that implements the `IAdmissible` interface, may be added using either a direct call of their implemented `add` function, or an indirect call of that function, via the overloaded operator(+). See Sect. 6.3 for an actual example.

4.2 Implements statement

The language must allow not only for named abstract interfaces to conform to other named abstract interfaces (cf. Sect. 4.1.3), but also for both language-intrinsic and user-defined types to do

the same. That is, it must be possible, for *all* these types, to provide actual implementations of procedures, whose signatures are contained within named abstract interfaces. Moreover, this should be possible even retroactively, that is without having to touch any original type definitions. Otherwise, wrapper types would, in general, need to be written, once an already defined type would need to implement some new abstract interface (e.g. in cases where the original type definition is inaccessible). All of this means that user-provided methods need to be allowed also for *language-intrinsic* types, and that the provision of implementations should be possible regardless of any original type definition site¹.

The `implements` statement, that is described in this section, provides these capabilities. This statement is modeled after the “`extension`” feature of Swift, where it is used to enable retroactive implementation of new methods, additional constructors, and especially new protocols/interfaces for types, in order to *dynamically* change a subtyping (i.e. interface inheritance) hierarchy, and thus achieve utmost code flexibility. Swift’s `extension` blocks fulfill essentially the same purpose as Rust’s `impl` blocks. They have been somewhat simplified here (for a first implementation), and adjusted to Fortran’s idiosyncracies and syntax, that binds methods to types through declaration blocks, rather than by including the actual implementation bodies themselves into such a block (the implementations need to be supplied as module procedures, as is usual in Fortran).

The syntax of the feature is similar to that of the `contains` sections of derived type definitions. That is, `implements` statements make use of type-bound procedure declarations. The latter employ mostly the same attributes (i.e. `public`, `private`, `pass[(arg-name)]`, `nopass`, and `non_overridable`) which are also allowed in derived type definitions. “Generic” type-bound procedure declarations (as the Fortran standard calls them) can be used as well. The main difference to derived type definitions is that the `implements` statement cannot be used with abstract derived types, and that it hence also does not accept deferred type-bound procedure declarations. If interoperability with abstract types is required, the later to be discussed `implements` derived type attribute (cf. Sect. 4.4), that has none of these restrictions, can be used instead.

As a final disclaimer, notice that the `implements` statement has absolutely *no* relation to subclassing, i.e. one derived type being extended into another through (rigid) implementation inheritance. Rather, this is a feature that adds new capabilities to a single, *given* type.

4.2.1 Binding methods to a type

The simplest use of the `implements` statement is to bind some new methods to a type. The following example shows how to bind two methods, named `add` and `sub`, to a derived type with name `Calculable`, and how to provide their implementations as module procedures:

```
1 module basic
2   ...
3   type :: Calculable
4     private
5     real :: a
6   end type Calculable
7
8   implements :: Calculable
9     procedure, public, pass(self) :: add, sub
10  end implements Calculable
11
```

¹Exempting only the case of abstract derived types, as it is discussed below.

```

12 contains
13
14 function add(self,other) result(res)
15     class(Calculable), intent(in) :: self, other
16     class(Calculable)           :: res
17     res%a = self%a + other%a
18 end function add
19
20 function sub(self,other) result(res)
21     class(Calculable), intent(in) :: self, other
22     class(Calculable)           :: res
23     res%a = self%a - other%a
24 end function sub
25
26 end module basic

```

4.2.2 Implementing a trait

Assume now, that the purpose of binding the previous two methods to the `Calculable` type was to actually make this type compatible with settings where conformance to the `ICalculable` interface of Sect. 4.1.3 is required. So far, we have provided both the binding of the required methods, and their implementations, but we haven't made `Calculable` pluggable yet into code that is written in terms of the `ICalculable` interface. To fix this, we can simply acknowledge (even from a different module, as in the following example), that the `Calculable` type already has all of the required functionality to implement the `ICalculable` interface:

```

1 module enhanced
2     ...
3     use basic, only: Calculable
4
5     implements (ICalculable) :: Calculable
6     end implements Calculable
7
8 end module enhanced

```

4.2.3 Implementing multiple traits

It is crucial, for flexibility, that the subtyping (i.e. interface inheritance) mechanism, that is provided through the `implements` statement, allow for a type to implement *multiple* different interfaces. For instance, if one wouldn't have defined the interface `ICalculable`, and would nevertheless need to use objects of type `Calculable` in settings that require conformance to either one of the `IAddable` or `ISubtractable` interfaces of Sect. 4.1.1, then the language must allow one to acknowledge `Calculable`'s conformance to these interfaces using a (comma-separated) list of interface names to the `implements` statement, as follows:

```

1 module enhanced
2     ...
3     use basic, only: Calculable

```

```

4
5  implements (IAddable,ISubtractable) :: Calculable
6  end implements Calculable
7
8 end module enhanced

```

The last two examples have demonstrated how interfaces can actually be implemented after their necessary method implementations were already bound to some type. Of course, it is also possible to do all of this at one fell swoop, as in the following alternative example of the aforegiven module basic:

```

1 module basic
2   ...
3   type :: Calculable
4     private
5       real :: a
6   end type Calculable
7
8   implements (IAddable,ISubtractable) :: Calculable
9     procedure, public, pass(self) :: add, sub
10  end implements Calculable
11
12 contains
13
14  function add(self,other) result(res)
15    class(Calculable), intent(in) :: self, other
16    class(Calculable)              :: res
17    res%a = self%a + other%a
18  end function add
19
20  function sub(self,other) result(res)
21    class(Calculable), intent(in) :: self, other
22    class(Calculable)              :: res
23    res%a = self%a - other%a
24  end function sub
25
26 end module basic

```

We could have also employed two separate implements statements, in order to implement one interface at a time, like so

```

1 module basic
2   ...
3   type :: Calculable
4     private
5       real :: a
6   end type Calculable
7
8   implements IAddable :: Calculable
9     procedure, public, pass(self) :: add

```



```

10  end implements Calculable
11
12  implements ISubtractable :: Calculable
13      procedure, public, pass(self) :: sub
14  end implements Calculable
15
16  contains
17      ...
18  end module basic

```

where we have skipped, for brevity, the implementation of the actual methods, that would be done exactly as in the preceding example. The effect would have been the same. Such splitting of `implements` statements can be useful to improve code readability, as it makes the association between the interfaces and the actual procedures, that are to be implemented for any one of them, immediately obvious. These two statements (together with the actual implementations), could then have been distributed even among different modules and files. Notice, also, how parentheses around interface lists in `implements` statements are optional, but not required.

4.2.4 Implementing a trait for an intrinsic type

The `implements` statement can be used to provide methods, and to implement traits, also for types that are *intrinsic* to the language – much in the same fashion as it was already demonstrated for derived types. The following listing shows how to code the Swift example of Listing 3.3, by implementing a single method named `output`, that performs printouts for variable instances of Fortran’s `real(real64)` type:

```

1  module real64_module
2
3      use, intrinsic :: iso_fortran_env, only: real64
4
5      abstract interface :: IPrintable
6          subroutine output(self)
7              deferred(self), intent(in) :: self
8          end subroutine output
9      end interface IPrintable
10
11     implements IPrintable :: real(real64)
12         procedure :: output
13     end implements real(real64)
14
15     contains
16
17     subroutine output(self)
18         real(real64), intent(in) :: self
19         write(*,*) "I_am_", self
20     end subroutine output
21
22 end module real64_module
23

```

```

24 program printx
25
26   use, intrinsic :: iso_fortran_env, only: real64
27
28   real(real64) :: x
29
30   x = 4.9d0
31   call x%output()
32
33 end program printx

```

Notice, that once an `implements` statement has been used to augment the functionality of an existing (intrinsic or derived) type (e.g. the `real(real64)` type in the above example), then this new functionality will be available on *all* instances of that type, even if they were created before this augmentation was defined. As it is the case in the Swift language [16], this prevents the occurrence of any ambiguities regarding the available functionality of such type instances. For the same reason, it is sufficient to import into client code (as usual) either a named constant corresponding to an intrinsic type’s kind parameter (like the `real64` constant in the above `printx` example program), or a derived type’s (original) definition, in order to make use of the thus augmented functionality.

4.3 Declaration of run-time polymorphic variables (trait objects)

Named abstract interfaces/traits are types in their own right. Their main purpose is to allow the programmer to declare (and thereby constrain) polymorphic variables in terms of them. Either directly, i.e. as objects of such interfaces/traits in run-time polymorphism, as demonstrated in this section. Or indirectly, as constraints on generic type parameters in compile-time polymorphism (see Sects. 5.1.1 and 5.3 for examples of the latter).

4.3.1 Class specifier using a single trait

In order to use abstract interfaces to support run-time polymorphic objects through subtyping, Fortran’s `class` specifier for polymorphic variable declarations needs to be enhanced to accept named abstract interfaces/traits, like in the declarations of the following two variables (that make use of the `IAddable` interface of Sect. 4.1.1)

```

1 class(IAddable), allocatable :: adder
2 class(IAddable), pointer :: adderptr

```

or the following declaration of a procedure argument:

```

1 class(IAddable), intent(inout) :: adder

```

The semantics here are that whenever a named abstract interface appears within the `class` specifier of an object’s declaration, then all the public methods of that object whose signatures are prescribed by the adopted interface (like `IAddable` in the above examples), will make use of late binding. That is, their calls will be resolved by the run-time system of the language (e.g. through a virtual method table).

In accordance with how objects that make use of run-time polymorphism through subclassing (i.e. implementation inheritance) are declared in the present Fortran standard, also “trait objects”

(like `adder`, and `adderptr` in the examples above) must either be declared using the `allocatable`, or the pointer attribute, or they must be arguments of a procedure. The proposed additions are therefore backwards compatible with the functionality that is already available in the present language. See also Sect. 5.1.2 for further information on the class specifier when used with abstract interfaces.

4.3.2 Class specifier using trait combinations

The `ICalculable` interface of Sect. 4.1.3 derives from the `IAddable` and `ISubtractable` interfaces that were discussed in Sect. 4.1.1. Objects that require the combined functionality of both these latter interfaces could thus be declared in terms of the `ICalculable` interface as follows:

```
1 class(ICalculable), allocatable :: addsub
```

It happens often, though, that one needs to declare objects that conform to multiple interfaces, but where one would like to avoid having to code some (otherwise unneeded) intermediary interface, that derives from these. In such cases, it should be possible to express an object's declaration more directly, by providing a list of interface names to the class specifier that expresses a traits combination, as in the following example

```
1 class(IAddable,ISubtractable), allocatable :: addsub
```

in which object `addsub` again conforms to both the `IAddable` and `ISubtractable` traits.

4.4 Derived types

For reasons of both regularity in the language, and full interoperability with Fortran's legacy (i.e. present) OO model, it should be possible to implement interfaces also directly from within derived type definitions. This can be accomplished by allowing an `implements attribute` within such definitions, that – in contrast to the previously discussed `implements` statement – features also interoperability with abstract derived types. Notice, however, that by its nature of being bound to derived type definitions, the former attribute (necessarily) lacks the latter statement's retroactive implementation capabilities, that are important for fully exploiting the flexibility that the modern traits-based OO programming model offers.

4.4.1 Implements and sealed attributes

The following code shows how the “one fell swoop” example of Sect. 4.2.3, that was written there in terms of an `implements` statement, can be reformulated to use an `implements` derived type attribute:

```
1 module basic
2   ...
3   type, sealed, implements(IAddable,ISubtractable) :: Calculable
4     private
5     real :: a
6   contains
7     procedure, public, pass(self) :: add, sub
8   end type Calculable
9
```

```

10 contains
11
12 function add(self,other) result(res)
13     type(Calculable), intent(in) :: self, other
14     type(Calculable)           :: res
15     res%a = self%a + other%a
16 end function add
17
18 function sub(self,other) result(res)
19     type(Calculable), intent(in) :: self, other
20     type(Calculable)           :: res
21     res%a = self%a - other%a
22 end function sub
23
24 end module basic

```

This example also demonstrates the use of another, new, derived type attribute, the sealed attribute. Derived types that are sealed are inextensible by type extension, i.e. implementation inheritance. Thus, there's also no need for the passed-object dummy arguments of such types to be polymorphic, and hence for the programmer to declare them with the class specifier. One can use the type specifier, instead, as it is demonstrated in the implementation of the add and sub methods of this example.

4.4.2 Implements and extends attributes

The previous Sect. 4.4.1 has illustrated how the implements derived type attribute would be typically used in modern code (that avoids the use of implementation inheritance, or subclassing). However, the multiple interface inheritance (or subtyping) features, that are provided by the implements derived type attribute², can actually be combined with the type extension capabilities which are already present in the language. For instance, code examples like the following are possible:

```

1  type :: Parent
2  contains
3      procedure :: method1
4      procedure :: method2
5  end type Parent
6
7  type, extends(Parent), implements(IChild) :: Child
8  contains
9      procedure :: method3
10     procedure :: method4
11 end type Child
12 ...

```

Here, a Child type is defined, that inherits two methods (method1 and method2) from a Parent type, and adds two further methods of its own (method3 and method4), in order to conform to an

²And, with some restrictions, also the implements statement.

interface, `ICChild`, that consists of the signatures of all four of these methods. In such use cases, the `implements` attribute shall always precede the `implements` attribute.

Since the `implements` derived type attribute allows for inheritance of *multiple* interfaces, this also fixes present Fortran's single inheritance limitations, *without* introducing the potential ambiguities that multiple inheritance of implementation would cause (which are also known as "The Diamond Problem").

4.4.3 Implements attribute and abstract derived types

The examples of the `implements` derived type attribute, that were presented in the previous two Sects. 4.4.1 and 4.4.2, could have been coded equally well using `implements` statements. Yet, the `implements` attribute provides one capability that the `implements` statement lacks, namely implementation of interfaces even if the implementing derived type is abstract. In such cases, the implementing type is allowed to provide an only partial implementation of the interface(s) that it adopts. Any non-abstract type that extends this abstract type through type extension must, though, provide a full implementation.

For instance, it is valid for the following abstract type, which is named `Base`,

```
1 type, abstract, implements(Interface) :: Base
2 contains
3   ! declarations for all but one method of Interface
4   ...
5 end type Base
6 ! method implementations
7 ...
```

to implement all but one of the methods of the specific `Interface` that it adopts. On the other hand, the following type

```
1 type, extends(Base) :: Extended
2 contains
3   ! declaration for the remaining method of Interface
4   ...
5 end type Extended
6 ! method implementation
7 ...
```

that extends the `Base` type, must then implement the `Interface`'s remaining method. Thus, the `implements` attribute makes interface inheritance interoperable even with abstract derived types and their type extensions (and hence with present Fortran's full OO model), in a manner that is equivalent to how such interoperability works in the Java language [15].

Notice, however, that if the programmer decides to forego the use of abstract derived types altogether, then the previously discussed `implements` *statement* offers not only all of the capabilities of the `implements` attribute, but also adds to these the possibility of retroactive implementation. Since the thus resulting programming style is a much more flexible, complete, modern-day alternative to both abstract types and type extension, our recommendation would be to use the interoperability features that are described in Sect. 4.4.2 and the present Sect. 4.4.3, only within legacy applications, and to otherwise avoid the use of both abstract types and type extension in new code.

Chapter 5

Fortran extensions II: Generics

The new subtyping features that were discussed in the previous chapter are required in order to uniformly express and support both run-time and compile-time polymorphism in Fortran. We will now proceed with discussing further enhancements that are specifically needed in order to support compile-time polymorphism, i.e. generics.

5.1 Enhancements to interfaces

5.1.1 Generic procedure signatures

Abstract interfaces should be allowed to contain signatures of generic procedures, as in the Swift language. The approach taken in Go and Rust, to instead parameterize the abstract interfaces themselves, is not as attractive from a user’s perspective (cf. Sect. 3.3). The following code shows, as an example, an abstract interface called `ISum` that contains the signature intended for a generic type-bound procedure (i.e. generic method), named `sum`:

```
1  abstract interface :: ISum
2      function sum{INumeric :: T}(self,x) result(s)
3          deferred(self), intent(in) :: self
4          type(T),          intent(in) :: x(:)
5          type(T)              :: s
6      end function sum
7  end interface ISum
```

The example illustrates the use of a generic type parameter, that is simply called `T` here, in terms of which the regular function arguments are declared. A significant difference of generic type parameters, as compared to regular function arguments, is that the former will be substituted by an actual type argument at compile time, in a process called instantiation.

A similarity is that, in the same way that regular function arguments need to be constrained by a provided type, type parameters need to be constrained by a provided meta-type. This (meta-type) constraint must either be a single abstract interface name (like `INumeric` in the present example), or a comma separated list of abstract interface names that expresses a traits combination (whose use was already illustrated in the run-time polymorphic example of Sect. 4.3.2). The proposed Fortran generics thus support “strong concepts”, and can be fully type-checked by the compiler.

Both, the type parameter and the constraint which precedes it, and is separated from it by a double colon, are part of a generic type parameter list that is enclosed in curly braces, and follows

immediately behind the procedure's name. Notice that the syntax that is used here, which deviates slightly from how Fortran's regular function arguments are declared, appears justified, as it reflects that, despite some similarities, in type parameters one is dealing with different entities.

5.1.2 Type sets

In order to make the interface based generics facility easy to use together with intrinsic types (and the multitude of intrinsic procedures that Fortran supports for these types), it must be possible, as in the Go language, to define generics constraints by means of abstract interfaces that are expressed in terms of type sets.

Unions of types

The following example shows the simplest form of such a type set, by defining an interface `INumeric`, for use as a generics constraint in the example of Sect. 5.1.1, in order to admit for the type parameter `T`, that was given there, only the (32 bits wide) default integer intrinsic data type:

```
1  abstract interface :: INumeric
2      integer
3  end interface INumeric
```

The above example is actually a special case of specifying entire *unions* of member types as a type set. A type set consisting of such a union of types is demonstrated in the following example

```
1  abstract interface :: INumeric
2      integer | real(real64)
3  end interface INumeric
```

that redefines interface `INumeric` such as to admit either the default 32 bit integer, or the 64 bit real type as a generics constraint.

The semantics of such a type set construct are that it implicitly defines a *set of function signatures*, namely the signatures of the intersecting (common) set of all the operations and intrinsic functions (also called methods in the following) that work with all the member types of the type set. This can also be restated, by saying that a type `T` *implements* an interface consisting of such a type set, if (and only if) it is a member type of this set. For instance, the `integer` and `real(real64)` types implement the interface `INumeric` (as it is given above) because they are member types of its type set. In contrast, Fortran's various complex types do not implement this interface, because they do not belong to its set of member types. The validity of the latter statement can also be easily checked, by considering that the complex types do not support, i.e. implement, the relational operators (`<`) and (`>`). Implementations of the latter are required for conformance to this interface, because these operators *are* implemented by both the `integer` and `real(real64)` member types.

Assumed kind parameters

Expanding on the previous example, an `INumeric` interface that might be even more useful as a generics constraint, for a number of tasks, could be coded as follows:

```
1  abstract interface :: INumeric
2      integer(*) | real(*) | complex(*)
3  end interface INumeric
```

Notice how this makes use of both unions of types, and assumed (i.e. wildcard) kind parameters for types, to include *all* integer, real, and complex types, that are admitted by the language, in a single abstract interface constraint.

The use of assumed kind parameters is here merely syntactic sugar that allows one to avoid having to write out a type set for all the possible kinds of a type. For instance, if the particular Fortran implementation supports `real(real32)` and `real(real64)` as its only real types, then `real(kind=*)`, or `real(*)` for short, is understood to mean the type set “`real(real32) | real(real64)`”. Notice, also, that the more types are added to an interface in this fashion, the smaller the set of intersecting methods will usually become.

Empty interface

In the limit of adding all possible types to a type set, there won’t be any common methods left that are implemented by all its types. This results in the important case of the empty interface, that matches all types (since any type has at least zero methods):

```
1  abstract interface :: IAnyType
2  end interface IAnyType
```

Implicit notation

For simple use cases, it should be optionally possible for the programmer to employ a shorter notation for declaring type constraints for generics, than having to explicitly define an interface of type sets, like `INumeric` above, and to then use it as in Sect. 5.1.1. The following modification of interface `ISum`’s original declaration of Sect. 5.1.1, provides such an example:

```
1  abstract interface :: ISum
2      function sum{integer | real(real64) :: T}(self,x) result(s)
3          deferred(self), intent(in) :: self
4          type(T),          intent(in) :: x(:)
5          type(T)           :: s
6  end function sum
7  end interface ISum
```

The notation within the generic type parameter list in curly braces defines a type set interface implicitly, to be used as a type constraint for type `T`. In this particular case, to admit only the 32 bit integer, or 64 bit real type for `T`, as discussed above.

Predefined constraints

The facilities that are described in this section are flexible enough for the user to be able to construct frequently required generics constraints himself, the way he needs them. Nevertheless, the language should ideally also supply a collection of predefined, commonly used generics constraints in the form of abstract interfaces that are contained in a language-intrinsic module, tentatively called `generics_constraints` here. The list of such predefined interfaces could include

- an empty interface of the name `IAnyType` (as shown above),
- some predefined numeric interfaces allowing for different numeric operations, but also

- some predefined interfaces to allow for the use of relational operators with different intrinsic types.

Such interfaces could then be imported from user code through a use statement like in the following example that assumes the existence of a language defined interface `INumeric`:

```

1 module user_code
2
3   use, intrinsic :: generics_constraints, only: INumeric
4
5   abstract interface :: ISum
6     function sum{INumeric :: T}(self,x) result(s)
7       deferred(self), intent(in) :: self
8       type(T),          intent(in) :: x(:)
9       type(T)           :: s
10    end function sum
11  end interface ISum
12
13 end module user_code

```

Present restrictions and possible future extensions

Interfaces that are formulated in terms of type sets are presently *exclusively* intended for use as generics constraints. Hence, a compiler must ensure that they are *not* used for any other purpose. In particular, they are not intended to be implemented by derived (i.e. user-defined) types¹, and they cannot be used in variable declarations that involve the class specifier (as described in Sect. 4.3).

As it was demonstrated by the aforegiven examples, interfaces that are formulated in terms of type sets typically boil down to a set of language-intrinsic types. Since the class declaration specifier is undefined for intrinsic types, and since – moreover – the semantics of this specifier allow for late binding of methods, which is incompatible with intrinsic types, a compiler will need to ensure that the class declaration specifier is not used in conjunction with interfaces that are formulated in terms of type sets. This includes the empty interface, `IAnyType`.

In a future language revision, such interfaces could, however, be admitted for use with the type declaration specifier, in order to enable compile-time polymorphism through union (also called sum) types [13, 14]. This would take the present design to its logical conclusion, offer an alternative to generic type parameters for certain use cases, and fill a present gap in the later to be discussed Table 6.1.

5.2 Generic casts for intrinsic types

A language like Fortran, that is intended for numeric use, where conversions between different intrinsic types are required rather frequently, must allow conversions between such types to be done also in generic code, in a similarly user-friendly fashion as it is the case in the Go language.

¹One of the problems here is that any new intrinsic function that would be added to the language for some intrinsic type, would change the set of methods of all the type sets of which this type is a member. This would break any user-defined types that would implement interfaces which are based on these type sets.

That is, by means that are built into the language, without having to rely exclusively on external library functionality.

For instance, generic routines will often have to initialize the result of reduction operations, as it is, e.g., the case in the test problem implementation of Sect. 2.2. There, a reduction variable for summation, *s*, needs to be initialized to the zero constant of type *T*. In Go, it is easily possible to express this initialization by transforming the (typeless) zero constant, 0, into the corresponding constant of type *T*, i.e. by simply writing *s* = *T*(0).

If, for instance, *T* is then instantiated at compile time with the `float64` type, the expression *T*(0) will be transformed by the compiler into `float64(0)`, i.e. a call to the correct conversion function. In Fortran’s case, the compiler would have to translate the above expression into the intrinsic function call `real(0,kind=real64)`, which should actually be easy to do, also for all other cases where such conversion is indeed possible. Otherwise, the compiler should emit an error message, and abort compilation at the generics instantiation step.

5.3 Generic procedures, methods, and derived types

As already mentioned, Fortran’s basic generics design should allow both ordinary and type-bound procedures (i.e. methods), and derived types to be given their own generic type parameters.

5.3.1 Generic procedures

Using the syntax proposed in this document, an implementation of a Fortran function for array summation that is generic over the type of its input array argument would look as follows²:

```
1  function sum{INumeric :: T}(x) result(s)
2      type(T), intent(in) :: x(:)
3      type(T)                :: s
4      integer :: i
5      s = T(0)
6      do i = 1, size(x)
7          s = s + x(i)
8      end do
9  end function sum
```

The following examples will assume that the `INumeric` generics constraint, that is used here, is provided by an interface that admits both the `integer` and `real(real64)` types for *T* (see Sect. 5.1.2).

To actually use the *sum* generic function, one simply needs to pass to it (via its regular arguments list) an argument of one of the admitted “numeric” types, e.g. as in the following two calls:

```
1  integer      :: integer_total
2  real(real64) :: float_total
3
4  integer_total = sum([1,2,3,4,5])
5  float_total   = sum([1.d0,2.d0,3.d0,4.d0,5.d0])
```

Here, the compiler will *automatically* instantiate appropriate versions of the *sum* generic function by using type inference on the regular arguments lists. That is, the programmer can straightforwardly use his generic routine on any of the admitted types.

²Generic subroutines can be coded completely analogously.

Although the present design *doesn't* typically require it, the programmer can accomplish generic instantiation also manually, by the additional provision of generic type arguments in curly braces, as in the following two calls:

```
1 integer_total = sum{integer}([1,2,3,4,5])
2 float_total   = sum{real(real64)}([1.d0,2.d0,3.d0,4.d0,5.d0])
```

Such manual instantiation of generic procedures is only needed in the relatively rare cases where a regular arguments list is either unavailable or outright inappropriate, e.g. for use in procedure pointer assignments, or in associate statements, as in the next example:

```
1 real(real64) :: dtot(2)
2 real(real32) :: stot(2)
3 procedure(sum_real64), pointer :: dsum
4
5 abstract interface
6     function sum_real64(x) result(s)
7         real(real64), intent(in) :: x(:)
8         real(real64)              :: s
9     end function sum_real64
10 end interface
11
12 dsum => sum{real(real64)}
13
14 dtot(1) = dsum([1.d0,2.d0,3.d0,4.d0,5.d0])
15 dtot(2) = dsum([2.d0,4.d0,6.d0,8.d0])
16
17 associate( ssum => sum{real(real32)} )
18     stot(1) = ssum([1.,2.,3.,4.,5.])
19     stot(2) = ssum([2.,4.,6.,8.])
20 end associate
```

5.3.2 Generic methods

The same summation algorithm as that of Sect. 5.3.1, when implemented as a generic method, `sum`, that is bound to a derived type named `SimpleSum`, which implements the interface `ISum` as it was given in Sect. 5.1.1, would instead look as follows:

```
1 module simple_library
2     ...
3
4     type, public :: SimpleSum
5     end type SimpleSum
6
7     implements ISum :: SimpleSum
8         procedure :: sum
9     end implements SimpleSum
10
11 contains
12
```

```

13  function sum{INumeric :: T}(self,x) result(s)
14      class(SimpleSum), intent(in) :: self
15      type(T),           intent(in) :: x(:)
16      type(T)
17      integer :: i
18      s = T(0)
19      do i = 1, size(x)
20          s = s + x(i)
21      end do
22  end function sum
23
24  end module simple_library

```

We have used here the *implements statement* for implementing an interface by a derived type, and will continue to consistently do so for the remainder of this chapter. The alternative way of employing the *implements derived type attribute* for the same purpose, is demonstrated on the same examples in Chapter 6.

Generic methods are used completely analogously to generic procedures. The automatic instantiation use case of Sect. 5.3.1 would, for instance, be written

```

1  type(SimpleSum) :: simple
2
3  integer_total = simple%sum([1,2,3,4,5])
4  float_total  = simple%sum([1.d0,2.d0,3.d0,4.d0,5.d0])

```

whereas the corresponding manual instantiation case would take the form:

```

1  type(SimpleSum) :: simple
2
3  integer_total = simple%sum{integer}([1,2,3,4,5])
4  float_total  = simple%sum{real(real64)}([1.d0,2.d0,3.d0,4.d0,5.d0])

```

5.3.3 Generic derived types

In addition to procedures and methods, generic type parameter lists must also be allowed for derived type definitions, as in the following example, in which the *ISum* interface of Sect. 5.1.1 is implemented by another derived-type, named *PairwiseSum*:

```

1  module pairwise_library
2      ...
3      type, public :: PairwiseSum{ISum :: U}
4          private
5              type(U) :: other
6          end type PairwiseSum
7
8      implements ISum :: PairwiseSum{ISum :: U}
9          procedure :: sum
10     end implements PairwiseSum
11
12 contains

```

```

13
14 function sum{INumeric :: T}(self,x) result(s)
15     class(PairwiseSum{U}), intent(in) :: self
16     type(T),                intent(in) :: x(:)
17     type(T)                  :: s
18     integer :: m
19     if (size(x) <= 2) then
20         s = self%other%sum(x)
21     else
22         m = size(x) / 2
23         s = self%sum(x(:m)) + self%sum(x(m+1:))
24     end if
25 end function sum
26
27 end module pairwise_library

```

Notice, how type `PairwiseSum` depends on a generic type parameter, `U`, that is used within type `PairwiseSum` itself in order to declare a field variable of type `U`, which is named `other`. As it is indicated by the type constraint on `U`, object `other` conforms to the `ISum` interface itself, and therefore contains its own implementation of the `sum` procedure.

The above example furthermore demonstrates, how a derived type's generic parameters are brought into the scope of its type-bound procedures via the latter's passed-object dummy arguments. In this example, type `PairwiseSum`'s method, `sum`, has a passed-object dummy argument, `self`, that is declared being of `class(PairwiseSum{U})`. Hence, method `sum` can now access `PairwiseSum`'s generic parameter `U`. This allows the method to make use of two independently defined generic type parameters, `T` and `U`, which grants it increased flexibility. This also means that there is *no* implicit mechanism of bringing generic parameters of a derived type into the scope of its methods. If a type-bound procedure needs to access the generic parameters of its derived type, it must be provided with a passed-object dummy argument.

Notice, that the declaration `class(PairwiseSum{U})` does not imply any ambiguities or contradictions with respect to compile-time vs. run-time polymorphism, because substitution semantics apply. At compile time, the compiler will substitute a set of different type arguments for the generic parameter `U`. Hence, the notation `PairwiseSum{U}` really refers to a set of multiple, related, but *different* `PairwiseSum` types, whose *only* commonality is that they all implement the `ISum` interface (and furthermore contain different field components that do the same). Of course, passed-object dummy arguments of any of the different `PairwiseSum` types of this set can then be run-time polymorphic, in exactly the same manner that passed-object dummy arguments of other derived types that implement the same interface can be run-time polymorphic.

The following code snippet finally shows how an object of the `PairwiseSum` type could be declared and instantiated statically (by substituting its generic type parameter `U` by the `SimpleSum` type of Sect. 5.3.2), and how its generic `sum` method would be employed using automatic type inference:

```

1  type(PairwiseSum{SimpleSum}) :: pairwise
2
3  integer_total = pairwise%sum([1,2,3,4,5])
4  float_total   = pairwise%sum([1.d0,2.d0,3.d0,4.d0,5.d0])

```

5.4 Updated structure constructors

If a derived type is parameterized over a generic type, as in the example of Sect. 5.3.3, then its structure constructor must also be assumed to be parameterized over the same generic type. Hence, calls of structure constructors, with specific argument types substituting the generic type parameters of their derived types, must be valid. The following (in this particular case redundant, as demonstrated above) run-time instantiation of object `pairwise` by the structure constructor of type `PairwiseSum` provides such an example:

```
1  type(PairwiseSum{SimpleSum}) :: pairwise
2
3  pairwise = PairwiseSum{SimpleSum}()
4
5  integer_total = pairwise%sum([1,2,3,4,5])
6  float_total   = pairwise%sum([1.d0,2.d0,3.d0,4.d0,5.d0])
```

As in Sect. 5.3.2, `SimpleSum` is a derived type that implements the `ISum` interface, but (in contrast to the `PairwiseSum` type, as implemented in Sect. 5.3.3) is not parameterized by any generic type parameter itself.

We also propose to make a further, small, but extremely important addition to structure constructors: namely to introduce the notion that a structure constructor is implicitly defined *within the same scope that hosts the definition of its derived type*. This would make it possible for the structure constructor to access even private components of its derived type, through host association. Meaning that if the derived type is hosted within the specification part of a module, such components would be accessible, and thus could be initialized, even by calls to the structure constructor that are being performed from outside this module's scope. Which would be in complete analogy to how user-defined constructors work in Fortran.

In this way, it would become possible to initialize private, allocatable derived type components by structure constructors, which is absolutely crucial for concise OO programming. Since such an extension would merely add to the capabilities of the language, it would be fully backwards compatible. The elegance of the aforementioned Go and Swift code versions, but also of the Fortran code examples that are presented in the next chapter is largely due to the use of such constructors. Lacking these, user-defined constructors would have to be employed, leading to overly complex implementations, as e.g. in the Rust example code given in Listing 3.2.

5.5 Extensibility to rank genericity

Fortran's special role, as a language that caters to numerical programming, demands that any generics design for the language must allow for the possibility to also handle genericity of array rank. The present design offers a lot of room in this respect, but since this is a largely orthogonal issue, and since we consider a discussion of such functionality to be non-essential for the purpose of a very first prototype implementation of the generics features that are described here, we defer it to a separate document.

Chapter 6

Fortran examples

In order to comprehensively illustrate how the new features, that were discussed in the last two chapters, would be used in practice, we will give in the present chapter several worked out examples. The first two subsections contain both complete functional and OO Fortran code versions of the standard test problem that is used throughout this document. While the last two subsections deal with how to formulate general generics constraints, and how to use associated types with generics.

6.1 Functional versions of the standard test problem

Fortran presently lacks support for advanced functional programming capabilities, like closures and variables of higher-order functions, that are, e.g., available in Go and other modern languages. In contrast to the Go version of the standard test problem that is given in Sect. 2.2, the functional Fortran code versions that are presented in this section therefore make no attempt to eliminate rigid dependencies on user-defined function implementations, and content themselves with demonstrating how the new generics features can be used to eliminate rigid dependencies on language-intrinsic types.

6.1.1 Automatic instantiation of generic procedures

Listing 6.1 shows a straightforward generic functional implementation of the test problem, that uses automatic type inference by the compiler. The following additional remarks should help to better understand this code:

- To express type genericity for the arguments and return values of our different generic functions, we make use of a type constraint expressed by the interface `INumeric`, that is implemented as the type set `integer | real(real64)`.
- Interface `INumeric` is defined by the user himself. Thus, there is no need for an external dependency.
- Any required conversions to generic types are done using explicit casts, as in Go.
- *All* the required instantiations of generic procedures are done automatically by the compiler, based on type inference of the regular arguments that are passed to these procedures.

Listing 6.1: Fortran version of the array averaging problem with automatic generics instantiation.

```

1  module averager_library
2
3      use, intrinsic :: iso_fortran_env, only: real64
4
5      implicit none
6      private
7
8      public :: simple_average, pairwise_average
9
10     abstract interface :: INumeric
11         integer | real(real64)
12     end interface INumeric
13
14 contains
15
16 function simple_sum{INumeric :: T}(x) result(s)
17     type(T), intent(in) :: x(:)
18     type(T)                :: s
19     integer :: i
20     s = T(0)
21     do i = 1, size(x)
22         s = s + x(i)
23     end do
24 end function simple_sum
25
26 function pairwise_sum{INumeric :: T}(x) result(s)
27     type(T), intent(in) :: x(:)
28     type(T)                :: s
29     integer :: m
30     if (size(x) <= 2) then
31         s = simple_sum(x)
32     else
33         m = size(x) / 2
34         s = pairwise_sum(x(:m)) + pairwise_sum(x(m+1:))
35     end if
36 end function pairwise_sum
37
38 function simple_average{INumeric :: T}(x) result(a)
39     type(T), intent(in) :: x(:)
40     type(T)                :: a
41     a = simple_sum(x) / T(size(x))
42 end function simple_average
43
44 function pairwise_average{INumeric :: T}(x) result(a)
45     type(T), intent(in) :: x(:)
46     type(T)                :: a

```



```

47     a = pairwise_sum(x) / T(size(x))
48     end function pairwise_average
49
50 end module averager_library
51
52
53 program main
54
55     ! dependencies on intrinsic constants
56     use, intrinsic :: iso_fortran_env, only: real64
57
58     ! dependencies on implementations
59     use averager_library, only: simple_average, pairwise_average
60
61     implicit none
62
63     ! declarations
64     integer,      parameter :: xi(5) = [1, 2, 3, 4, 5]
65     real(real64), parameter :: xf(5) = [1.d0, 2.d0, 3.d0, 4.d0, 5.d0]
66
67     integer :: key
68
69     write(*,'(a)') 'Simple__sum_average:_1'
70     write(*,'(a)') 'Pairwise_sum_average:_2'
71     write(*,'(a)',advance='no') 'Choose_an_averaging_method:_ '
72     read(*,*) key
73
74     select case (key)
75     case (1)
76         print '(i8)',    simple_average(xi)
77         print '(f8.5)',  simple_average(xf)
78     case (2)
79         print '(i8)',    pairwise_average(xi)
80         print '(f8.5)',  pairwise_average(xf)
81     case default
82         stop 'Case_not_implemented!'
83     end select
84
85 end program main

```

The example demonstrates that using the new generics features together with a functional (or procedural) programming style is easy, that the syntax is economical, and that type inference by the compiler should be straightforward and therefore reliable. Hence, we believe that the generics features, that are described here, will place no burden on the programmer.

6.1.2 Manual instantiation of generic procedures

It is actually possible to make the Fortran code version that was given in Listing 6.1, resemble the Go code version of Listing 2.2 a bit closer, by having two procedure pointers stand in, within the `select case` statement of the main program, for the closures that were used in the Go code. As this is a good example for demonstrating how generics can be instantiated manually by the programmer, we give in Listing 6.2 an alternative form of the main program of Listing 6.1 that makes use of both procedure pointers and such manual instantiation (as it was discussed in Sect. 5.3.1).

Listing 6.2: Main program using procedure pointers and manual generics instantiation.

```
1 program main
2
3   ! dependencies on intrinsic constants
4   use, intrinsic :: iso_fortran_env, only: real64
5
6   ! dependencies on implementations
7   use averager_library, only: simple_average, pairwise_average
8
9   implicit none
10
11  ! declarations
12  integer, parameter :: xi(5) = [1, 2, 3, 4, 5]
13  real(real64), parameter :: xf(5) = [1.d0, 2.d0, 3.d0, 4.d0, 5.d0]
14
15  integer :: key
16  procedure(average_integer), pointer :: avi
17  procedure(average_real64), pointer :: avf
18
19  abstract interface
20    function average_integer(x) result(a)
21      integer, intent(in) :: x(:)
22      integer :: a
23    end function average_integer
24    function average_real64(x) result(a)
25      real(real64), intent(in) :: x(:)
26      real(real64) :: a
27    end function average_real64
28  end interface
29
30  write(*, '(a)') 'Simple_sum_average:_1'
31  write(*, '(a)') 'Pairwise_sum_average:_2'
32  write(*, '(a)', advance='no') 'Choose_an_averaging_method:_ '
33  read(*, *) key
34
35  select case (key)
36  case (1)
37    avi => simple_average{integer}
38    avf => simple_average{real(real64)}
```

```

39  case (2)
40      avi => pairwise_average{integer}
41      avf => pairwise_average{real(real64)}
42  case default
43      stop 'Case_not_implemented!'
44  end select
45
46  print '(i8)',  avi(xi)
47  print '(f8.5)', avf(xf)
48
49  end program main

```

6.2 Object-oriented versions of the standard test problem

The present section will demonstrate that being able to use the new generics features seamlessly and easily even within a modern-day OO programming setting is one of the great strengths of the present design.

6.2.1 Dynamic method dispatch

Listing 6.3 shows our encapsulated (OO) Fortran code version of the standard test problem, that corresponds closest to the code versions that were presented in Chapter 3 for all the other languages.

- As in all these other versions, three interfaces are used to manage all the source code dependencies in the problem: `INumeric`, `ISum`, and `IAverager`. Interface `INumeric` is defined by the user himself as a type set, similar to the corresponding Go code.
- In contrast to the Go and Rust versions (Listings 3.1 and 3.2), none of the aforementioned interfaces is parameterized itself, since we followed Swift’s basic model of generics.
- Interface inheritance is expressed through the presence of the `implements` attribute in derived-type definitions (equivalent to the Swift version, Listing 3.4). Alternatively, `implements` statements could be used (cf. Sects. 4.2, 5.3.2, and 5.3.3).
- All our derived types are sealed, which makes them inextensible to implementation inheritance, and thus enables us to declare the passed-object dummy arguments of their type-bound procedures with the type specifier.
- The example code makes use, in the main program, of the new structure constructors, with their enhancements that were discussed in Sect. 5.4, for the classes `Averager`, `SimpleSum`, and `PairwiseSum`.
- This Fortran version makes use of modules and use statements with only clauses, in order to make explicit the source code dependencies of the different defined classes/ADTs.

Listing 6.3: Proposed encapsulated Fortran version of the array averaging example.

```

1  module interfaces

```

```

2
3 use, intrinsic :: iso_fortran_env, only: real64
4
5 implicit none
6 private
7
8 public :: INumeric, ISum, IAverager
9
10 abstract interface :: INumeric
11     integer | real(real64)
12 end interface INumeric
13
14 abstract interface :: ISum
15     function sum{INumeric :: T}(self,x) result(s)
16         deferred(self), intent(in) :: self
17         type(T),          intent(in) :: x(:)
18         type(T)           :: s
19     end function sum
20 end interface ISum
21
22 abstract interface :: IAverager
23     function average{INumeric :: T}(self,x) result(a)
24         deferred(self), intent(in) :: self
25         type(T),          intent(in) :: x(:)
26         type(T)           :: a
27     end function average
28 end interface IAverager
29
30 end module interfaces
31
32
33 module simple_library
34
35     use interfaces, only: ISum, INumeric
36
37     implicit none
38     private
39
40     public :: SimpleSum
41
42     type, sealed, implements(ISum) :: SimpleSum
43     contains
44         procedure :: sum
45     end type SimpleSum
46
47 contains
48
49     function sum{INumeric :: T}(self,x) result(s)

```

```

50     type(SimpleSum), intent(in) :: self
51     type(T),           intent(in) :: x(:)
52     type(T)            :: s
53     integer :: i
54     s = T(0)
55     do i = 1, size(x)
56         s = s + x(i)
57     end do
58     end function sum
59
60 end module simple_library
61
62
63 module pairwise_library
64
65     use interfaces, only: ISum, INumeric
66
67     implicit none
68     private
69
70     public :: PairwiseSum
71
72     type, sealed, implements(ISum) :: PairwiseSum
73     private
74     class(ISum), allocatable :: other
75     contains
76     procedure :: sum
77     end type PairwiseSum
78
79 contains
80
81     function sum{INumeric :: T}(self,x) result(s)
82     type(PairwiseSum), intent(in) :: self
83     type(T),           intent(in) :: x(:)
84     type(T)            :: s
85     integer :: m
86     if (size(x) <= 2) then
87         s = self%other%sum(x)
88     else
89         m = size(x) / 2
90         s = self%sum(x(:m)) + self%sum(x(m+1:))
91     end if
92     end function sum
93
94 end module pairwise_library
95
96
97 module averager_library

```

```

98
99 use interfaces, only: IAverager, ISum, INumeric
100
101 implicit none
102 private
103
104 public :: Averager
105
106 type, sealed, implements(IAverager) :: Averager
107     private
108     class(ISum), allocatable :: drv
109 contains
110     procedure :: average
111 end type Averager
112
113 contains
114
115 function average{INumeric :: T}(self,x) result(a)
116     type(Averager), intent(in) :: self
117     type(T),          intent(in) :: x(:)
118     type(T)           :: a
119     a = self%drv%sum(x) / T(size(x))
120 end function average
121
122 end module averager_library
123
124
125 program main
126
127 ! dependencies on abstractions
128 use interfaces,      only: IAverager
129
130 ! dependencies on implementations
131 use simple_library,  only: SimpleSum
132 use pairwise_library, only: PairwiseSum
133 use averager_library, only: Averager
134
135 implicit none
136
137 ! declarations
138 integer :: key
139 class(IAverager), allocatable :: avs, avp, av
140
141 ! use of enhanced structure constructors
142 avs = Averager(drv = SimpleSum())
143 avp = Averager(drv = PairwiseSum(other = SimpleSum()))
144
145 write(*,'(a)') 'Simple_{}_sum_average:_1'

```

```

146 write(*,'(a)') 'Pairwise_sum_average:_2'
147 write(*,'(a)',advance='no') 'Choose_an_averaging_method:_ '
148 read(*,*) key
149
150 select case (key)
151 case (1)
152     ! simple sum case
153     av = avs
154 case (2)
155     ! pairwise sum case
156     av = avp
157 case default
158     stop 'Case_not_implemented!'
159 end select
160
161 print '(i8)', av%average([1, 2, 3, 4, 5])
162 print '(f8.5)', av%average([1.d0, 2.d0, 3.d0, 4.d0, 5.d0])
163
164 end program main

```

The most important point to note, in Listing 6.3, is how this code makes use of trait objects (cf. Sect. 4.3) and generics to avoid rigid dependencies on user-defined implementation and language-intrinsic types, respectively, and to thus realize a plugin architecture that allows for a maximum of code reuse. Notice, in particular, how the main program is the only part of the code that (necessarily) depends on implementations. The *entire* rest of the code is decoupled, i.e. it depends merely on abstract interfaces (see the use statements in the above modules).

Notice also, that all the abstract interfaces that are required for this purpose are defined by the user himself, without any need to depend on external libraries. The OO Fortran version that is presented here is therefore as clean as the Go implementation of Listing 3.1 with respect to dependency management, and as easy to use, and to read and understand, as the Swift implementation of Listing 3.4.

6.2.2 Static method dispatch

One of the greatest benefits of the present design is that, through the use of generics, polymorphic methods in OO programming can be made to use static dispatch¹. This will enable inlining of polymorphic methods by the compiler, to potentially improve code performance. Listing 6.4 gives a *minimally* changed version of the OO Fortran code of Listing 6.3, to effect static dispatch of the different sum methods.

The required changes are confined to a parameterization of the PairwiseSum and Averager derived types, by generic type parameters that conform to the ISum interface and are named U. These type parameters are then used in order to declare the field objects other and drv of the PairwiseSum and Averager types, respectively, by means of the type specifier. Taken together, these changes signify compile-time polymorphism for the objects other and drv to the compiler, and hence static dispatch of their methods (whose interfaces were declared in the ISum interface). Contrast this with

¹This can be achieved in a compiler by implementing generic polymorphism through the (compile-time) technique of monomorphization, that relies on static method dispatch. In contrast to traditional (run-time) polymorphism, that relies on virtual method tables and dynamic method dispatch.

the class specifier, that was used previously for the declaration of these former trait objects, in order to effect run-time polymorphism and dynamic dispatch of their methods. See also Table 6.1 for a summary of rules regarding method dispatch.

Everything else, especially the declaration of these object variables as allocatables and their instantiation using chaining of constructor calls, was kept the same in order to demonstrate that static method dispatch does *not* mean that the actual object instances that contain the methods must be initialized and their memory allocated at compile-time (although in this particular case this is possible, as demonstrated in the next section, given that these objects do not contain any other allocatable data fields, like arrays). Notice, also, that none of the source code dependencies in the use statements have changed. That is, the code is *still* fully decoupled, despite making now use of static dispatch!

Listing 6.4: Demonstrates static method dispatch for the sum methods.

```

1  module interfaces
2
3      use, intrinsic :: iso_fortran_env, only: real64
4
5      implicit none
6      private
7
8      public :: INumeric, ISum, IAverager
9
10     abstract interface :: INumeric
11         integer | real(real64)
12     end interface INumeric
13
14     abstract interface :: ISum
15         function sum{INumeric :: T}(self,x) result(s)
16             deferred(self), intent(in) :: self
17             type(T),          intent(in) :: x(:)
18             type(T)           :: s
19         end function sum
20     end interface ISum
21
22     abstract interface :: IAverager
23         function average{INumeric :: T}(self,x) result(a)
24             deferred(self), intent(in) :: self
25             type(T),          intent(in) :: x(:)
26             type(T)           :: a
27         end function average
28     end interface IAverager
29
30 end module interfaces
31
32
33 module simple_library
34
35     use interfaces, only: ISum, INumeric

```



```

36
37  implicit none
38  private
39
40  public :: SimpleSum
41
42  type, sealed, implements(ISum) :: SimpleSum
43  contains
44      procedure :: sum
45  end type SimpleSum
46
47  contains
48
49  function sum{INumeric :: T}(self,x) result(s)
50      type(SimpleSum), intent(in) :: self
51      type(T),          intent(in) :: x(:)
52      type(T)           :: s
53      integer :: i
54      s = T(0)
55      do i = 1, size(x)
56          s = s + x(i)
57      end do
58  end function sum
59
60  end module simple_library
61
62
63  module pairwise_library
64
65      use interfaces, only: ISum, INumeric
66
67      implicit none
68      private
69
70      public :: PairwiseSum
71
72      type, sealed, implements(ISum) :: PairwiseSum{ISum :: U}
73      private
74          type(U), allocatable :: other
75      contains
76          procedure :: sum
77      end type PairwiseSum
78
79  contains
80
81  function sum{INumeric :: T}(self,x) result(s)
82      type(PairwiseSum{U}), intent(in) :: self
83      type(T),              intent(in) :: x(:)

```

```

84     type(T)                                :: s
85     integer :: m
86     if (size(x) <= 2) then
87         s = self%other%sum(x)
88     else
89         m = size(x) / 2
90         s = self%sum(x(:m)) + self%sum(x(m+1:))
91     end if
92     end function sum
93
94 end module pairwise_library
95
96
97 module averager_library
98
99     use interfaces, only: IAverager, ISum, INumeric
100
101     implicit none
102     private
103
104     public :: Averager
105
106     type, sealed, implements(IAverager) :: Averager{ISum :: U}
107     private
108     type(U), allocatable :: drv
109     contains
110     procedure :: average
111     end type Averager
112
113 contains
114
115     function average{INumeric :: T}(self,x) result(a)
116         type(Averager{U}), intent(in) :: self
117         type(T),           intent(in) :: x(:)
118         type(T)            :: a
119         a = self%drv%sum(x) / T(size(x))
120     end function average
121
122 end module averager_library
123
124
125 program main
126
127     ! dependencies on abstractions
128     use interfaces,      only: IAverager
129
130     ! dependencies on implementations
131     use simple_library,  only: SimpleSum

```

```

132 use pairwise_library, only: PairwiseSum
133 use averager_library, only: Averager
134
135 implicit none
136
137 ! declarations
138 integer :: key
139 class(IAverager), allocatable :: avs, avp, av
140
141 ! use of enhanced structure constructors
142 avs = Averager(drv = SimpleSum())
143 avp = Averager(drv = PairwiseSum(other = SimpleSum()))
144
145 write(*,'(a)') 'Simple_sum_average:_1'
146 write(*,'(a)') 'Pairwise_sum_average:_2'
147 write(*,'(a)',advance='no') 'Choose_an_averaging_method:_ '
148 read(*,*) key
149
150 select case (key)
151 case (1)
152     ! simple sum case
153     av = avs
154 case (2)
155     ! pairwise sum case
156     av = avp
157 case default
158     stop 'Case_not_implemented!'
159 end select
160
161 print '(i8)', av%average([1, 2, 3, 4, 5])
162 print '(f8.5)', av%average([1.d0, 2.d0, 3.d0, 4.d0, 5.d0])
163
164 end program main

```

object declaration	dynamic dispatch	static dispatch
class(Interface)	always	never
class(DerivedType)	if DerivedType is extended	if DerivedType is unextended ²
type(DerivedType)	never	always
type(Interface)	—	—

Table 6.1: Correspondence between object declarations and method dispatch strategies that would be typically employed by an optimizing Fortran compiler according to the present document. A dash indicates that the case in question is presently undefined, but could be used for a future extension as discussed in Sect. 5.1.2.

²Or the method is declared as non_overridable.

The `av` object of `IAverager` type, in Listing 6.4, still needs to make use of run-time polymorphism, because it is initialized in a `select` case statement by the main program. This object cannot be made to employ compile-time polymorphism, as it is initialized within a statement that performs a run-time decision.

As in corresponding Swift and Rust implementations of the standard test problem, that are not fully reproduced here (see the accompanying Code directory), the instantiation (in the main program) of the objects `other` and `drv` through constructor calls, has the benefit that the compiler can infer the correct type arguments, that are required to automatically instantiate all the involved generic derived types. Hence (and similarly to both the functional and OO code versions of Listings 6.1 and 6.3, respectively), not a single manual instantiation of generics is necessary, anywhere, in Listing 6.4.

As a final remark, we'd like to emphasize that, on readability grounds, a coding style as that given in Listing 6.3 is generally preferable over that of Listing 6.4. The use of numerous generic type parameters can quickly make code unreadable. We'd therefore recommend the default use of run-time polymorphism for managing dependencies on user-defined types, and the employment of generics for this latter task merely in cases where profiling has shown that static dispatch would significantly speed up a program's execution (by allowing method inlining by the compiler). Of course, the use of generics to manage dependencies on language-intrinsic types remains unaffected by this recommendation.

6.2.3 Static method dispatch and static object declarations

In the present simple example, and taking Listing 6.4 as a baseline, it is actually possible to even avoid some of the run-time memory allocation overhead of the program, by having the field objects `other` and `drv`, that are contained within the `PairwiseSum` and `Averager` types, be statically declared. To accomplish this, the two lines

```
1  type(U), allocatable :: other
2  type(U), allocatable :: drv
```

in Listing 6.4 need to be replaced by the following two code lines:

```
1  type(U) :: other
2  type(U) :: drv
```

Notice, that even in this code version, and because the compiler should be able to automatically infer generic type arguments from the types of regular arguments in constructor calls, the object instantiations that are to be carried out from the main program can remain the same, that is:

```
1  avs = Averager(drv = SimpleSum())
2  avp = Averager(drv = PairwiseSum(other = SimpleSum()))
```

Thus making manual instantiation of generic types unnecessary.

Alternatively, though, the two calls of `Averager`'s structure constructor could be manually given generic type arguments to confirm the types of the regular arguments (deleting here the latter's keywords):

```
1  avs = Averager{SimpleSum}(SimpleSum())
2  avp = Averager{PairwiseSum{SimpleSum}}(PairwiseSum(SimpleSum()))
```

Or one could delete the regular arguments to the constructor altogether, and provide only the generic type arguments, as follows (see Sect. 5.4):

```

1  avs = Averager{SimpleSum}()
2  avp = Averager{PairwiseSum{SimpleSum}}()

```

The full source code for this version of the test problem can be found in the Code subdirectory that is accompanying this document.

6.3 Employing general generics constraints

In the previous examples of this chapter, we made predominant use of generics constraints that, in the form of the `INumeric` interface, were formulated as type set unions of intrinsic types. While this is a very frequently occurring use case, it is of course possible to formulate generics constraints in a more general manner, through explicit specification of procedure signatures in abstract interfaces. One such example was already given in the form of the `ISum` interface, that was used as a generics constraint for derived types in Listing 6.4.

Another example is shown in the following Listing 6.5. This listing provides a main program with an internal `generic_sum` function, whose generic type parameter, `T`, is constrained by the interface `IAdmissible` of Sect. 4.1.4. This interface requires not only numeric functionality in the form of a custom addition operator, but also custom type-conversion functionality. Listing 6.5 provides also two modules that implement both this `IAdmissible` interface, and the `IPrintable` interface of Sect. 4.1.1, for two different data types. A derived type with name `MyType` illustrates the procedure of implementing these interfaces for a user-defined type, while the same is also shown for the real type, as an intrinsic type example. This enables the use of *both* these types with the `generic_sum` function, but also with a second internal subroutine, called `printout`, that requires I/O functionality.

Listing 6.5: Example of using non-type-set interfaces as generics constraints for procedures.

```

1  module my_type
2
3      use interfaces, only: IAdmissible, IPrintable
4
5      implicit none
6      private
7
8      type, public, sealed :: MyType
9          private
10         integer :: n
11     end type MyType
12
13     implements (IAdmissible,IPrintable) :: MyType
14         procedure :: add, cast, output
15     end implements MyType
16
17 contains
18
19     function add(self,other) result(res)
20         type(MyType), intent(in) :: self, other
21         type(MyType)                :: res
22         res%n = self%n + other%n

```

```

23 end function add
24
25 function cast(self,i) result(res)
26     type(MyType), intent(in) :: self
27     integer,        intent(in) :: i
28     type(MyType)          :: res
29     res%n = i
30 end function cast
31
32 subroutine output(self)
33     type(MyType), intent(in) :: self
34     write(*,'(a,i8)') "I_am:_", self%n
35 end subroutine output
36
37 end module my_type
38
39 module real_type
40
41     use interfaces, only: IAdmissible, IPrintable
42
43     implicit none
44     private
45
46     implements (IAdmissible,IPrintable) :: real
47         procedure :: add, cast, output
48     end implements real
49
50 contains
51
52     function add(self,other) result(res)
53         real, intent(in) :: self, other
54         real              :: res
55         res = self + other
56     end function add
57
58     function cast(self,i) result(res)
59         real,    intent(in) :: self
60         integer, intent(in) :: i
61         real      :: res
62         res = real(i)
63     end function cast
64
65     subroutine output(self)
66         real, intent(in) :: self
67         write(*,'(a,f8.5)') "I_am:_", self
68     end subroutine output
69
70 end module real_type

```

```

71
72 program test
73
74 use interfaces, only: IAdmissible, IPrintable
75 use my_type, only: MyType
76
77 implicit none
78
79 integer :: i
80 real :: arr_r(10) ! array of real type
81 real :: res_r ! result of real type
82 type(MyType) :: arr_t(10) ! array of MyType
83 type(MyType) :: res_t ! result of MyType
84
85 ! initializations
86 arr_r = [(real(i),i=1,10)]
87 arr_t = [(MyType(i),i=1,10)] ! use of enhanced structure constructor
88
89 ! use generic_sum with arrays of both types; no manual instantiations needed!
90 res_r = generic_sum(arr_r)
91 res_t = generic_sum(arr_t)
92
93 ! print the results
94 call printout(res_r,res_t)
95
96 contains
97
98 pure function generic_sum{IAdmissible :: T}(arr) result(res)
99 type(T), intent(in) :: arr(:)
100 type(T) :: res
101 integer :: n, i
102 n = size(arr)
103 res = res%cast(0)
104 if (n > 0) then
105 res = arr(1)
106 do i = 2, n
107 res = res + arr(i)
108 end do
109 end if
110 end function generic_sum
111
112 subroutine printout{IPrintable :: T,U}(res1,res2)
113 type(T), intent(in) :: res1
114 type(U), intent(in) :: res2
115 call res1%output()
116 call res2%output()
117 end subroutine printout
118

```

6.4 Associated type usage with generics

Lastly, we will demonstrate, now, how associated types in abstract interfaces (as they were introduced in Sect. 4.1.2) can be used in conjunction with implementing types that are parameterized over some generic parameter. The following example defines an abstract interface, `IAppendable`, that requires any implementing type to provide functionality for appending some item to itself. This interface contains two associated types: an alias, `deferred(self)`, for the type of the passed-object dummy argument of an implementing derived type, and a placeholder, `deferred(item)`, for the type of the appendable items, that needs to be supplied by an actual implementation of that interface.

```

1 module vector_library
2
3   implicit none
4   private
5
6   abstract interface :: IAnyType
7   end interface IAnyType
8
9   abstract interface :: IAppendable
10    subroutine append(self,item)
11      deferred(self), intent(inout) :: self
12      deferred(item), intent(in)    :: item
13    end subroutine append
14  end interface IAppendable
15
16  type, public, implements(IAppendable) :: Vector{IAnyType :: U}
17    private
18    type(U), allocatable :: elements(:)
19    contains
20      procedure :: append
21    end type Vector
22
23  contains
24
25    subroutine append(self,item)
26      class(Vector{U}), intent(inout) :: self
27      type(U),          intent(in)    :: item
28      self%elements = [self%elements,item]
29    end subroutine append
30
31 end module vector_library

```

An example of such an implementation of interface `IAppendable` is provided here by the derived type `Vector`, that stores an `elements` array of generic type `U`, where the type parameter `U` conforms to the `IAnyType` constraint. Notice, that in order to maintain consistency between the

type of the `elements` array, and any new `item` that we wish to append to it, we must force the `item` argument of method `append` to be of type `U` as well, as it is shown in this method's actual implementation. In order to accomplish this enforcement without contradicting interface `IAppendable`'s definition (that doesn't know anything about type `U`), we made use of the placeholder (i.e. associated) type `deferred(item)` in the latter interface. Given method `append`'s implementation, the compiler will then infer `deferred(item)` to be of the actual type `U`.

Associated types thus allow one to make interfaces generic, without having to parameterize either them or their method signatures by actual generic type parameters, and to thereby entail the instantiation needs of the latter, and the potential duplication of client code that this can result in (cf. Sect. 3.1.1). Associated types are therefore an important element of the present generics design, and its philosophy not to burden the user with superfluous manual instantiations of generics.

Chapter 7

Conclusions

The Fortran extensions, that are described in this document for both run-time and compile-time polymorphism, resulted from the consistent application of orthogonal language design. That is, significant new capabilities are provided through the simple extension of already existing Fortran features, and their mutual interaction, rather than the addition of (inappropriate and superfluous) new constructs. Indeed, only in one single case (the implements statement of Sect. 4.2) did we find it necessary to introduce a new language feature. But even then, it was ensured that the new feature would be of *equal* utility in supporting both compile-time and run-time polymorphism, so that orthogonality, again, prevailed.

The end result is a language that is both much easier to use, and much more powerful, than that of competing approaches which prefer to abandon the proven philosophy of orthogonal design. The presented extensions are fully backwards compatible with the present Fortran language, and (in stark contrast to competing approaches) allow for a consistent use of fully type-checked generics not just in procedural and functional, but *also* in OO programming settings, *including* the support of static method dispatch (which can improve the performance of polymorphic methods by facilitating inlining via the compiler).

Taken together, these capabilities will allow for the uniform management of source code dependencies on both language-intrinsic and user-defined types in Fortran. They will enable the Fortran programmer to write unprecedented modular code, that is on par with the most modern languages in terms of reusability, and moreover does not sacrifice any computational performance. The present design achieves all of this largely *without* requiring manual instantiations of generics, and it furthermore provides a solid foundation for a number of possible future extensions, like support for array-rank genericity, and compile time polymorphism via union (sum) types.

Bibliography

- [1] Luca Cardelli and Peter Wegner. “On understanding types, data abstraction, and polymorphism”. In: *ACM Computing Surveys* 17 (4 Dec. 1985), pages 471–523. URL: <https://dl.acm.org/doi/pdf/10.1145/6041.6042>.
- [2] Brad J. Cox, Steve Naroff, and Hansen Hsu. “The Origins of Objective-C at PPI/Stepstone and Its Evolution at NeXT”. In: *Proceedings of the ACM on Programming Languages* 4 (June 2020), pages 1–74. URL: <https://dl.acm.org/doi/10.1145/3386332>.
- [3] Ole-Johan Dahl. “The Birth of Object Orientation: the Simula Languages”. In: *From Object-Orientation to Formal Methods*. Edited by O. Owe, S. Krogdahl, and T. Lyche. Berlin: Springer, 2004. URL: <https://www.mn.uio.no/ifi/english/about/ole-johan-dahl/bibliography/the-birth-of-object-orientation-the-simula-languages.pdf>.
- [4] Erich Gamma et al. *Design Patterns. Elements of Reusable Object-Oriented Software*. Prentice Hall, 1994.
- [5] Magne Haverlaen, Jaakko Järvi, and Damian Rouson. *Reflecting on Generics for Fortran*. July 2019. URL: <https://j3-fortran.org/doc/year/19/19-188.pdf>.
- [6] Allen Holub. *Holub on Patterns, Learning Design Patterns by Looking at Code*. Apress, 2004. ISBN: 1-59059-388-X.
- [7] Steve Klabnik and Carol Nichols. *The Rust Programming Language*. URL: <https://doc.rust-lang.org/stable/book/ch10-02-traits.html#implementing-a-trait-on-a-type>.
- [8] George Lyon. URL: <https://gist.github.com/GeorgeLyon/c7b07923f7a800674bc9745ae45ddc7f>.
- [9] Robert C. Martin. *Agile Software Development: Principles, Patterns, and Practices*. Prentice Hall, 2002. ISBN: 978-0-135-97444-5.
- [10] Robert C. Martin. *OO vs FP*. Nov. 2014. URL: <https://blog.cleancoder.com/uncle-bob/2014/11/24/FPvsOO.html>.
- [11] Robert C. Martin. *Clean Architecture: A Craftsman’s Guide to Software Structure and Design*. Boston, MA: Prentice Hall, 2017. ISBN: 978-0-13-449416-6.
- [12] Nicholas D Matsakis and Felix S Klock II. “The Rust language”. In: *ACM SIGAda Ada Letters*. Volume 34. ACM. 2014, pages 103–104.
- [13] Benjamin C. Pierce. *Programming with intersection types, union types, and polymorphism*. URL: <https://api.semanticscholar.org/CorpusID:118907019>.
- [14] Ian Lance Taylor. URL: <https://github.com/golang/go/issues/45346>.
- [15] *The Java Tutorials. Abstract Methods and Classes*. URL: <https://docs.oracle.com/javase/tutorial/java/IandI/abstract.html>.

- [16] *The Swift Programming Language (6 beta). Extensions*. URL: <https://docs.swift.org/swift-book/documentation/the-swift-programming-language/extensions/>.
- [17] Wolfgang Weck and Clemens Szyperski. *Do We Need Inheritance?* URL: https://www.researchgate.net/publication/2297653_Do_We_Need_Inheritance.

Acknowledgements

The present extensions for Fortran could not have been devised without the guidance we had through the prior work of countless developers on the Swift, Go, Rust, Carbon, and Java languages. We wish to thank all these developers, even though we can only mention here a few of them explicitly by name, who have particularly influenced the present effort. Chris Lattner's design of the Swift language provided the foundational ideas upon which we built the majority of our language extensions for Fortran, including much of the modern OO features, and the generics. While Robert Griesemer's, Ian Lance Taylor's, and the entire Go team's inspirational work on type sets accounts largely not only for our remaining generics design, but also for its ease of use. Robert also kindly provided the original code version from which Listing 2.2 was derived, and answered the many questions we had regarding Go's generics. We are also grateful to everyone who provided feedback to our work, like Magne Haverlaen and Brad Richardson, whose comments helped us to improve this paper.