

# **Traits, Generics, and Modern-day OOP for Fortran**

Konstantinos Kifonidis, Ondrej Certik, Derick Carnazzola

June 22, 2025

## Abstract

Based on conclusions drawn from a survey of modern languages, a traits system for Fortran is developed that allows for the uniform expression of run-time and compile-time polymorphism in the language, and thus for the uniform management of source code dependencies on user-defined and language-intrinsic types. The feature set that is described here is small enough to facilitate a first prototype implementation in an open source compiler (like LFortran, LLVM Flang, or GNU Fortran), but at the same time comprehensive enough to already endow Fortran with polymorphic capabilities that widely equal those of modern programming languages like Swift, Rust, Go, Carbon, or Mojo. The discussed extensions are fully backwards compatible with present Fortran, and enable modern-day, traits-based, object-oriented programming, and powerful, easy to use, fully type-checked, generics. The latter support seamlessly both the procedural, functional, and object-oriented programming styles, and they largely get by *without* requiring manual instantiations by the user. The presented design could also be naturally extended to support rank-genericity for arrays, structural subtyping, and compile-time polymorphic union/sum types. Its modern capabilities are expected to transform the way both Fortran applications and libraries will be written in the future. Decoupled software plugin architectures with enormously improved source code flexibility, reusability, maintainability, and reliability will become possible, without any need for run-time type inspections, and without any loss in computational performance.

# Chapter 1

## Introduction

Polymorphism was discovered in the 1960ies by Kristen Nygaard and Ole-Johan Dahl during their development of Simula 67, the world's first object-oriented (OO) language [3]. Their work introduced into programming what is nowadays known as “virtual method table” (i.e. function-pointer) based run-time polymorphism, which is both the first focus of this document, and the decisive feature of all OO languages. Several other forms of polymorphism are known today, the most important of them being parametric polymorphism [1], also known as “generics”, which is the second focus of this document, and which has historically developed disjointly from run-time polymorphism since it makes use of compile-time mechanisms.

### 1.1 The purpose of polymorphism

But what is the purpose of polymorphism in a programming language? What is polymorphism actually good for? One of the more comprehensive answers to this question was given by Robert C. Martin in numerous books (e.g. [13]), as well as in the following quotation from his blog [12]:

“There really is only one benefit to polymorphism; but it's a big one. It is the inversion of source code and run time dependencies.

In most software systems when one function calls another, the runtime dependency and the source code dependency point in the same direction. The calling module depends on the called module. However, when polymorphism is injected between the two there is an inversion of the source code dependency. The calling module still depends on the called module at run time. However, the source code of the calling module does not depend upon the source code of the called module. Rather both modules depend upon a polymorphic interface.

This inversion allows the called module to act like a plugin. Indeed, this is how all plugins work.”

Notice the absence of the words “code reuse” in these statements. The purpose of polymorphism, according to Martin, is the “inversion” (i.e. replacement, or management) of rigid source code dependencies by means of particular abstractions, i.e. polymorphic interfaces (or protocols/traits, as they are also known today). The possibility to reuse code is then merely the logical consequence of such proper dependency management in a polymorphism-based software plugin architecture.

## 1.2 Source code dependencies in statically typed languages

Which then are the source code dependencies that polymorphism helps us manage? It has been customary to make the following distinction when answering this question:

- Firstly, most larger programs that are written in statically typed languages (like Fortran) have dependencies on user-defined procedures and data types. If the programmer employs encapsulation of both a program's procedures and its data, i.e. its state, both these dependencies can actually be viewed as dependencies on *user-defined abstract data types*. These are the dependencies that Martin is concretely referring to in the above quotation, and it is these dependencies on (volatile) user-defined data and implementations that are particularly troublesome, because they lead to rigid coupling between the various different *parts* of an application. Their results are recompilation cascades, the non-reusability of higher-level source code units, the impossibility to comprehend a large application incrementally, and fragility of such an application as a whole.
- Secondly, every program, that is written in a statically typed language, also has dependencies on abstract data types that are provided by the language itself. Fortran's integer, real, complex, etc., intrinsic types are examples of *language-intrinsic abstract data types*. While hard-wired dependencies on such intrinsic types do not couple different parts of a program (because the implementations of these types are supplied by the language), they nevertheless make a program's source code rigid with respect to the data that it can be used on.

The most widely used approaches to manage dependencies on language-intrinsic types have so far been through generics, while dependency management of user-defined (abstract data) types has so far been the task of OO programming and OO design patterns. Martin [13] has, for instance, defined object-orientation as follows:

"OO is the ability, through the use of polymorphism, to gain absolute control over every source code dependency in [a software] system. It allows the architect to create a plugin architecture, in which modules that contain high-level policies are independent of modules that contain low-level details. The low-level details are relegated to plugin modules that can be deployed and developed independently from the modules that contain high-level policies."

## 1.3 Modern developments

Notice how Martin's modern definition of object-orientation, that emphasizes source code decoupling, is the antithesis to the usually taught "OO" approaches of one class rigidly inheriting implementation code from another. Notice also how his definition does not require some specific type of polymorphism for the task of dependency management, as long as (according to Martin's first quotation) the mechanism is based on polymorphic interfaces.

Martin's statements on the purpose of both polymorphism and OO simply reflect the two crucial developments that have taken place in these fields over the last decades. Namely, the realizations that

- run-time polymorphism should be freed from the conflicting concept of implementation inheritance (to which it was originally bound given its Simula 67 heritage), and be formulated exclusively in terms of conformance to polymorphic interfaces, i.e. function signatures, or purely procedural abstractions, and that

- compile-time polymorphism should be formulated in exactly the same way as well.

These two developments, taken together, have recently opened up the possibility to treat polymorphism, and hence the dependency management of both user-defined and language-intrinsic types, uniformly in a programming language. As a consequence, it has become possible to use the potentially more efficient (but also less flexible) mechanism of compile-time polymorphism also for certain tasks that have traditionally been reserved for run-time polymorphism (in the form of OO programming), and to mix and match the two polymorphism types inside a single application to better satisfy a user's needs for both flexibility and efficiency.

## 1.4 Historical background

The road towards these realizations was surprisingly long. Over the last five decades, a huge body of OO programming experience first had to demonstrate that the use of (both single and multiple) implementation inheritance breaks encapsulation in OO languages, and therefore results in extremely tightly coupled, rigid, fragile, and non-reusable code. This led to an entire specialized literature on OO design patterns [4, 7, 11], that aimed at avoiding such rigidity by replacing the use of implementation inheritance with the means to formulate run-time polymorphism that are discussed below. It also led to the apprehension that implementation inheritance (but *not* run-time polymorphism) should be abandoned [22]. In modern languages, implementation inheritance is either kept solely for backwards compatibility reasons (e.g. in the Swift and Carbon languages), or it is foregone altogether (e.g. in Rust, and Go).

The first statically typed mainstream programming language that offered a proper separation of run-time polymorphism from implementation inheritance was Objective-C. It introduced “protocols” (i.e. polymorphic interfaces) in the year 1990 [2]. Protocols in Objective-C consist of pure function signatures, that lack implementation code. Objective-C provided a mechanism to implement multiple such protocols by a class, and to thus make classes conform to protocols. This can be viewed as a restricted form of multiple inheritance, namely inheritance of object *specification*, which is also known as *subtyping*.

Only a few years later, in 1995, the Java language hugely popularized these ideas using the terms “interfaces” and “interface inheritance” [2]. Today, nearly all modern languages support polymorphic interfaces/protocols, and the basic mechanism of multiple interface inheritance that was introduced to express run-time polymorphism in Objective-C, often in even improved, more flexible, manifestations. The only negative exceptions in this respect being modern Fortran, and C++, which both still stick to the obsolescent Simula 67 paradigm.

A similarly lengthy learning process, as that outlined for run-time polymorphism, also took place in the field of compile-time/parametric polymorphism. Early attempts, notably templates in C++, to render function arguments and class parameters polymorphic, did not impose any constraints on such arguments and parameters, that could be checked by C++ compilers. With the known results on compilation times and cryptic compiler error messages [6]. Surprisingly, Java, the language that truly popularized polymorphic interfaces in OO programming, did not provide an interface based mechanism to constrain its generics. Within the pool of mainstream programming languages, this latter realization was only made with the advent of Rust [14].

Rust came with a traits (i.e. polymorphic interfaces) system with which it is possible for the user to uniformly and transparently express both generics (i.e. compile-time) and run-time polymorphism in the same application, and to relatively easily switch between the two, where possible. Rust's traits are an improved form of protocols/interfaces in that the user can implement them for

a type without having these implementations be coupled to the type's actual definition. Thus, existing types can be made to retroactively implement new traits, and hence be used in new settings (with some minor restrictions on user ownership of either the traits or the types).

Rust's main idea was quickly absorbed by almost all other mainstream modern languages, most notably Swift, Go, and Carbon, with the difference that these latter languages tend to leave the choice between static and dynamic binding of procedures to the compiler, or language implementation, rather than the programmer. C++ is in the process of adopting generics constraints for its "templates" under the term "strong concepts", but without implementing the greater idea to uniformly express *all* the polymorphism in the language through traits. An implementation of this latter idea must today be viewed as a prerequisite in order to call a language design "modern". The purpose of this document is to describe extensions to Fortran, that aim to provide the Fortran language with such modern capabilities.

## Chapter 2

# Case study: Calculating the average value of a numeric array

To illustrate the advanced features and capabilities of some of the available modern programming languages with respect to polymorphism, and hence dependency management, we will make use here of a case study: the simple test problem of calculating the average value of a set of numbers that is stored inside a one-dimensional array. In the remainder of this chapter, we will first provide an account and some straightforward monomorphic (i.e. rigidly coupled) functional implementation of this test problem, followed by a functional implementation that makes use of both run-time and compile-time polymorphism to manage rigid source code dependencies. In the survey of programming languages that is presented in Chapter 3, we will then recode this standard test problem in an encapsulated fashion, to highlight how the source code dependencies in this problem can be managed in different languages even in more complex situations, that require OO techniques.

### 2.1 Monomorphic functional implementation

We have chosen Go here as a language to illustrate the basic ideas. Go is easily understood, even by beginners, and is therefore well suited for this purpose (another good choice would have been the Swift language). The code in the following Listing 2.1 should be self explanatory for anyone who is even only remotely familiar with the syntax of C family languages. So, we'll make only a few remarks regarding syntax.

- While mostly following a C like syntax, variable declarations in Go are essentially imitating Pascal syntax, where a variable's name precedes the declaration of the type.
- Go has two assignment operators. The usual `=` operator, as it is known from other languages, and the separate `:=` operator, that employs automatic type inference for combined declaration and initialization of a variable.
- Go has array slices that most closely resemble those of Python's Numpy (which exclude the upper bound of an array slice).

Our basic algorithm for calculating the average value of an array of integer elements employs two different implementations for averaging. The first makes use of a "simple" summation of all the array's elements, in ascending order of their array index. While the second sums in a "pair-wise" manner, dividing the array in half to carry out the summations recursively, and switching

to the “simple” method once subdivision is no longer possible. In both cases, the resulting sum is then divided by the array’s number of elements, to obtain the desired average.

Listing 2.1: Monomorphic functional version of the array averaging example in Go.

```
1 package main
2
3 import "fmt"
4
5 func simple_sum(x []int32) int32 {
6     s := int32(0)
7     for i := 0; i < len(x); i++ {
8         s += x[i]
9     }
10    return s
11 }
12
13 func pairwise_sum(x []int32) int32 {
14     if len(x) <= 2 {
15         return simple_sum(x)
16     } else {
17         m := len(x) / 2
18         return pairwise_sum(x[:m+1]) + pairwise_sum(x[m+1:])
19     }
20 }
21
22 func simple_average(x []int32) int32 {
23     return simple_sum(x) / int32(len(x))
24 }
25
26 func pairwise_average(x []int32) int32 {
27     return pairwise_sum(x) / int32(len(x))
28 }
29
30 // .....
31 // main program
32 // .....
33
34 func main() {
35     var key int32
36
37     fmt.Println("Simple_sum_average:_1")
38     fmt.Println("Pairwise_sum_average:_2")
39     fmt.Print("Choose_an_averaging_method:_")
40     fmt.Scan(&key)
41
42     xi := []int32{1,2,3,4,5}
43
44     switch key {
45     case 1:
46         fmt.Printf("%d\n", simple_average(xi))
47     case 2:
48         fmt.Printf("%d\n", pairwise_average(xi))
49     default:
50         fmt.Println("Case_not_implemented!")
51     }
52
53 }
```

An inspection of Listing 2.1 will readily reveal that this code has three levels of rigid (i.e. hard-wired) dependencies. Namely,

1. function `pairwise_sum` depending on function `simple_sum`'s implementation,
2. functions `simple_average` and `pairwise_average` depending on functions' `simple_sum` and `pairwise_sum` implementations, respectively, and
3. the entire program depending rigidly on the `int32` data type in order to declare both the arrays that it is operating on, and the results of its summation and averaging operations.

The first two items are dependencies on user-defined implementations, while the third is a typical case of rigid dependency on a language-intrinsic type, which renders the present code incapable of being applied to arrays of any other data type than `int32s`. Given that we are dealing with three levels of dependencies, three levels of polymorphism will accordingly be required to remove all these dependencies.

## 2.2 Polymorphic functional implementation

Listing 2.2 gives an implementation of our test problem, that employs Go's generics and functional features to eliminate the last two of the rigid dependencies that were listed in Sect. 2.1. The code makes use of Go's generics to admit arrays of both the `int32` and `float64` types as arguments to all functions, and to express the return values of the latter. It also makes use of the run-time polymorphism inherent in Go's advanced functional features, namely closures and variables of higher-order functions, to replace the two previous versions of function `average` (that depended on specific implementations), by a single polymorphic version. Only the rigid dependency of function `pairwise_sum` on function `simple_sum` has not been removed, to keep the code more readable. In the OO code versions, that will be presented in Chapter 3, even this dependency is eliminated.

A few remarks are in order for a better understanding of Listing 2.2's code:

- In Go, generic type parameters to a function, like the parameter `T` here, are provided in a separate parameter list, that is enclosed in brackets `[ ]`.
- Generic type parameters have a constraint that follows their declared name. Go exclusively uses interfaces as such constraints (like the interface `INumeric` in the following code).
- Interfaces consist of either explicit function signatures, or *type sets*, like `int32 | float64` in the present example. The latter actually signify a set of function signatures, too, namely the signatures of the intersecting set of all the operations and intrinsic functions for which the listed types provide implementations.
- The code makes use of type conversions to the generic type `T`, where required. For instance, `T(0)` converts the (typeless) constant `0` to the corresponding zero constant of type `T`.
- The code instantiates closures and stores these by value in two variables named `avi` and `avf` for later use (Fortran and C programmers should note that `avi` and `avf` are *not* function pointers!).

Listing 2.2: Polymorphic functional version of the array averaging example in Go.

```

1 package main
2
3 import "fmt"
4
5 type INumeric interface{ int32 | float64 }
6
7 func simple_sum[T INumeric](x []T) T {
8     s := T(0)
9     for i := 0; i < len(x); i++ {
10         s += x[i]
11     }
12     return s
13 }
14
15 func pairwise_sum[T INumeric](x []T) T {
16     if len(x) <= 2 {
17         return simple_sum(x)
18     }
19     m := len(x) / 2
20     return pairwise_sum(x[:m+1]) + pairwise_sum(x[m+1:])
21 }
22
23 func average[T INumeric](sum func([]T) T, x []T) T {
24     return sum(x) / T(len(x))
25 }
26
27 func main() {
28     var avi func([]int32) int32
29     var avf func([]float64) float64
30     var key int32
31
32     fmt.Println("Simple_sum_average:", 1)
33     fmt.Println("Pairwise_sum_average:", 2)
34     fmt.Print("Choose_an_averaging_method:_")
35     fmt.Scan(&key)
36
37     switch key {
38     case 1:
39         avi = func(x []int32) int32 {
40             return average(simple_sum[int32], x)
41         }
42         avf = func(x []float64) float64 {
43             return average(simple_sum[float64], x)
44         }
45     case 2:
46         avi = func(x []int32) int32 {
47             return average(pairwise_sum[int32], x)
48         }
49         avf = func(x []float64) float64 {
50             return average(pairwise_sum[float64], x)
51         }
52     default:
53         fmt.Println("Case_not_implemented!")
54         return
55     }
56
57     xi := []int32{1, 2, 3, 4, 5}

```

```
58     xf := []float64{1, 2, 3, 4, 5}
59
60     fmt.Printf("%d\n",  avi(xi))
61     fmt.Printf("%.1f\n",avf(xf))
62 }
```

Notice how, in order to instantiate the closures `avi` and `avf` (see the `switch` statement), manual instantiations of the `simple_sum` and `pairwise_sum` generic functions are required – with the arguments `int32` or `float64` being substituted for the generic type parameter, `T`, of these functions.

The motivation to code the example as in Listing 2.2 is that once the two closures, `avi` and `avf`, are properly instantiated, they could then be passed from the main program to any other client code that would need to make use of the particular averaging algorithm that was selected by the user. This latter client code would *not* have to be littered with `switch` statements itself, and it would *not* have to depend on any specific implementations. It would merely depend on the closures' interfaces. The same holds for the OO code versions that are discussed in the next chapter, with objects replacing the closures (both being merely slightly different realizations of the same idea).

## Chapter 3

# Survey of modern languages

In the present chapter, we give encapsulated (i.e. OO) code versions of the test problem in various modern languages. As in the functional code version that was presented in Sect. 2.2, we employ run-time polymorphism to manage the dependencies on user-defined implementations (in this case abstract data types), and generics in order to manage the dependencies on language-intrinsic types. This serves to illustrate how both run-time and compile-time polymorphism can be typically used for dependency management in an OO setting in these modern languages. The survey also aims to highlight the many commonalities but also some of the minor differences in the approaches to polymorphism that were taken in these different languages. As a final disclaimer, we do not advocate to code problems in an OO manner that can be easily coded in these languages in a functional way (as it is the case for this problem). However, in more complex cases, where many more nested functions would need to be used, and where state would have to be hidden, the OO programming style would be the more appropriate one. Hence, our test problem will stand in, in this chapter, for emulating such a more complex problem, that would benefit from the OO coding style.

### 3.1 Go

Go has supported run-time polymorphism through (polymorphic) “interfaces” (and thus modern-day OO programming) since its inception. In Go, encapsulation is done by storing state in a `struct` and by binding procedures, that need to use that state, to this same `struct`. Thereby creating a user-defined abstract data type (or ADT) with methods. Go allows the programmer to implement multiple polymorphic interfaces for such a type (i.e. to use multiple interface inheritance), even though it offers no explicit language statement for this purpose.

Instead, a user-defined type is implicitly assumed to implement an interface whenever it provides implementations of all the interface’s function signatures. This structural way of implementing interfaces also merely requires an object reference of the type to be passed to its methods (by means of a separate parameter list, in front of a method’s actual name). It is otherwise decoupled from the type’s (i.e. the ADT’s `struct`) definition. Go, finally, makes it explicit in its syntax that interfaces (like `structs`) are types in their own right, and that hence polymorphic variables (i.e. objects) can be declared in terms of them.

Restrictions in Go are that language-intrinsic types cannot have user-provided methods, and that methods and interfaces cannot be directly implemented for user-defined types whose definitions are located in other packages. That is, the programmer has to write wrappers in the latter case.

Since version 1.18, Go also supports compile-time polymorphism through generics. Go's generics make use of "strong concepts", since they are bounded by constraints that are expressed through interfaces. Hence, the Go compiler will fully type-check generic code. In Go, structures, interfaces, and functions, but not methods, can all be given their own generic type parameters.

### 3.1.1 Encapsulated version coded in Go

Listing 3.1 gives an encapsulated version of the test problem coded in Go. The two different implementations of the sum function have been encapsulated in two different ADTs named `SimpleSum` and `PairwiseSum`, whereas a third ADT named `Averager` encapsulates the functionality that is required to perform the actual averaging. The latter two ADTs contain the lower-level objects other and `drv` of `ISum[T]` type as components, to which they delegate calls to these objects' `sum` methods. Notice, how the use of the polymorphic interface `ISum[T]`, for the declarations of these objects, enables them to be initialized with either `SimpleSum` or `PairwiseSum` instances.

A second interface, named `IAverager`, is used to enable polymorphism for different averaging algorithms. Finally, there's a third interface, `INumeric`, that serves exactly the same purpose as in the functional polymorphic version that was given in Sect. 2.2, namely to make all function arguments and return values polymorphic, by admitting as input and output parameters both the `int32` and `float64` intrinsic types.

Hence, three polymorphic interfaces were required in this code, to eliminate the three levels of rigid dependencies that were listed in Sect. 2.1. Notice also that, exempting `INumeric`, all the interfaces and all the user-defined ADTs need to take in generic type parameters in this example. In Go, this is required to enable all the sum and average methods to use such generic type parameters.

Listing 3.1: Encapsulated Go version of the array averaging example.

```
1 package main
2
3 import "fmt"
4
5 // .....
6 // Interfaces
7 // .....
8
9 type INumeric interface {
10     int32 | float64
11 }
12
13 type ISum[T INumeric] interface {
14     sum(x []T) T
15 }
16
17 type IAverager[T INumeric] interface {
18     average(x []T) T
19 }
20
21 // .....
22 // SimpleSum ADT
23 // .....
24
25 type SimpleSum[T INumeric] struct {
26 }
27
28 func (self SimpleSum[T]) sum(x []T) T {
29     s := T(0)
```

```

30     for i := 0; i < len(x); i++ {
31         s += x[i]
32     }
33     return s
34 }
35
36 // .....
37 // PairwiseSum ADT
38 // .....
39
40 type PairwiseSum[T INumeric] struct {
41     other ISum[T]
42 }
43
44 func (self PairwiseSum[T]) sum(x []T) T {
45     if len(x) <= 2 {
46         return self.other.sum(x)
47     } else {
48         m := len(x) / 2
49         return self.sum(x[:m+1]) + self.sum(x[m+1:])
50     }
51 }
52
53 // .....
54 // Averager ADT
55 // .....
56
57 type Averager[T INumeric] struct {
58     drv ISum[T]
59 }
60
61 func (self Averager[T]) average(x []T) T {
62     return self.drv.sum(x) / T(len(x))
63 }
64
65 // .....
66 // main program
67 // .....
68
69 func main() {
70     var avi IAverager[int32]
71     var avf IAverager[float64]
72     var key int32
73
74     fmt.Println("Simple_sum_average:_1")
75     fmt.Println("Pairwise_sum_average:_2")
76     fmt.Print("Choose_an_averaging_method:_")
77     fmt.Scan(&key)
78
79     switch key {
80     case 1:
81         avi = Averager[int32]{ SimpleSum[int32]{} }
82         avf = Averager[float64]{ SimpleSum[float64]{} }
83     case 2:
84         avi = Averager[int32]{ PairwiseSum[int32]{ SimpleSum[int32]{} } }
85         avf = Averager[float64]{ PairwiseSum[float64]{ SimpleSum[float64]{} } }
86     default:
87         fmt.Println("Case_not_implemented!")
88         return

```

```

89     }
90
91     xi := []int32{1,2,3,4,5}
92     xf := []float64{1.,2.,3.,4.,5.}
93
94     fmt.Printf("%d\n", avi.average(xi))
95     fmt.Printf("%.1f\n", avf.average(xf))
96 }

```

The main program makes use of Go’s built-in structure constructors, and chaining of their calls, to instantiate objects of the required ADTs. In particular, it instantiates run-time polymorphic Averager objects (depending on whether simple or pairwise sum averaging is to take place), and it does so for both the `int32` and `float64` types separately, to then use these objects on `int32` and `float64` data, respectively.

That *two* such objects are required (one for each language-intrinsic data type) is connected to the aforementioned fact that in order to make methods use generic type parameters in Go, one has to parameterize interfaces, and instantiate these with different actual data types, as in `func main`’s first two code lines. A single (i.e. unparameterized) `IAverager` interface therefore doesn’t suffice, which is unfortunate from the user’s perspective, as some code duplication in client code cannot be avoided in this way.

It should also be noted, that Go’s intuitive way of expressing the generics constraints of interface `INumeric` in terms of a type set, besides having obvious advantages in terms of code clarity and conciseness, comes also at a price. Namely its only *partial* conformance to the Open/Closed Principle (OCP) of OO programming [10]. The latter principle demands that only new code be added to an application for extending its functionality, and any already written code to not be changed.

It can be seen, from Listing 3.1, that this principle can only be partly, but not strictly (i.e. fully), complied with in generic object-oriented Go code. Because for the programmer to make this listing work also with, e.g., the `float32` type, he would have to add this type to the type set of interface `INumeric`, and to thus touch already written code. Notice, though, that *none* of the implementation code of Listing 3.1 would be affected. Only the interface `INumeric` would need to be changed, that the remaining client code depends upon as a generics constraint.

If the client code were distributed among different packages or modules, the consequences of such a (weak) violation of the OCP would be confined to the occurrence of a recompilation cascade in the dependent modules. In actual practice, one may quite often regard this as an acceptable (minor) inconvenience, rather than a fundamental maintenance issue – especially in situations where the need to add new types is a relatively rare one.

## 3.2 Rust

Like Go, Rust supports both run-time and compile-time polymorphism through polymorphic interfaces, which Rust calls “traits”. Unlike Go, Rust has its programmers implement traits in a nominal manner, by using explicit `impl` code blocks to provide a trait’s method implementations. These same `impl` blocks can also be used to bind so-called “associated functions” to a type, which aren’t methods, i.e. which don’t take in a `self` (passed-object dummy) argument. A typical example for this are user-defined constructors. See the functions that are named `new` in the following code Listing 3.2, and are called using `::` syntax.

In contrast to Go, Rust allows the programmer to implement traits for both user-defined *and* language-intrinsic types, and to do so for types that are located in external libraries (called “crates”

in Rust), as long as the traits themselves are defined in the programmer’s own crate. The reverse, namely implementing an external trait for a user-owned type, is also possible. Only the (edge) case of implementing an external trait for an external type is not allowed (this is called the “orphan rule” [8]). The latter case requires the use of wrappers.

Comparable to Go, Rust’s generics model allows for the generic parameterization of functions, traits, and user-defined types like structs. Rust does not explicitly forbid generic methods. However, if one defines such a method’s signature within a trait, then this will make the trait unusable for the declaration of any “trait objects” [9], i.e. for the employment of run-time polymorphism. Thus, the Rust programmer will in general (need to) parameterize traits and structs rather than any methods themselves. Rust generics are fully type-checked at compilation time, i.e. Rust supports “strong concepts”.

### 3.2.1 Encapsulated version coded in Rust

The encapsulated Rust version of our test problem, that is given in the following Listing 3.2, is in its outline quite similar to the corresponding Go version<sup>1</sup>. There are, however, a few differences, that are listed in the following notes, and in our concluding remarks.

- Rust uses angled brackets, `< >`, to indicate generic parameter lists.
- Generics constraints in Rust are typically enforced by specifying the required traits in `impl` blocks using `where` statements.
- Rust does not offer an equivalent to Go’s type set syntax and semantics. In order to enable numeric operations on generic types, Rust instead provides a `Num` trait via an external `num` crate (i.e. library). However, the use of this external dependency makes available neither the `AddAssign` (`+=`) operator, nor casts to generic types within generic routines (as they are provided in Go).
- Since the employed algorithm relies on these features, the following code uses a homemade `INumeric` trait, that derives from Rust’s `Num` and `AddAssign` traits, and extends them by the requirement to also implement a constructor via an associated function, `new`, that provides the needed type casting (notice that `Self` stands in here for the implementing type). The conformance of the `i32` and `f64` intrinsic types to the `INumeric` trait is then acknowledged via `impl` blocks, that furthermore implement any of the still outstanding functionality of this trait for these two types.
- Where necessary, use of the `Copy` trait is also made, to work around Rust’s default move semantics.
- To help make all of the source code dependencies explicit, our Rust version employs `modules`, and `use` statements to import the required functionality.
- Rust’s default structure constructors exhibit the same restrictions as Fortran’s. That is, they are unable to initialize from an external scope, structure components that are private to their

---

<sup>1</sup>Notice that the present Rust version makes universal use of dynamic method dispatch via trait objects, to correspond most closely to all the other implementations that we provide in both the present chapter, and in Sect. 6.2.1. An alternative, more idiomatic, Rust version that is equivalent to the Fortran version which we’ll give in Sect. 6.2.2, and that effects static dispatch of the various `sum` methods through the use of generics, can be found in the `Code` subdirectory that is accompanying this document.

module. As in Fortran, use of user-defined constructors must be made instead (cf. the new functions that are defined in separate impl blocks for the ADTs PairwiseSum and Averager).

- To declare run-time polymorphic variables, one has to put so-called “trait objects” into “Boxes”, i.e. to declare smart pointers of them, for dynamic instantiation and heap memory allocation (this is the Rust equivalent to using allocatable polymorphic objects in Fortran).

Listing 3.2: Encapsulated Rust version of the array averaging example.

```
1 pub mod interfaces {
2
3     use num::Num;
4     use std::ops::AddAssign;
5
6     // .....
7     // Interfaces
8     // .....
9
10    pub trait INumeric: Num + AddAssign {
11        fn new(n: usize) -> Self;
12    }
13
14    pub trait ISum<T> {
15        fn sum(&self, x: &[T]) -> T;
16    }
17
18    pub trait IAverager<T> {
19        fn average(&self, x: &[T]) -> T;
20    }
21 }
22
23 pub mod intrinsics {
24
25     use crate::interfaces::INumeric;
26
27     impl INumeric for i32 {
28         fn new(n: usize) -> i32 {
29             return n as i32
30         }
31     }
32
33     impl INumeric for f64 {
34         fn new(n: usize) -> f64 {
35             return n as f64
36         }
37     }
38 }
39
40 pub mod simple_library {
41
42     use crate::interfaces::{INumeric, ISum};
43
44     // .....
45     // SimpleSum ADT
46     // .....
47
48     pub struct SimpleSum;
49 }
```

```

50 impl<T> ISum<T> for SimpleSum where T: INumeric + Copy {
51     fn sum(&self, x: &[T]) -> T {
52         let mut s = T::new(0);
53         for i in 0 .. x.len() {
54             s += x[i];
55         }
56         return s
57     }
58 }
59 }
60
61 pub mod pairwise_library {
62
63     use crate::interfaces::{INumeric,ISum};
64
65     // .....
66     // PairwiseSum ADT
67     // .....
68
69     pub struct PairwiseSum<T> {
70         other: Box<dyn ISum<T>>,
71     }
72
73     impl<T> PairwiseSum<T> where T: INumeric {
74         pub fn new(other: Box<dyn ISum<T>>) -> PairwiseSum<T> {
75             PairwiseSum{
76                 other: other,
77             }
78         }
79     }
80
81     impl<T> ISum<T> for PairwiseSum<T> where T: INumeric {
82         fn sum(&self, x: &[T]) -> T {
83             if x.len() <= 2 {
84                 return self.other.sum(x);
85             } else {
86                 let m = x.len() / 2;
87                 return self.sum(&x[..m+1]) + self.sum(&x[m+1..]);
88             }
89         }
90     }
91 }
92
93 pub mod averager_library {
94
95     use crate::interfaces::{INumeric,ISum,IAverager};
96
97     // .....
98     // Averager ADT
99     // .....
100
101     pub struct Averager<T> {
102         drv: Box<dyn ISum<T>>,
103     }
104
105     impl<T> Averager<T> where T: INumeric {
106         pub fn new(drv: Box<dyn ISum<T>>) -> Averager<T> {
107             Averager{
108                 drv: drv,

```

```

109     }
110 }
111 }
112
113 impl<T> IAverager<T> for Averager<T> where T: INumeric {
114     fn average(&self, x: &[T]) -> T {
115         return self.drv.sum(&x) / T::new(x.len());
116     }
117 }
118 }
119 // .....
120 // main program
121 // .....
122 #[macro_use] extern crate text_io;
123
124 fn main() {
125     use std::io;
126     use std::io::Write;
127     use crate::interfaces::IAverager;
128     use crate::simple_library::SimpleSum;
129     use crate::pairwise_library::PairwiseSum;
130     use crate::averager_library::Averager;
131
132     let avi: Box<dyn IAverager::<i32>>;
133     let avf: Box<dyn IAverager::<f64>>;
134     let key: i32;
135
136     println!("Simple_sum_average:_1");
137     println!("Pairwise_sum_average:_2");
138     print!("Choose_an_averaging_method:_");
139
140     io::stdout().flush().unwrap();
141     scan!("{}",key);
142
143     let ssi = SimpleSum{};
144     let ssf = SimpleSum{};
145     let psi = PairwiseSum::new(Box::new(SimpleSum{}));
146     let psf = PairwiseSum::new(Box::new(SimpleSum{}));
147
148     match key {
149         1 => { avi = Box::new(Averager::new(Box::new(ssi)));
150               avf = Box::new(Averager::new(Box::new(ssf))); }
151         2 => { avi = Box::new(Averager::new(Box::new(psi)));
152               avf = Box::new(Averager::new(Box::new(psf))); }
153         _ => { println!("Case_not_implemented!");
154               return; }
155     }
156
157     let xi: [i32;5] = [1,2,3,4,5];
158     let xf: [f64;5] = [1.,2.,3.,4.,5.];
159
160     println!("{}", avi.average(&xi));
161     println!("{:.1}", avf.average(&xf));
162 }

```

The Rust program version is somewhat longer than the corresponding Go version, because user-defined constructors had to be provided, to implement the aforementioned custom type conversion functionality and to initialize module-hosted (opaque) abstract data types, and because

dependencies from external libraries, but also from modules of the user’s own crate, had to be imported into the main function (as it would be necessary in realistic situations). The main function’s logic could also not be expressed as concisely as in the Go version, because the various trait objects required (manual) “boxing” for their instantiation, and because Rust’s default move semantics had to be worked around.

The most important difference between the two languages, though, is that Rust uses traits (i.e. interfaces) as generics constraints whose function signatures are always specified explicitly, whereas Go’s are typically specified implicitly (through type sets). Together with Rust’s capability (that Go lacks) to explicitly implement traits even for language-intrinsic types, this allows Rust code to fully (i.e. strictly) conform to OO programming’s Open/Closed Principle. In order to enable, for instance, the code of Listing 3.2 to accept also the `f32` type, one would simply need to add a further module, that implements the `INumeric` trait for this new intrinsic type. None of the pre-existing code would have to be changed.

### 3.3 Swift

Being a successor language to Objective-C, Swift differs slightly from the languages considered so far in that it opted to retain implementation inheritance for backwards compatibility to Objective-C, whereas both Go and Rust do not support implementation inheritance *by design*. Swift therefore supports “classical” classes, but it also allows one to bind methods to structures (which, in contrast to classes, are value types in Swift).

Like Go and Rust, Swift (furthermore) supports a traits system to implement both run-time and compile-time polymorphism through polymorphic interfaces, that are called “protocols” in Swift. If the Swift programmer chooses to ignore implementation inheritance and classes, he can therefore very much program with structures and protocols in Swift as he would with structures and interfaces/traits in Go and Rust, respectively.

Given Swift’s backwards compatible design, implementation of a protocol (i.e. interface inheritance) is usually done as in classical OO languages, i.e. within a structure’s or a class’s definition. A colon (:) followed by one or more interface names must be supplied for this purpose after the structure’s or class’s own name. However, a very powerful facility for types to implement protocols retroactively is also provided, through so-called extensions, that work even if the types’ source code is inaccessible (because one is, e.g., working with a library in binary form). In order to support strict conformance to the Open/Closed Principle of OO programming, this same facility allows protocols to be implemented also for language-intrinsic types. For instance, the following little program, given by Listing 3.3, prints out “I am 4.9”.

Listing 3.3: Swift example of implementing a protocol for an intrinsic data type.

```
1 protocol Printable {
2     func output()
3 }
4 extension Float64: Printable {
5     func output() {
6         print("I_am_\(self)")
7     }
8 }
9 var y: Float64 = 4.9
10 y.output()
```

Swift generics support “strong concepts”, and are thus fully type-checked at compile time, and their capabilities are at least on par with those of Go and Rust. In several aspects they are actually

superior. For instance, Swift allows function overloading to be used together with generics [5, 9]. Swift also allows for parameterized *methods*, instead of parameterized protocols. The latter point has some interesting, positive implications for the Swift programmer, that will be discussed in detail below.

### 3.3.1 Encapsulated version coded in Swift

Listing 3.4 gives an example of how the encapsulated version of the array averaging test problem can be programmed in Swift. See the following remarks, to understand this code.

- By default, function and method calls in Swift need to make use of mandatory keyword arguments.
- Array slices are not arrays themselves. Hence, an explicit conversion of such slices via an `Array()` constructor is required to use them as arrays.
- Like Rust, Swift uses angled brackets to indicate generic parameter lists. Type constraints are formulated within these lists by supplying a protocol name after a generic type parameter (separated by a colon).
- Similar to Rust, Swift provides a `Numeric` protocol (that is defined in its standard library) for use as a generics constraint for numeric types, which however does *not* include the division operation!
- It is easy to extend this protocol to support also the division operator, similar to the way we've extended traits in the Rust code example. However, the following code goes a step further than that, and foregoes any dependence on library functionality whatsoever, by defining a custom `INumeric` protocol that prescribes the interfaces of all the required operators and type initializers (i.e. constructors) explicitly.
- The present Swift code makes use of the language's default, built-in, initializers for the intrinsic types and the `SimpleSum` struct that it employs, and user-defined initializers for all the structs that need to initialize private field variables.
- Type conversion into a generic type, `T`, is effected in Swift by calling an appropriate initializer of this type. The call syntax is equivalent to that of Go's generic casts, but is a bit peculiar in that, e.g., Go's `T(0)` is written as `T(exactly:0)!` in Swift (making use of both the [default] initializer's mandatory keyword, `exactly`, and the `!` operator to unwrap the returned optional type). Notice, how the initializer's signature needs to be prescribed in protocol `INumeric`, in order to be able to use this functionality with some conforming type `T`.
- Swift's `Int32` and `Float64` intrinsic types already implement all of protocol `INumeric`'s functionality by default. It therefore suffices to simply acknowledge this fact through empty extension statements for these two types.

Listing 3.4: Encapsulated Swift version of the array averaging example.

```

1 // .....
2 // Interfaces
3 // .....
4
5 protocol INumeric {
6     init?(exactly: Int)
7     static func += (lhs: inout Self, rhs: Self)
8     static func + (lhs: Self, rhs: Self) -> Self
9     static func / (lhs: Self, rhs: Self) -> Self
10 }
11
12 protocol ISum {
13     func sum<T: INumeric>(x: [T]) -> T
14 }
15
16 protocol IAverager {
17     func average<T: INumeric>(x: [T]) -> T
18 }
19
20 // .....
21 // Intrinsics
22 // .....
23
24 extension Int32: INumeric {}
25 extension Float64: INumeric {}
26
27 // .....
28 // SimpleSum ADT
29 // .....
30
31 struct SimpleSum: ISum {
32     func sum<T: INumeric>(x: [T]) -> T {
33         var s = T(exactly:0)!
34         for i in 0 ..< x.count {
35             s += x[i]
36         }
37         return s
38     }
39 }
40
41 // .....
42 // PairwiseSum ADT
43 // .....
44
45 struct PairwiseSum: ISum {
46     private let other: any ISum
47
48     init(other: any ISum) {
49         self.other = other
50     }
51
52     func sum<T: INumeric>(x: [T]) -> T {
53         if x.count <= 2 {
54             return other.sum(x: x)
55         } else {
56             let m = x.count / 2
57             return sum(x: Array(x[..<m])) + sum(x: Array(x[m...]))

```

```

58     }
59 }
60 }
61
62 // .....
63 // Averager ADT
64 // .....
65
66 struct Averager: IAverager {
67     private let drv: any ISum
68
69     init(drv: any ISum) {
70         self.drv = drv
71     }
72
73     func average<T: INumeric>(x: [T]) -> T {
74         return drv.sum(x: x) / T(exactly: x.count)!
75     }
76 }
77
78 // .....
79 // main function
80 // .....
81
82 func main() {
83     let av: any IAverager
84     let key: Int32?
85
86     print("Simple_sum_average:_1")
87     print("Pairwise_sum_average:_2")
88     print("Choose_an_averaging_method:_", terminator: "")
89     key = Int32(readLine()!)
90
91     switch key {
92     case 1:
93         av = Averager(drv: SimpleSum())
94     case 2:
95         av = Averager(drv: PairwiseSum(other: SimpleSum()))
96     default:
97         print("Case_not_implemented!")
98         return
99     }
100
101     let xi: [Int32] = [1,2,3,4,5]
102     let xf: [Float64] = [1.0,2.0,3.0,4.0,5.0]
103
104     print( av.average(x: xi) )
105     print( av.average(x: xf) )
106 }
107
108 // execute main function
109 main()

```

Even a casual glance at the Swift version will show that the Swift code is the easiest to read and understand among the generic object-oriented implementations that were presented in this chapter. This is largely the result of Swift supporting generic methods, and hence not requiring the programmer to parameterize and instantiate any generic interfaces (protocols), in contrast to both Go and Rust. The consequences are

- that method genericity for an ADT's objects can be expressed using only a single, as opposed to multiple protocols,
- that therefore merely a *single* object instance of that same protocol is required, to be able to operate on many different language-intrinsic data types, and
- that this also (largely) *obviates the need for manual instantiations of generics in Swift* (because generic functions/methods are easier to instantiate automatically by the compiler, as it can almost always infer the required types by checking the regular arguments that are passed to a function/method)!

As an example, consider the object `av` in the above Swift code that contains the functionality for array averaging. This object supports two different levels of polymorphism: Firstly, given that it is an instance of the `IAverager` protocol, it can be polymorphically assigned different averaging algorithms (see the `switch` statement). Secondly, because it contains an `average` method that is generic, it can be used on data of different intrinsic types, like `Int32` and `Float64` here.

Notice that the `main` function in the Swift code needs to declare merely a *single* such object variable of `IAverager` type, to make use of all these capabilities. This is a direct consequence of there being only a single (i.e. unparameterized) version of the `IAverager` protocol, and of parameterizing the protocol's method signatures by generic types rather than the protocol itself.

Contrast this with Go's and Rust's model, where not only separate objects of `IAverager` type are required for *every* different intrinsic data type that the programmer wishes to use these objects with. But where also *manual* instantiations of corresponding versions of the generically parameterized `IAverager` interface/trait are required from the programmer, for declaring these objects. Swift's generics model gets rid of all of that complexity, and therefore vastly simplifies client code. We consider this a very significant advantage of the generics approach that is taken in Swift vs. that of Go and Rust.

## 3.4 Conclusions

The use of run-time polymorphism by means of (polymorphic) interfaces is rather similar in all the languages that were considered here. The most significant difference in this respect is that Go has stricter limitations than the other languages regarding code ownership, when it comes to retroactively implementing interfaces for existing types. Whereas Rust (with some minor restrictions), and Swift allow the implementation of an interface by some type to be accomplished effectively without regard to the type's definition site. Rust and Swift thereby overcome Haverdaen et al.'s critique [6] of Java regarding this point. In fact, it is *interface inheritance* which makes the uniform polymorphic treatment of both intrinsic and user-defined types possible in the first place in Rust and Swift, that Haverdaen et al. seem to also (rightly) demand. All the considered languages are also quite similar in that they support fully type checked generics via the mechanism of interfaces. In the following, we will thus focus on summarizing the most significant differences in these languages' generics features.

### 3.4.1 Go

Go's basic model to implement generics allows structures, interfaces, and ordinary functions, but not methods, to be given their own generic type parameters. The lack of true generic methods makes some duplication of instantiation code in clients unavoidable. Nevertheless, generic Go

code is quite easy to read and to understand. Go features built-in, easy to use support for conversion to generic types. Yet, where Go truly differs from both Rust and Swift are the unique, and extremely useful, implicit mechanisms for conformance to interfaces that it supports, namely structural subtyping, and the brilliant new notion of formulating interfaces in terms of type sets, to which the member types of these sets conform by definition.

The latter idea, along with the corresponding syntax to support it, enables the Go programmer to formulate generics constraints very naturally and concisely, without having to explicitly implement interfaces for this purpose. This is what makes the use of generics in Go pleasant. It is, moreover, essential for supporting concise generic procedural and functional programming. While type sets retain their advantages also in generic object-oriented programming, their disadvantage in this setting is that they allow for only partial, but not full, conformance to object-orientation's Open/Closed Principle.

### 3.4.2 Rust

Rust's basic model for generics is similar to Go's in that it allows for parameterization of structures, interfaces, and ordinary functions, but not necessarily methods. Hence, what has been said above for Go in this respect holds also for Rust. Rust has, unfortunately, some quirks which render its use for the management of all types of dependencies through polymorphism somewhat sub-optimal when compared to the other languages considered here. The language is unpleasant to use, because of its "borrow checker", its employment of move semantics by default, its *excessive* obsession with type safety, and its overall C++-like philosophy to copiously rely on external dependencies, even for the most basic tasks.

Because of the latter two points, the Rust version of our test case is marred by dependencies on external libraries, which is quite contrarian to the purpose of programming in a polymorphic fashion, namely to avoid rigid dependencies. Even with the functionality provided by these external dependencies, casts to generic types within generic routines weren't outright possible, and had to be achieved, instead, through some work on the programmer's side. The points we like most about the language are its idea to decouple trait (i.e. interface) implementations from a struct's definition through `impl` blocks that work with both user-defined and intrinsic types, the full conformance to the Open/Closed Principle that this enables, and the complete control over the use of dynamic vs. static method dispatch (via trait objects and generics, respectively) that Rust affords the programmer. These are the features of Rust that, in our opinion, Fortran should borrow in some form.

### 3.4.3 Swift

Swift's basic model of implementing generics by allowing parameterized structures, functions, and methods (but not parameterized interfaces) is both the easiest to read, and the easiest to use from a programmer's perspective. In this model, if a programmer knows how to write generic functions, his knowledge automatically translates into coding generic methods. Swift's generics design supports casts to generic types, regardless of whether these types are user- or language-defined. Moreover, the Swift compiler is able to instantiate generics largely automatically, through type inference of the regular arguments that are passed to functions, methods, and (structure or class) constructors. In contrast to the other languages, in Swift, the user almost never has to bother with instantiating any generics.

Like Rust, Swift enables explicit and retroactive implementation of interfaces for both user-defined and intrinsic types, and the full conformance to the Open/Closed Principle of OOP that

this affords one. Unlike Rust and Go, Swift offers generics that are interoperable with both class inheritance and ad hoc polymorphism (i.e. function overloading), which are two features that are also an integral part of Fortran. Out of all the considered modern languages, Swift therefore shares the largest similarities with Fortran. Moreover, there is notable documentation available [5, 16] on the model with which Swift’s generics are implemented in the Swift compiler. Finally, and very importantly, with this implementation model, compilation times scale only with the amount of generic code *written*, rather than the amount of code *instantiated* [5]. This is in stark contrast to templated approaches (as used e.g. in C++).

For all these reasons, we consider Swift’s generics to be the by far most attractive model to base Fortran’s generics capabilities on. The fact that Swift is a language that does not put emphasis on numerics, and whose present standard library therefore does not provide a truly useful Numeric protocol/interface, is of absolutely *no* consequence for adopting Swift’s generics design as a baseline for Fortran. We believe, though, that this baseline design would benefit substantially from being supplemented with Go’s implicit mechanisms for interface conformance.

Especially for properly supporting generic procedural and functional programming in Fortran, i.e. for making them *sufficiently concise*, Go-like type sets should be added. These would also be useful in generic object-oriented programming, even though Swift’s design already comes with good support for the latter, right out of the box. Its single minor disadvantage in this respect, namely the need to always provide (possibly empty) explicit extension blocks to implement a new interface, even if the involved types already command over all of the functionality that is prescribed by that interface, could be overcome by allowing for Go-like structural, in addition to nominal, subtyping. A conclusion that seems to have been arrived at also in the recent Mojo language [15]. See Sects. 6.2.3 and 7.2 for some further discussion of this point.

## Chapter 4

# Fortran extensions I: Traits for types

The present and the next chapter, describe a number of simple extensions to Fortran, that we consider essential for enabling dependency management through polymorphism at a level of functionality that is on par with modern languages like Swift, Rust, Go, Carbon, or Mojo. The present chapter adds general subtyping capabilities to Fortran, while the next chapter aims at providing specific support for generics. The extensions that are related to subtyping concern abstract interfaces, derived types, and the `class` specifier for variable declarations. They also encompass an important, new, Fortran feature: the `implements` statement.

### 4.1 Traits (named abstract interfaces)

The most essential of all the following extensions is the capability to define named abstract interfaces, or traits (i.e. named collections of procedure signatures), in order to suitably constrain (and hence to type-check) the declarations of polymorphic variables. Traits are the crucial feature that is required to uniformly and properly express both run-time and compile-time polymorphism (i.e. modern-day OO and generic programming, respectively) in the language, and to thereby enable a uniform management of dependencies on both user-defined *and* language-intrinsic types.

#### 4.1.1 Definition

Fortran already allows the programmer to define unnamed abstract interfaces. But to use these as traits, named versions of them are required. Moreover, and in order to support all of the functionality that is discussed in the next sections, these interfaces must admit signatures of not only ordinary type-bound procedures, but also initializers (i.e. constructors), and operators (including assignments). The following example defines three such traits, `ICastable`, `ICalculable`, and `IPrintable`, that are intended as abstract blueprints for actual type-bound implementations of, respectively, an initializer, `init`, a function, `saxpy`, and a subroutine, `output`:

```
1  abstract interface :: ICastable
2      function init(n)
3          integer, intent(in) :: n
4      end function init
5  end interface ICastable
6
7  abstract interface :: ICalculable
8      function saxpy(a,x,y) result(res)
9          real, intent(in) :: a, x(:), y(:)
10         real :: res(:)
```

```

11  end function saxpy
12  end interface ICalculable
13
14  abstract interface :: IPrintable
15      subroutine output()
16      end subroutine output
17  end interface IPrintable

```

Examples of traits with signatures of type-bound operators are given in Sects. 5.1.3, 6.2.3, and 6.3. Notice, how initializers within traits are declared as *special* functions, that have a reserved name, `init`, and for which no return type is specified. Because the latter is considered to be already fully determined by the initialized type. This is equivalent to how initializers are declared in Swift protocols.

It should also be pointed out that traits contain the signatures of *messages*, which in OO and generic programming may be sent to different receivers. A well-designed traits facility should therefore *not* require such message signatures to contain passed-object dummy arguments, or their types, i.e. any details concerning the receivers. Because in order to be sendable to different receivers, a message must not know anything about its receiver. This is clearly exhibited by the present examples, and also the Go and Swift examples<sup>1</sup> of Listings 3.1 and 3.4. But it is violated in present Fortran practice by deferred type-bound procedures. We will elaborate further on this in Sect. 4.3.3. There, we will also show how to fix this problem, while achieving interoperability of the new traits facility with the inheritance feature that is already available in the present language.

### 4.1.2 Overloading

In the examples that were considered so far, traits were shown to contain uniquely named message signatures. But since Fortran supports ad hoc polymorphism through “generic” overloading, traits must also admit *overloaded* versions of initializer, method, operator, and assignment signatures – that share a generic name, but differ in their arguments lists. The following trait, `IConstructable`, illustrates this for a type-bound initializer, by overloading the reserved name `init` with two such signatures:

```

1  abstract interface :: IConstructable
2      function init(n)
3          integer, intent(in) :: n
4      end function init
5      function init(a)
6          real, intent(in) :: a
7      end function init
8  end interface IConstructable

```

Traits for all the other cases can be formulated analogously – through overloaded use of a method name, or the reserved identifiers operator(`op`), and assignment(`=`), respectively (with `op` as per [18]). An example of an overloaded multiplication operator signature is given in Sect. 6.3.

### 4.1.3 Extends attribute

Abstract interface definitions must allow the programmer to define new traits that inherit procedure signatures from *multiple* simpler traits (multiple interface inheritance). In the following

<sup>1</sup>Notice, especially, that despite the fact that Go (like Fortran) uses explicitly passed receiver objects, these appear only in the actual *implementations* of methods, and even there they are clearly *syntactically separated* from the latter’s interfaces.

example, the trait `IAdmissible` inherits the procedure signatures that are contained in both the `ICastable` and `IPrintable` traits of Sect. 4.1.1, making `IAdmissible` at the same time a *subtype* of both these simpler traits:

```
1 abstract interface, extends(ICastable + IPrintable) :: IAdmissible
2 end interface IAdmissible
```

That is, objects that implement (or adopt) the `IAdmissible` trait (i.e. conform to it), can also be used in settings that require conformance to either the `ICastable`, or `IPrintable` traits.

The next example finally shows an extended version of the `IAdmissible` trait, that not only inherits the aforementioned procedure signatures, but also adds the signature of a further method, `input`, to them:

```
1 abstract interface, extends(ICastable + IPrintable) :: IAdmissible
2   subroutine input(filename)
3     character(*), intent(in) :: filename
4   end subroutine input
5 end interface IAdmissible
```

## 4.2 Implements statement

The language must allow not only for traits to conform to other traits (cf. Sect. 4.1.3), but also for both language-intrinsic and user-defined types to do the same. That is, it must be possible for these types to acknowledge their conformance to a trait, if they already command over implementations of all its functionality, or to provide any still outstanding implementations that are required for such conformance. Moreover, this must be possible even retroactively, i.e. regardless of (and without having to touch) any original type definition site<sup>2</sup>. Otherwise, wrapper types would, in general, need to be written, once an already defined type would need to implement some new trait (e.g. in cases where the original type definition is inaccessible). This also means that the provision of user-defined functionality needs to be allowed even for *language-intrinsic* types<sup>3</sup>. The `implements` statement, that is described in this section, provides these capabilities.

### 4.2.1 Overview

The `implements` statement is modeled after the “extension” feature of Swift, to enable retroactive implementation of new methods, initializers, operators, and especially traits for types, in order to *dynamically* change a subtyping (i.e. interface inheritance) hierarchy, and to thus achieve utmost code flexibility. Swift’s extension blocks fulfill essentially the same purpose as Rust’s `impl` blocks. They have been somewhat simplified here, and adjusted to Fortran’s idiosyncracies and syntax, that binds procedures to types through declaration blocks, rather than by including the actual implementations themselves into such a block (the implementations need to be supplied as module procedures, as it is usual in Fortran).

The syntax of the feature is an extended subset of that for the contains sections of derived type definitions. That is, `implements` statements make use of type-bound procedure declarations. These employ the already well-known attributes `public`, `private`, `pass[(arg-name)]`, `nopass`, and `non_overridable`, along with the new attribute `static`. Some additional syntax for these declarations is provided, to enable direct overloading of method names, and to thus make it less verbose than through the use of generic type-bound procedure declarations (as the Fortran standard calls

<sup>2</sup>Exempting only the case of abstract derived types, as it is discussed below.

<sup>3</sup>With some restrictions. See Sect. 4.2.4.

them). However, the latter are supported as well. Furthermore, a new type of initial procedure declarations is available, for specifying type-bound initializers (i.e. constructors).

The main difference to derived type definitions is that the `implements` statement does not accept declarations of deferred type-bound procedures (which we propose to retire partly or entirely, see Sect. 4.3.3), and that it cannot be used with abstract derived types. If interoperability with abstract types is required, the later to be discussed `implements` derived type attribute (cf. Sect. 4.3) can be used instead. As a final disclaimer, notice that the `implements` statement has *no* relation to subclassing, i.e. one derived type being extended into another through (rigid) implementation inheritance. Rather, this is a feature that adds new capabilities to a single, *given*, type.

## 4.2.2 Binding functionality to a type

The most basic use of the `implements` statement is to (retroactively) bind procedures of varying functionality to a type. The following example shows how to bind the names of an overloaded, user-provided, initializer, and an ordinary method to a derived type with name `MyType`, and how to provide their actual implementations as module procedures:

```

1 module basic
2   ...
3   type :: MyType
4     private
5     real :: a
6   end MyType
7
8   implements :: MyType
9     initial :: init => init_by_int, init_by_real
10    procedure, pass :: output
11  end implements MyType
12
13 contains
14
15  function init_by_int(n) result(res)
16    integer, intent(in) :: n
17    type(MyType)        :: res
18    res%a = real(n)
19  end function init_by_int
20
21  function init_by_real(a) result(res)
22    real, intent(in) :: a
23    type(MyType)      :: res
24    res%a = a
25  end function init_by_real
26
27  subroutine output(self)
28    class(MyType), intent(in) :: self
29    write(*,*) "I_am_", self%a
30  end subroutine output
31
32 end module basic

```

Notice the use of the new declaration of `initial` type, that instructs the compiler to overload the implementing type's structure constructor with, in this case, multiple type-bound initializer implementations of the (reserved) generic name `init`. Notice, also, that the return type of a module function that actually implements an initializer is always the concrete type that it initializes.

The above example also demonstrates how an ordinary method of name `output` can be bound

to MyType, in quite the same manner as this is accomplished in the present language. That is, by declaring the passed-object dummy argument, `self`, with the `class` declaration specifier (as it is required for my MyType to be usable as a base type for implementation inheritance).

### 4.2.3 Implementing traits

#### Retroactive implementation

Assume now, that the purpose of binding the previous procedures to MyType was to actually make this type compatible with settings where conformance to the IPrintable and IConstructable traits of, respectively, Sects. 4.1.1 and 4.1.2 is required. So far, we have provided both the bindings of the required initializers and method, as well as their implementations – making sure that the signatures of the latter match those that are contained in the IConstructable and IPrintable traits, by ignoring any passed-object arguments (which a compiler will simply skip, in the present design, when checking for trait, i.e. interface, conformance of an implementation).

However, we haven't made MyType pluggable yet into code that is written in terms of these traits. To fix this, we can simply acknowledge (even from a different module, as in the following example) that MyType already has all of the required functionality to implement these traits:

```
1 module enhanced
2   ...
3   use basic, only: MyType
4
5   implements (IConstructable + IPrintable) :: MyType
6   end implements MyType
7
8 end module enhanced
```

It is crucial for flexibility, that the subtyping (i.e. interface inheritance) mechanism, that the `implements` statement provides, allow for a type to implement *multiple* different traits – as it is demonstrated in this example.

#### Collective implementation

The last code example has shown how traits can be implemented after the procedure implementations, that they require, were already bound to some type. Of course, it is also possible to do all of this at one fell swoop, as in the following alternative version of the module `basic`, whose original version was given in Sect. 4.2.2:

```
1 module basic
2   ...
3   type :: MyType
4     private
5     real :: a
6   end MyType
7
8   implements (IConstructable + IPrintable) :: MyType
9     initial :: init => init_by_int, init_by_real
10    procedure, pass :: output
11  end implements MyType
12
13 contains
14   ...
15 end module basic
```

To avoid needless repetition, we have omitted here the implementation of the actual procedures, that would be done exactly as it was shown in Sect. 4.2.2.

### Split implementation

We could have also employed two separate `implements` statements, in order to implement one trait at a time, and achieve the same effect:

```
1 module basic
2   ...
3   type :: MyType
4     private
5     real :: a
6   end MyType
7
8   implements IConstructable :: MyType
9     initial :: init => init_by_int, init_by_real
10  end implements MyType
11
12  implements IPrintable :: MyType
13    procedure, pass :: output
14  end implements MyType
15
16 contains
17   ...
18 end module basic
```

Here we have skipped again, for brevity, the actual implementations. Such a splitting of `implements` statements can be useful to improve code readability, as it makes the association between the traits and the actual procedures, that are to be implemented for any one of them, immediately obvious. These two statements (together with the actual implementations) could then have been distributed even among different modules and files. Notice, also, how parentheses around traits lists in `implements` statements are optional, but not required.

#### 4.2.4 Implementing traits for intrinsic types

The `implements` statement can be used to provide new functionality also for types that are *intrinsic* to the language – in quite the same fashion as it was already demonstrated for derived types. However, the present language features some restrictions in this context, that we deem essential to retain. In particular, present Fortran doesn't allow one to overload, i.e. to change, the implementations of the intrinsic operators that the language defines for intrinsic types. Which is what makes numeric Fortran code *predictable*. Any present restrictions that serve this latter purpose should therefore continue to be honored.

The following listing, which shows how to code the Swift example of Listing 3.3, is a case that is unaffected by these considerations. It implements the `IPrintable` trait, with its single method `output`, for Fortran's `real(real64)` type, to perform printouts for variables of this type:

```
1 module real64_module
2
3   use, intrinsic :: iso_fortran_env, only: real64
4
5   abstract interface :: IPrintable
6     subroutine output()
7     end subroutine output
8   end interface IPrintable
```

```

9
10 implements IPrintable :: real(real64)
11   procedure, pass :: output
12 end implements real(real64)
13
14 contains
15
16   subroutine output(self)
17     real(real64), intent(in) :: self
18     write(*,*) "I_am_", self
19   end subroutine output
20
21 end module real64_module
22
23 program printy
24
25   use, intrinsic :: iso_fortran_env, only: real64
26
27   real(real64) :: y
28
29   y = 4.9d0
30   call y%output()
31
32 end program printy

```

Notice, that once an `implements` statement has been used to augment the functionality of an existing (intrinsic or derived) type (like the `real(real64)` type in the above example), then this new functionality will be available on *all* instances of that type, even if they were created before this augmentation was defined. As it is the case in the Swift language [21], this prevents the occurrence of any ambiguities regarding the available functionality of such type instances. For the same reason, it is sufficient to import into client code (as usual) either a named constant corresponding to an intrinsic type's kind parameter (like the `real64` constant in the above `printy` example program), or a derived type's (original) definition, in order to make use of the thus augmented functionality.

### 4.3 Implements derived type attribute

For reasons of both regularity in the language, and interoperability with Fortran's legacy (i.e. present) OO model, it should be possible to implement traits also directly from within derived type definitions. This can be accomplished by allowing an `implements` *attribute* within such definitions, that – in contrast to the `implements` statement – features also interoperability with abstract derived types, and their extensions by implementation inheritance.

By its very nature of being bound to derived type definitions, the former attribute, though, necessarily lacks the latter statement's retroactive implementation capabilities, that are important for exploiting the full possibilities of the traits-based OO programming model. The latter model is a much more flexible, complete, modern-day alternative to both abstract types and type extension (i.e. implementation inheritance). Hence, our recommendation is to use the interoperability features, that are described in the following Sects. 4.3.2 and 4.3.3, mainly as migration tools (for modernizing legacy applications incrementally, through refactoring), and to otherwise avoid the use of both abstract types and type extension in new programs.

### 4.3.1 Use with the sealed attribute

The following code shows how the “Collective implementation” example of Sect. 4.2.3, that was written there in terms of an `implements` statement, can be reformulated to use an `implements` derived type attribute:

```
1 module basic
2   ...
3   type, sealed, implements(IConstructable + IPrintable) :: MyType
4   private
5     real :: a
6   contains
7     initial :: init => init_by_int, init_by_real
8     procedure, pass :: output
9   end implements MyType
10
11 contains
12
13   function init_by_int(n) result(res)
14     integer, intent(in) :: n
15     type(MyType)        :: res
16     res%a = real(n)
17   end function init_by_int
18
19   function init_by_real(a) result(res)
20     real, intent(in) :: a
21     type(MyType)     :: res
22     res%a = a
23   end function init_by_real
24
25   subroutine output(self)
26     type(MyType), intent(in) :: self
27     write(*,*) "I_am_", self%a
28   end subroutine output
29
30 end module basic
```

This example also demonstrates the use of a second, new, derived type attribute, the `sealed` attribute. Derived types that are sealed are inextensible by type extension, i.e. implementation inheritance. Thus, there’s also no need for the passed-object dummy arguments of such types to be polymorphic, and hence for the programmer to declare them with the `class` specifier. One can use the type specifier, instead, as it is demonstrated by the implementation of method `output` in the above example.

### 4.3.2 Use with the extends attribute

The previous Sect. 4.3.1 has illustrated how the `implements` derived type attribute would be typically used in modern code (that avoids the use of implementation inheritance, or subclassing). However, the multiple interface inheritance (or subtyping) features, that are provided by the `implements` derived type attribute<sup>4</sup>, can actually be combined with the type extension capabilities which are already present in the language.

In the following example, a `Parent` type is defined, that implements a `method1`. The implementation of the latter is then inherited by a `Child` type, that, in turn, implements an own `method2`. In this way, the `Child` type now commands over implementations of both `method1`, and `method2`, as

---

<sup>4</sup>And, with some restrictions, also the `implements` statement.

it is required for it to conform to both the IParent and IChild traits, respectively, that it adopts. Which consist of the signatures of these methods.

```

1  module extends_parent
2
3      abstract interface :: IParent
4          subroutine method1(n)
5              integer, intent(in) :: n
6          end subroutine method1
7      end interface IParent
8
9      abstract interface :: IChild
10         subroutine method2(res)
11             real, intent(out) :: res
12         end subroutine method2
13     end interface IChild
14
15     type, implements(IParent) :: Parent
16     contains
17         procedure, pass :: method1 => implementation1
18     end type Parent
19
20     type, extends(Parent), implements(IParent + IChild) :: Child
21     contains
22         procedure, pass :: method2 => implementation2
23     end type Child
24
25 contains
26
27     subroutine implementation1(self,n)
28         class(Parent), intent(in) :: self
29         integer,          intent(in) :: n
30     end subroutine implementation1
31
32     subroutine implementation2(self,res)
33         class(Child), intent(in)  :: self
34         real,          intent(out) :: res
35     end subroutine implementation2
36
37 end module extends_parent

```

In cases like the present one, where the extends and implements derived type attributes are used in combination, extends shall always precede implements. Since the latter attribute allows for the inheritance of *multiple* interfaces/traits, as it is demonstrated in this example, this also fixes present Fortran’s single inheritance limitations, *without* introducing the potential ambiguities that multiple inheritance of implementation would cause (which are also known as “The Diamond Problem”).

### 4.3.3 Use with the abstract attribute

The two previous examples of the implements derived type attribute, that were presented in Sects. 4.3.1 and 4.3.2, can, in fact, be coded equally well using implements statements. Yet, the implements attribute provides one capability that the implements statement lacks, namely implementation of traits even if the implementing derived type is abstract.

It must be noted, in this context, that the implements attribute provides a new, far more general, facility for requiring derived types to provide implementations of abstract methods, than the

present language's *flawed* (non-generalizable) mechanism of declaring methods with the deferred attribute. The latter feature's way of prescribing abstract method signatures differs from the traits based implements mechanisms that are proposed here, in that it generally requires these signatures to contain entirely *superfluous* passed-object dummy arguments, which unnecessarily and wrongly *intermix* the receiver of a message (i.e. the passed-object), with the interface of the message itself.

This is a fundamental design flaw in the present language that requires correction, and couldn't be simply carried over to the design of the traits based facilities, because it would have compromised all the new features that are based upon traits. In order to avoid having (in the long run) two subtly different ways of expressing abstract method signatures in the language, we propose to first make the specification of passed-object dummy arguments, in abstract interfaces of deferred methods, optional rather than mandatory, and to subsequently declare obsolescent such mandatory specification. Alternatively, Fortran's deferred type-bound procedure declarations could be retired in their entirety, in favor of using traits, as it is demonstrated in the following example:

```

1  module abstract_new
2
3      abstract interface :: IDeferred
4          subroutine method2(res)
5              real, intent(out) :: res
6          end subroutine method2
7      end interface IDeferred
8
9      type, abstract, implements(IDeferred) :: Parent
10     contains
11         procedure, pass :: method1 => implementation1
12     end type Parent
13
14     type, extends(Parent) :: Child
15     contains
16         procedure, pass :: method2 => implementation2
17     end type Child
18
19 contains
20
21     subroutine implementation1(self,n)
22         class(Parent), intent(in) :: self
23         integer,          intent(in) :: n
24     end subroutine implementation1
25
26     subroutine implementation2(self,res)
27         class(Child), intent(in)  :: self
28         real,          intent(out) :: res
29     end subroutine implementation2
30
31 end module abstract_new

```

Here, the abstract type Parent provides an implementation of a method1, but it also obliges any type that derives from it to implement the abstract method2. In present Fortran, this would be expressed through a deferred type-bound procedure declaration for method2 within Parent, along with the provision of an unnamed abstract interface. Whereas in the above example this is formulated in the alternative way of using a trait: Parent is made to implement the IDeferred trait, that contains the abstract signature of method2. But since Parent is declared, at the same time, with the abstract attribute, it can pass this obligation to its children. Hence, the Child type, that extends(Parent), is now expected to provide an actual implementation of this method.

More generally, a type that is declared abstract, is allowed to provide partial or no imple-

mentations of the traits that it, or its ancestors, adopt. Any non-abstract type that extends such an abstract type through type extension must, however, provide full implementations. This is equivalent to how interfaces work in conjunction with abstract classes in the Java language [20].

## 4.4 Trait objects (run-time polymorphic variables)

The main purpose of named abstract interfaces/traits is to allow the programmer to declare (and thereby constrain) polymorphic variables in terms of them. This can happen either directly, by traits functioning as types in run-time polymorphism (as it is demonstrated in this section), or indirectly, by them acting as constraints on generic type parameters in compile-time polymorphism (see, e.g., Sect. 5.1.1).

### 4.4.1 Class specifier using a single trait

In order to use traits to declare run-time polymorphic objects through subtyping, Fortran’s class specifier for polymorphic variable declarations needs to be enhanced to accept traits; like in the declarations of the following two variables (that make use of the `IPrintable` trait of Sect. 4.1.1)

```
1 class(IPrintable), allocatable :: printer
2 class(IPrintable), pointer      :: printerptr
```

or the following declaration of a procedure argument:

```
1 class(IPrintable), intent(in) :: printer
```

The semantics here are that whenever a trait appears within the class specifier of an object’s declaration, then all the public methods of that object whose signatures are prescribed by the adopted trait (like `IPrintable` in the above examples), will make use of dynamic binding. That is, their calls will be resolved by the run-time system of the language (e.g. through a virtual method table). See also Sect. 5.1.5 for further information on the class specifier when used with traits.

In accordance with how objects that make use of run-time polymorphism through subclassing (i.e. implementation inheritance) are declared in the present Fortran standard, also “trait objects” (like `printer`, and `printerptr` in the examples above) must either be declared using the `allocatable`, or the `pointer` attribute, or they must be arguments of a procedure. The proposed extensions are therefore backwards compatible with the functionality that is already available in the present language.

### 4.4.2 Class specifier using trait combinations

The `IAdmissible` trait of Sect. 4.1.3 derives from the `ICastable` and `IPrintable` traits of Sect. 4.1.1. An object that requires the combined functionality of both these latter traits could thus be declared in terms of the `IAdmissible` trait as follows:

```
1 class(IAdmissible), allocatable :: obj
```

It happens often, though, that one needs to declare objects that conform to multiple traits, but where one would like to avoid having to code some (otherwise unneeded) intermediary trait, that derives from these. In such cases, it should be possible to express an object’s declaration more directly – by providing a traits combination to the class specifier. As in the following example

```
1 class(ICastable + IPrintable), allocatable :: obj
```

in which `obj` again conforms to both the `ICastable` and `IPrintable` traits.

## Chapter 5

# Fortran extensions II: Generics

The new subtyping features that were discussed in the previous chapter are required to uniformly express and support both run-time and compile-time polymorphism in Fortran. We will now proceed with describing further enhancements that are specifically needed in order to support compile-time polymorphism, i.e. generics.

### 5.1 Enhancements to traits

#### 5.1.1 Traits with generic procedure signatures

Abstract interface definitions must be allowed to contain the signatures of generically parameterized procedures, as it is the case in the Swift language. Parameterizing the abstract interfaces themselves (as in Go or Rust) is not as attractive from a user’s perspective (cf. Sect. 3.3). As an example, the following code shows a trait that is called `ISum`, which contains a signature that is intended for a generic type-bound function (i.e. generic method) with name `sum`:

```
1  abstract interface :: ISum
2      function sum{INumeric :: T}(x) result(s)
3          type(T), intent(in) :: x(:)
4          type(T)                :: s
5      end function sum
6  end interface ISum
```

The example illustrates the use of a generic type parameter, that is simply called `T` here, in terms of which the regular function arguments are declared. A significant difference of generic type parameters, as compared to regular function arguments, is that the former will be substituted by actual *type* arguments at compile time, in a process that is called instantiation.

A similarity is that, in the same way that regular function arguments need to be constrained by a provided type, type parameters need to be constrained by a provided meta-type. This (meta-type) constraint must either be the name of a single trait (like `INumeric` in the present example), or a traits combination that is expressed in terms of the `+` operator<sup>1</sup>. If the constraint is omitted by the programmer, then the compiler will infer the corresponding type parameter to be constrained by the `IAnyType` (empty) trait of Sect. 5.1.2. The proposed Fortran generics thus support “strong concepts”, and can be fully type-checked by a compiler.

Both, the type parameter and the constraint which precedes it, and is separated from it by a double colon, are part of a generic type parameters list that is enclosed in curly braces, and

---

<sup>1</sup>For usage examples of traits combinations in run-time and compile-time polymorphism, see Sects. 4.4.2 and 6.3, respectively.

follows immediately behind the procedure's name. Notice that the syntax that is used here, which deviates slightly from how Fortran's regular function arguments are declared, appears justified, as it reflects that, despite some similarities, in type parameters one is dealing with different entities.

### 5.1.2 Traits of type sets

To make Fortran's traits based generics facility easy to use, the language should allow for generics constraints to be expressed in an *implicit* fashion, as abstract interfaces/traits that are formulated in terms of type sets, as it is possible in the Go language. This is particularly important for admitting the use of selected sets of intrinsic types in generic code (and the multitude of intrinsic procedures that Fortran supports for these types). It is also essential for supporting concise generic procedural and functional programming, where one might wish to avoid having to explicitly implement interfaces and methods for types (which is, arguably, an OO technique).

#### Unions of types

The following example shows the simplest form of such a type set. It defines an `INumeric` trait, for use as a generics constraint in the example of Sect. 5.1.1, in order to admit for the type parameter, `T`, that was given there, only the (32 bits wide) default integer intrinsic data type:

```
1  abstract interface :: INumeric
2      integer
3  end interface INumeric
```

The above example is actually a special case of specifying entire *unions* of member types as a type set. A type set consisting of such a union of types is demonstrated in the following example

```
1  abstract interface :: INumeric
2      integer | real(real64)
3  end interface INumeric
```

that redefines the `INumeric` trait such as to admit either the default 32 bit integer, or the 64 bit real type as a generics constraint.

The semantics of such a type set construct are that it implicitly defines a *set of procedure signatures*, namely the signatures of the intersecting (common) set of all the operations and intrinsic procedures (also called methods in the following) that work with all the member types of the type set. This can also be restated, by saying that a type `T` *implements* a trait consisting of such a type set, if (and only if) it is a member type of this set. For instance, the `integer` and `real(real64)` types implement the `INumeric` trait (as it is given above) because they are member types of its type set. In contrast, Fortran's various complex types do not implement this trait, because they do not belong to its set of member types. The validity of the latter statement can be easily checked, by considering that the complex types do not support, i.e. implement, the relational operators (`<`) and (`>`). Implementations of the latter are required for conformance to this trait, because these operators *are* implemented by both the `integer` and `real(real64)` member types.

#### Assumed kind parameters

Expanding on the previous example, an `INumeric` trait that might be even more useful as a generics constraint, for a number of numeric tasks, could be coded as follows:

```
1  abstract interface :: INumeric
2      integer(*) | real(*) | complex(*)
3  end interface INumeric
```

Notice how this makes use of both unions of types, and assumed (i.e. wildcard) kind parameters for types, to include *all* integer, real, and complex types, that are admitted by the language, in a single abstract interface constraint.

The use of assumed kind parameters is here merely syntactic sugar that allows one to avoid having to write out a type set for all the possible kinds of a type. For instance, if the particular Fortran implementation supports `real(real32)` and `real(real64)` as its only real types, then `real(kind=*)`, or `real(*)` for short, is understood to mean the type set “`real(real32) | real(real64)`”. Notice, also, that the more types are added to a trait in this fashion, the smaller the set of intersecting methods will usually become.

## Empty trait

In the limit of adding all possible types to a type set, there won’t be any common methods left that are implemented by all its types. This results in the important case of the empty trait, that matches all types (since any type has at least zero methods):

```
1  abstract interface :: IAnyType
2  end interface IAnyType
```

## Inline notation

For simple use cases, it should be optionally possible for the programmer to employ a shorter notation for declaring type constraints for generics, than having to separately define a trait of type sets, like `INumeric` above, and to then use it as in Sect. 5.1.1. The following modification of trait `ISum`’s original declaration of Sect. 5.1.1, provides such an example:

```
1  abstract interface :: ISum
2      function sum{integer | real(real64) :: T}(x) result(s)
3          type(T), intent(in) :: x(:)
4          type(T)                :: s
5  end function sum
6  end interface ISum
```

The inline use of a type set, within the type parameters list in curly braces, is an alternative notation for defining a type set trait for use as a generics constraint for type `T`. In this particular case, for admitting only the 32 bit integer, or 64 bit real type for `T`, as it was discussed above.

### 5.1.3 Traits with associated (deferred) types

In addition to supporting implicitly expressed generics constraints through type sets, the language must allow such constraints to be formulated also through *explicit* specification of the signatures of ordinary type-bound procedures, initializers, or operators in traits. The proper support of operator signatures, but also of other generic functionality, requires the use of so-called associated types within traits. Associated types (which are also available in Rust and Swift) are essentially aliases. They are employed within traits, in lieu of types whose actual value is not known at the time a trait is formulated, but will be known by a compiler once the programmer has actually implemented that trait for some (concrete derived or intrinsic) data type.

A typical use case for associated types is when a trait needs to refer to the implementing type itself. This use case is, in fact, so frequent that the aforementioned languages provide a predefined associated type, `Self`, for it. In Fortran (which is case-insensitive), we propose to call this type `itself`. The following version of an `INumeric` trait makes use of this predefined type, in order to

explicitly declare the signatures of an operator(+) for addition, and an operator(/) for division, that together with the signature of an explicitly declared initializer, init, are intended for use as a generics constraint within the ISum trait of Sect. 5.1.1:

```

1  abstract interface :: INumeric
2      function init(n)
3          integer, intent(in) :: n
4      end function init
5      function operator(+)(lhs,rhs) result(res)
6          type(itself), intent(in) :: lhs, rhs
7          type(itself)           :: res
8      end function operator(+)
9      function operator(/)(lhs,rhs) result(res)
10         type(itself), intent(in) :: lhs, rhs
11         type(itself)           :: res
12     end function operator(/)
13 end interface INumeric

```

For a more comprehensive usage example, see Sect. 6.2.3, where this trait is employed in the framework of our standard test problem, to strictly conform to OOP's Open/Closed Principle.

In general, associated types will have to be declared by the programmer himself, using deferred type-definition statements, as it is shown in the following example of the trait IAppendable:

```

1  abstract interface :: IAppendable
2      typedef, deferred :: Element
3      subroutine append(item)
4          type(Element), intent(in) :: item
5      end subroutine append
6  end interface IAppendable

```

This trait requires any implementing type to provide a subroutine for appending certain items (to the type). The type of these items is unknown upon trait formulation, and is therefore declared using the deferred (i.e. placeholder) type Element, that is left for inference by the compiler. A worked out example that illustrates the actual use of this trait is provided in Sect. 6.4.

#### 5.1.4 Predefined traits

The facilities that were described in the previous sections are flexible enough to enable the user to easily construct generics constraints himself, the way he needs them. Nevertheless, the language should ideally also supply a collection of already predefined, frequently used, generics constraints, in the form of traits that are contained in a language-intrinsic module (that we'll tentatively call generics here). A list of such predefined traits could include

- the empty trait IAnyType that admits any type (see Sect. 5.1.2),
- some predefined numeric traits allowing for different numeric operations, but also
- some predefined traits to allow for the use of relational operators with different types.

Such traits could then be imported from user code through a use statement, like in the following example, that assumes the existence of a language defined trait INumeric:

```

1  module user_code
2
3      use, intrinsic :: generics, only: INumeric
4
5      abstract interface :: ISum

```

```

6      function sum{INumeric :: T}(x) result(s)
7          type(T), intent(in) :: x(:)
8          type(T)                :: s
9      end function sum
10 end interface ISum
11
12 end module user_code

```

### 5.1.5 Present restrictions

The previous sections have shown examples of four kinds of traits, namely

- traits of type sets (or “TOTs”), like the `INumeric` traits of Sect. 5.1.2,
- traits with associated types (or “TATs”), like the `INumeric` and `IAppendable` traits of Sect. 5.1.3,
- “universal” traits, i.e. traits that do not fall into the previous two categories (examples of this are all the traits of Chapter 4, and the `ISum` trait of Sect. 5.1.1),
- and predefined traits, that can be any of the former three.

“Universal” traits have the least amount of restrictions connected to them. They can be used together with the `class` specifier in order to declare trait objects (cf. Sect. 4.4), i.e. run-time polymorphic variables, or instances of what in type theory are known as existential types. But they can *also* serve as generics constraints (as it will be shown in Sect. 5.4.3). These traits allow the programmer to switch from dynamic, to static dispatch of methods (see the later Sect. 6.2.2), and vice versa (hence their name). In contrast, both TOTs and TATs are presently (i.e. in the framework of a prototype implementation of this proposal) *exclusively* intended for use as generics constraints. Hence, a compiler must ensure that these two trait kinds are *not* used for any other purpose.

In particular, both TOTs and TATs cannot be used in variable declarations that make use of the `class` specifier, while TOTs cannot be implemented (in the sense of Sect. 4.2.3) by derived (i.e. user-defined) types<sup>2</sup>. In a future language revision, TOTs could, however, be admitted for use with the type declaration specifier to support compile-time polymorphism through union types [17, 19], as it is further elaborated on in Sect. 7.3. On the other hand, TATs could be made “universal”, and admitted for use with both the `class` and type specifiers, as in Swift 6.0, which makes use of the analogous keywords `any` and `some`, respectively, for the same purpose.

## 5.2 Casts to generic types

A language like Fortran, that is intended for numeric use, where conversions between different types are required rather frequently, must allow one to carry out type conversions also in generic code, in a manner that is similarly user-friendly as in the Swift and Go languages, without having to rely on external library functionality. Moreover, this should be possible *regardless* of whether the type to be cast to is an intrinsic or user-defined type. This requires language-built-in mechanisms (that were already discussed) for requiring the availability of initializer implementations for types, in a manner that is either implicit (see the Sect. 5.1.2 on type sets), or explicit (see the language features that were described in Sects. 4.1, and 4.2).

---

<sup>2</sup>One of the problems here is that any new intrinsic function for an intrinsic type, that would need to be added to the language, would change the set of methods of all the type sets of which this type is a member. This would break any user-defined types that would implement traits which are based on these type sets.

As an example, consider the test problem implementation of Sect. 2.2. There, a reduction variable for summation, *s*, needs to be initialized to the zero constant of a generic type, *T*. It should be possible, in Fortran, to effect this initialization by simply writing an expression like the following:

```
1 s = T(0)
```

Which (in this example) implies a cast of the zero constant of the integer type into the corresponding zero constant of type *T*, by a call to an appropriate initializer of this latter generic type.

Notice that if the generic type parameter, *T*, is instantiated at compile time with a language-intrinsic type, then the “initializer” to be called by the compiler would actually be Fortran’s built-in function for casts to this type. That is, if *T* is instantiated with, say, the `real(real64)` type, then the compiler would transform `s = T(0)` into the following, final, expression:

```
1 s = real(0, kind=real64)
```

If *T* is instantiated, instead, with a derived type, then an appropriate initializer of this type would need to be invoked. This could either be the type’s default “structure constructor”, or a user-provided initializer, either of which would need to take in an integer, and return an instance of type *T*. In the absence of an appropriate (language or user-provided) initializer implementation, the compiler should simply emit an error message, and abort compilation.

### 5.3 Type inference

As it will be shown in detail in the next sections, the present design requires a Fortran compiler to implement type inference, in order to make it easy for a programmer to use generic code. In such a framework, it is then only natural to support also combined automatic declaration and initialization of variables, as it is possible in the languages that were surveyed in Chapter 3. We propose to adopt Go’s `:=` assignment operator for this purpose, as it is syntactically light-weight, and could thus be easily employed not just for declaring local variables in a non-intrusive fashion within sections of regular executable statements, but also within loops (including the declaration of their indices), e.g. as follows:

```
1 real :: arr(10)
2 do i := 1, 10
3   tmp := real(i)
4   arr(i) = tmp
5 end do
```

This last (somewhat concocted) example could also be written as a single line of code, by embedding an implied do loop within an array constructor, and using the `:=` operator twice:

```
1 arr := [(real(i), i:=1,10)]
```

The `:=` operator is strictly intended for the declaration of *local* variables in procedures and main programs. Module variables, and variables that express function return values, aren’t considered local (in Fortran) and are therefore not allowed to be declared in this way.

### 5.4 Generic entities

The present Swift-inspired generics allow both ordinary and type-bound procedures (i.e. methods), as well as derived types and their constructors, to be given their own generic type parameters. Main programs *cannot* be generically parameterized, as they cannot be called by any other program unit. Generically parameterized modules are considered an equally moot concept, given that (in contrast to derived types) modules are not a part of the type system.

### 5.4.1 Generic procedures

Using the syntax that is proposed in this document, an implementation of a Fortran function for array summation that is generic over the type of its input array argument would look as follows<sup>3</sup>:

```
1  function sum{INumeric :: T}(x) result(s)
2      type(T), intent(in) :: x(:)
3      type(T)                :: s
4      s = T(0)
5      do i := 1, size(x)
6          s = s + x(i)
7      end do
8  end function sum
```

The following examples will assume that the INumeric generics constraint, that is used here for type T, is provided by a trait that prescribes either implicitly (see Sect. 5.1.2), or explicitly (see Sect. 5.1.3), an appropriate initializer and addition operator, that are implemented by the integer and real(real64) types.

To actually use the sum generic function, one simply needs to pass to it (via its regular arguments list) an argument of one of these two admitted “numeric” types, as in the following calls:

```
1  integer_total := sum([1,2,3,4,5])
2  double_total  := sum([1.d0,2.d0,3.d0,4.d0,5.d0])
```

The compiler will then use automatic type inference to instantiate appropriate versions of the sum generic function for each call, given the types of the provided regular arguments. The programmer can thus straightforwardly use his generic routine with any of the admitted types. In this particular example, type inference will actually be employed not just for calling sum, but also for instantiating the variables integer\_total and double\_total via the := operator.

Although the present design *doesn't* require it, the programmer can specify generic type arguments also manually. The following example demonstrates how both the aforegiven result variables can be declared, and generic type arguments to the sum function be provided, by hand. Notice the use of curly braces for the latter purpose:

```
1  integer      :: integer_total
2  real(real64) :: double_total
3
4  integer_total = sum{integer}([1,2,3,4,5])
5  double_total  = sum{real(real64)}([1.d0,2.d0,3.d0,4.d0,5.d0])
```

Manual instantiation of generic procedures is only needed in the relatively rare cases where a regular arguments list is either unavailable or outright inappropriate, e.g. for use in procedure pointer assignments, or in associate statements, as in the next example:

```
1  real(real64) :: dtot(2)
2  real(real32) :: stot(2)
3  procedure(sum_real64), pointer :: dsum
4
5  abstract interface
6      function sum_real64(x) result(s)
7          real(real64), intent(in) :: x(:)
8          real(real64)                :: s
9      end function sum_real64
10 end interface
11
12 dsum => sum{real(real64)}
```

---

<sup>3</sup>Generic subroutines can be coded completely analogously.

```

13 dtot(1) = dsum([1.d0,2.d0,3.d0,4.d0,5.d0])
14 dtot(2) = dsum([2.d0,4.d0,6.d0,8.d0])
15
16 associate( ssum => sum{real(real32)} )
17     stot(1) = ssum([1.,2.,3.,4.,5.])
18     stot(2) = ssum([2.,4.,6.,8.])
19 end associate

```

## 5.4.2 Generic methods

The same summation algorithm as that of Sect. 5.4.1, when implemented as a generic method, `sum`, that is bound to a derived type named `SimpleSum`, which implements the trait `ISum` as it was given in Sect. 5.1.1, would instead look as follows:

```

1 module simple_library
2     ...
3     type, public :: SimpleSum
4 end type SimpleSum
5
6 implements ISum :: SimpleSum
7     procedure, nopass :: sum
8 end implements SimpleSum
9
10 contains
11
12 function sum{INumeric :: T}(x) result(s)
13     type(T), intent(in) :: x(:)
14     type(T)                :: s
15     s = T(0)
16     do i := 1, size(x)
17         s = s + x(i)
18     end do
19 end function sum
20
21 end module simple_library

```

We have used here the `implements` statement for implementing a trait by a derived type, and will continue to consistently do so for the remainder of this chapter. The alternative way of employing the `implements` derived type attribute for the same purpose, is demonstrated on the same examples in Chapter 6. Generic methods are used completely analogously to generic procedures. The automatic instantiation use case of Sect. 5.4.1 would, for instance, be coded with the help of `SimpleSum`'s (default, structure) constructor as follows:

```

1 simple := SimpleSum()
2
3 integer_total := simple%sum([1,2,3,4,5])
4 double_total := simple%sum([1.d0,2.d0,3.d0,4.d0,5.d0])

```

Whereas Sect. 5.4.1's "manual" use case would take the form:

```

1 type(SimpleSum) :: simple
2
3 integer        :: integer_total
4 real(real64)   :: double_total
5
6 integer_total = simple%sum{integer}([1,2,3,4,5])
7 double_total  = simple%sum{real(real64)}([1.d0,2.d0,3.d0,4.d0,5.d0])

```

See also Sects. 5.4.3 and 5.4.4 for further instantiation examples.

### 5.4.3 Generic derived types

Generic type parameters lists are also allowed to appear in derived type definitions, as in the following example, in which the `ISum` trait of Sect. 5.1.1 is implemented by another derived-type, that is named `PairwiseSum`:

```
1 module pairwise_library
2   ...
3   type, public :: PairwiseSum{ISum :: U}
4     private
5       type(U) :: other
6   end type PairwiseSum
7
8   implements ISum :: PairwiseSum{ISum :: U}
9     initial :: init
10    procedure, pass :: sum
11  end implements PairwiseSum
12
13 contains
14
15  function init(other) result(res)
16    type(U), intent(in) :: other
17    type(PairwiseSum{U}) :: res
18    res%other = other
19  end function init
20
21  function sum{INumeric :: T}(self,x) result(s)
22    class(PairwiseSum{U}), intent(in) :: self
23    type(T),                intent(in) :: x(:)
24    type(T)                  :: s
25    if (size(x) <= 2) then
26      s = self%other%sum(x)
27    else
28      m := size(x) / 2
29      s = self%sum(x(:m)) + self%sum(x(m+1:))
30    end if
31  end function sum
32
33 end module pairwise_library
```

Notice, how type `PairwiseSum` depends on a generic type parameter, `U`, that is used within `PairwiseSum` in order to declare a field variable of type `(U)`, which is named `other`. As it is indicated by the type constraint on `U`, object `other` conforms to the `ISum` trait itself, and therefore contains its own implementation of the `sum` procedure.

The above example furthermore shows, that a derived type's generic parameters are accessible within the scopes of its type-bound initializers and procedures. For instance, type `PairwiseSum`'s method, `sum`, can access `PairwiseSum`'s generic parameter `U`, e.g. in order to declare its passed-object dummy argument, `self`, to be of `class(PairwiseSum{U})`. This allows method `sum` to make use of two independently defined generic type parameters, `T` and `U`, which grants it increased flexibility.

The declaration `class(PairwiseSum{U})` does not imply any ambiguities or contradictions with respect to compile-time vs. run-time polymorphism, because substitution semantics apply. At compile time, the compiler will substitute a set of different type arguments for the generic parameter `U`. Hence, the notation `PairwiseSum{U}` really refers to a set of multiple, related, but *different* `PairwiseSum` types (whose only commonality is that they all implement the `ISum` trait, and furthermore contain different field components that do the same). Of course, passed-object dummy arguments of any of the different `PairwiseSum` types of this set can then be run-time polymorphic.

In exactly the same manner that passed-object dummy arguments of other derived types that implement the same trait can be run-time polymorphic.

The following code snippet finally shows how an object of the `PairwiseSum` type could be declared and instantiated manually (by substituting its generic type parameter `U` by the `SimpleSum` type of Sect. 5.4.2), and how its generic `sum` method would be employed using automatic type inference:

```
1  type(PairwiseSum{SimpleSum}) :: pairwise
2
3  integer_total := pairwise%sum([1,2,3,4,5])
4  double_total  := pairwise%sum([1.d0,2.d0,3.d0,4.d0,5.d0])
```

#### 5.4.4 Generic constructors

If a derived type is parameterized over a generic type, as in the `PairwiseSum` type example of Sect. 5.4.3, then any constructors of this derived type must also be assumed to be implicitly parameterized over the same generic type. Hence, calls of these constructors, with specific argument types being provided in curly braces to substitute the generic type parameters of their derived types, must be valid for consistency reasons. This must actually hold for both structure *and* user-defined constructors/initializers – even if user-provided implementations of these initializers do *not* contain explicit, corresponding, generic parameters lists (see the example provided in Sect. 5.4.3). These lists are then understood to be implicitly provided by the compiler (compare the example calls below, and see also Sect. 6.4).

For an illustration of this, consider first the following variation of Sect. 5.4.3's last code snippet:

```
1  pairwise := PairwiseSum(SimpleSum())
2
3  integer_total := pairwise%sum([1,2,3,4,5])
4  double_total  := pairwise%sum([1.d0,2.d0,3.d0,4.d0,5.d0])
```

Notice, here, the lack of an explicit type declaration for object `pairwise`, the use of the `:=` operator, and the call to type `PairwiseSum`'s (user-defined) constructor on the right-hand-side of the first such assignment. The compiler will take advantage of this latter call to automatically instantiate object `pairwise` with the `PairwiseSum` type through type inference.

The above code actually makes use of constructor calls of all the types that are involved in the present example, namely `PairwiseSum` *and* `SimpleSum`, in the standard OO way of chaining such calls. The compiler will then use constructor `PairwiseSum`'s regular argument, `SimpleSum()`, in order to instantiate the generic field object `other`, that is embedded within the `PairwiseSum` type (cf. Sect. 5.4.3). By doing so, the compiler will also infer the generic type parameter of the `PairwiseSum` type, in complete analogy to how type inference works when calling a generic procedure – thus making manual instantiation of any of the involved generic types unnecessary.

However, it is also possible, within the last call, to provide `PairwiseSum`'s generic type argument manually in curly braces, for the mentioned consistency reasons, as follows:

```
1  pairwise := PairwiseSum{SimpleSum}(SimpleSum())
```

Hence, both of these constructor call variants are valid, and give the same result. Although the first variant is the one that we actually recommend for general use (possibly together with the employment of keyword arguments).

## Chapter 6

# Fortran examples

To comprehensively illustrate how the new features, that were introduced in the last two chapters, would be used in practice, we will give in the present chapter several worked out examples. The first two subsections contain both functional and OO Fortran code versions of the standard test problem that is used throughout this document. While the last two subsections deal with how to use operator overloading, and associated types with generics.

### 6.1 Functional versions of the standard test problem

Fortran presently lacks support for advanced functional programming capabilities, like closures and variables of higher-order functions, that are, e.g., available in Go and other modern languages. In contrast to the Go version of the standard test problem that is given in Sect. 2.2, the functional Fortran code versions that are presented in this section therefore make no attempt to eliminate rigid dependencies on user-defined function implementations, and content themselves with demonstrating how the new generics features can be used to eliminate rigid dependencies on language-intrinsic types.

#### 6.1.1 Automatic instantiation of generic procedures

Listing 6.1 shows a straightforward generic functional implementation of the standard test problem, that uses automatic type inference by the compiler to instantiate its generic functions. The following additional remarks should help to better understand this code:

- To express type genericity for the arguments and return values of all generic functions, a type constraint is employed that is implemented as the type set `integer | real(real64)` within a trait `INumeric`.
- Trait `INumeric` is defined by the user, obviating thereby the need for an external dependency.
- Any required conversions to generic types are done using explicit casts, as in Go (or Swift).
- *All* the required instantiations of generic procedures are performed automatically by the compiler, based on type inference of the latters' passed regular arguments.
- Use of type inference is also made in order to simultaneously declare and initialize local variables with the `:=` assignment operator, even within executable statement sections and loops.

Listing 6.1: Fortran version of the array averaging problem with automatic generics instantiation.

```

1  module averager_library
2
3      use, intrinsic :: iso_fortran_env, only: real64
4
5      implicit none
6      private
7
8      public :: simple_average, pairwise_average
9
10     abstract interface :: INumeric
11         integer | real(real64)
12     end interface INumeric
13
14     contains
15
16     function simple_sum{INumeric :: T}(x) result(s)
17         type(T), intent(in) :: x(:)
18         type(T)                :: s
19         s = T(0)
20         do i := 1, size(x)
21             s = s + x(i)
22         end do
23     end function simple_sum
24
25     function pairwise_sum{INumeric :: T}(x) result(s)
26         type(T), intent(in) :: x(:)
27         type(T)                :: s
28         if (size(x) <= 2) then
29             s = simple_sum(x)
30         else
31             m := size(x) / 2
32             s = pairwise_sum(x(:m)) + pairwise_sum(x(m+1:))
33         end if
34     end function pairwise_sum
35
36     function simple_average{INumeric :: T}(x) result(a)
37         type(T), intent(in) :: x(:)
38         type(T)                :: a
39         a = simple_sum(x) / T(size(x))
40     end function simple_average
41
42     function pairwise_average{INumeric :: T}(x) result(a)
43         type(T), intent(in) :: x(:)
44         type(T)                :: a
45         a = pairwise_sum(x) / T(size(x))
46     end function pairwise_average
47
48 end module averager_library
49
50 program main
51
52     ! dependencies on intrinsic constants
53     use, intrinsic :: iso_fortran_env, only: real64
54
55     ! dependencies on implementations
56     use averager_library, only: simple_average, pairwise_average
57

```

```

58  implicit none
59
60  ! declarations
61  integer :: key
62
63  write(*,'(a)') 'Simple_sum_average:_1'
64  write(*,'(a)') 'Pairwise_sum_average:_2'
65  write(*,'(a)',advance='no') 'Choose_an_averaging_method:_ '
66  read(*,*) key
67
68  xi := [1, 2, 3, 4, 5]
69  xf := [1.d0, 2.d0, 3.d0, 4.d0, 5.d0]
70
71  select case (key)
72  case (1)
73      print '(i1)', simple_average(xi)
74      print '(f3.1)', simple_average(xf)
75  case (2)
76      print '(i1)', pairwise_average(xi)
77      print '(f3.1)', pairwise_average(xf)
78  case default
79      stop 'Case_not_implemented!'
80  end select
81
82  end program main

```

The example demonstrates that using the new generics features together with a functional (or procedural) programming style is easy, that the syntax is concise, and that type inference by the compiler should be straightforward and therefore reliable. Hence, we believe that the generics features, that are described here, will place no burden on the programmer.

### 6.1.2 Manual instantiation of generic procedures

It is actually possible to make the Fortran code version that was given in Listing 6.1, resemble the Go code version of Listing 2.2 a bit closer, by having two procedure pointers stand in, within the `select case` statement of the main program, for the closures that were used in the Go code. As this is a good example for demonstrating how generics can be instantiated manually by the programmer, we give in Listing 6.2 an alternative form of the main program of Listing 6.1 that makes both use of procedure pointers, and such manual instantiation (as it was discussed in Sect. 5.4.1).

Listing 6.2: Main program using procedure pointers and manual generics instantiation.

```

1  program main
2
3      ! dependencies on intrinsic constants
4      use, intrinsic :: iso_fortran_env, only: real64
5
6      ! dependencies on implementations
7      use averager_library, only: simple_average, pairwise_average
8
9      implicit none
10
11     ! declarations
12     integer :: key
13     procedure(average_integer), pointer :: avi
14     procedure(average_real64), pointer :: avf
15

```

```

16  abstract interface
17      function average_integer(x) result(a)
18          integer, intent(in) :: x(:)
19          integer              :: a
20      end function average_integer
21      function average_real64(x) result(a)
22          real(real64), intent(in) :: x(:)
23          real(real64)              :: a
24      end function average_real64
25  end interface
26
27  write(*, '(a)') 'Simple_sum_average:_1'
28  write(*, '(a)') 'Pairwise_sum_average:_2'
29  write(*, '(a)', advance='no') 'Choose_an_averaging_method:_ '
30  read(*, *) key
31
32  select case (key)
33  case (1)
34      avi => simple_average{integer}
35      avf => simple_average{real(real64)}
36  case (2)
37      avi => pairwise_average{integer}
38      avf => pairwise_average{real(real64)}
39  case default
40      stop 'Case_not_implemented!'
41  end select
42
43  xi := [1, 2, 3, 4, 5]
44  xf := [1.d0, 2.d0, 3.d0, 4.d0, 5.d0]
45
46  print '(i1)', avi(xi)
47  print '(f3.1)', avf(xf)
48
49  end program main

```

## 6.2 Object-oriented versions of the standard test problem

The present section will demonstrate that being able to use the new generics features seamlessly and easily, even within a modern-day OO programming setting, is one of the great strengths of the present design.

### 6.2.1 Dynamic method dispatch

Listing 6.3 shows an encapsulated (i.e. OO) Fortran code version of the standard test problem, that makes use of dynamic method dispatch for all user-defined abstract data types (ADTs), and is therefore comparable to the code versions that were presented in Chapter 3 for all the other languages.

- As in all these other versions, three interfaces/traits are used to manage all the source code dependencies in the problem: INumeric, ISum, and IAverager. Trait INumeric is defined here by the user himself as a type set, similar to the corresponding Go code (for a version that does not use type sets, see the later Sect. 6.2.3).

- In contrast to the Go and Rust versions (Listings 3.1 and 3.2), none of the aforementioned interfaces/traits is parameterized itself, since we followed Swift’s basic model of generics.
- Interface inheritance is expressed through the presence of the `implements` attribute in derived-type definitions (equivalent to the Swift version, Listing 3.4). Alternatively, `implements` statements could be used (cf. Sects. 4.2, 5.4.2, and 5.4.3).
- All our derived types are sealed, which makes them inextensible to implementation inheritance, and thus enables us to declare the passed-object dummy arguments of their type-bound procedures with the type specifier.
- The example code employs both language- and user-defined constructors in its main program, in a manner that is fully equivalent to the Swift version, Listing 3.4.
- This Fortran version makes use of modules and `use` statements with `only` clauses, in order to make explicit the source code dependencies of the different user-defined ADTs/classes.

Listing 6.3: Fortran OO version of the array averaging example with dynamic method dispatch.

```

1  module interfaces
2
3      use, intrinsic :: iso_fortran_env, only: real64
4
5      implicit none
6      private
7
8      public :: INumeric, ISum, IAverager
9
10     abstract interface :: INumeric
11         integer | real(real64)
12     end interface INumeric
13
14     abstract interface :: ISum
15         function sum{INumeric :: T}(x) result(s)
16             type(T), intent(in) :: x(:)
17             type(T)                :: s
18         end function sum
19     end interface ISum
20
21     abstract interface :: IAverager
22         function average{INumeric :: T}(x) result(a)
23             type(T), intent(in) :: x(:)
24             type(T)                :: a
25         end function average
26     end interface IAverager
27
28 end module interfaces
29
30 module simple_library
31
32     use interfaces, only: ISum, INumeric
33
34     implicit none
35     private
36
37     public :: SimpleSum
38

```

```

39  type, sealed, implements(ISum) :: SimpleSum
40  contains
41      procedure, nopass :: sum
42  end type SimpleSum
43
44  contains
45
46  function sum{INumeric :: T}(x) result(s)
47      type(T), intent(in) :: x(:)
48      type(T)                :: s
49      s = T(0)
50      do i := 1, size(x)
51          s = s + x(i)
52      end do
53  end function sum
54
55  end module simple_library
56
57  module pairwise_library
58
59      use interfaces, only: ISum, INumeric
60
61      implicit none
62      private
63
64      public :: PairwiseSum
65
66      type, sealed, implements(ISum) :: PairwiseSum
67      private
68      class(ISum), allocatable :: other
69      contains
70          initial :: init
71          procedure, pass :: sum
72      end type PairwiseSum
73
74      contains
75
76      function init(other) result(res)
77          class(ISum), intent(in) :: other
78          type(PairwiseSum)        :: res
79          res%other = other
80      end function init
81
82      function sum{INumeric :: T}(self,x) result(s)
83          type(PairwiseSum), intent(in) :: self
84          type(T),                intent(in) :: x(:)
85          type(T)                :: s
86          if (size(x) <= 2) then
87              s = self%other%sum(x)
88          else
89              m := size(x) / 2
90              s = self%sum(x(:m)) + self%sum(x(m+1:))
91          end if
92      end function sum
93
94  end module pairwise_library
95
96  module averager_library
97

```

```

98  use interfaces, only: IAverager, ISum, INumeric
99
100 implicit none
101 private
102
103 public :: Averager
104
105 type, sealed, implements(IAverager) :: Averager
106     private
107     class(ISum), allocatable :: drv
108 contains
109     initial :: init
110     procedure, pass :: average
111 end type Averager
112
113 contains
114
115 function init(drv) result(res)
116     class(ISum), intent(in) :: drv
117     type(Averager)          :: res
118     res%drv = drv
119 end function init
120
121 function average{INumeric :: T}(self,x) result(a)
122     type(Averager), intent(in) :: self
123     type(T),          intent(in) :: x(:)
124     type(T)           :: a
125     a = self%drv%sum(x) / T(size(x))
126 end function average
127
128 end module averager_library
129
130 program main
131
132 ! dependencies on abstractions
133 use interfaces,      only: IAverager
134
135 ! dependencies on implementations
136 use simple_library,  only: SimpleSum
137 use pairwise_library, only: PairwiseSum
138 use averager_library, only: Averager
139
140 implicit none
141
142 ! declarations
143 class(IAverager), allocatable :: av
144 integer :: key
145
146 write(*,'(a)') 'Simple__sum_average:_1'
147 write(*,'(a)') 'Pairwise_sum_average:_2'
148 write(*,'(a)',advance='no') 'Choose_an_averaging_method:_ '
149 read(*,*) key
150
151 select case (key)
152 case (1)
153     av = Averager(drv = SimpleSum())
154 case (2)
155     av = Averager(drv = PairwiseSum(other = SimpleSum()))
156 case default

```

```

157     stop 'Case_not_implemented!'
158 end select
159
160 xi := [1, 2, 3, 4, 5]
161 xf := [1.d0, 2.d0, 3.d0, 4.d0, 5.d0]
162
163 print '(i1)', av%average(xi)
164 print '(f3.1)', av%average(xf)
165
166 end program main

```

The most important point to note, in Listing 6.3, is how this code makes use of trait objects (cf. Sect. 4.4) and generics to avoid rigid dependencies on user-defined implementation and language-intrinsic types, respectively, and to thus realize a plugin architecture that allows for a maximum of code reuse. In particular, the main program is the only part of the code that (necessarily) depends on implementations. The *entire* rest of the code is decoupled, that is it depends merely on traits (i.e. abstract interfaces, see the use statements in the above modules).

All the traits that are required for this purpose are defined by the user himself, without any need to depend on external libraries. The OO Fortran version, that is presented here, is therefore as clean with respect to dependency management, and as easy to use, read, and understand, as the Go and Swift implementations of Listings 3.1 and 3.4, respectively.

### 6.2.2 Static method dispatch

One of the greatest benefits of the present design is that, through the use of generics, polymorphic methods in OO programming can be made to use static dispatch<sup>1</sup>. This will enable inlining of polymorphic methods by Fortran compilers, to potentially improve code performance. Listing 6.4 gives minimally changed, alternative, implementations of the modules `pairwise_library` and `averager_library`, with which the original versions of these modules in Listing 6.3 would need to be replaced, to effect static dispatch of the different sum methods.

The required changes are confined to a parameterization of the `PairwiseSum` and `Averager` derived types, by generic type parameters, `U`, that conform to the (universal) `ISum` trait. These type parameters are then used to declare the field objects `other` and `drv` of the `PairwiseSum` and `Averager` types, respectively, by means of the type specifier. Taken together, these changes signify compile-time polymorphism for the objects `other` and `drv` to the compiler, and hence static dispatch of their methods (whose signatures were declared in the `ISum` trait). Contrast this with the `class` specifier, that was used previously for the declaration of these former trait objects, to effect run-time polymorphism and dynamic dispatch of their methods. See also Table 6.1 for a summary of rules regarding method dispatch.

Everything else, especially the declaration of these object variables as `allocatables` (and their instantiation using chaining of constructor calls in the main program), was kept the same in order to demonstrate that static method dispatch does *not* mean that the actual object instances that contain the methods must be initialized and their memory allocated at compile-time. (Although in this particular example this is possible, by simply deleting the `allocatable` attribute from their declarations.) Notice, that none of the source code dependencies in the use statements have changed. That is, the code is *still* fully decoupled, despite making now use of static dispatch!

<sup>1</sup>This can be achieved in a compiler by implementing generic polymorphism either through the technique of monomorphization, or through the use of “witness tables” for generics that are devirtualized at compile-time through specialization [5, 16]. In both cases, the result is static method dispatch.

Listing 6.4: Minimal required changes to effect static method dispatch of the sum methods.

```

1 module pairwise_library
2
3   use interfaces, only: ISum, INumeric
4
5   implicit none
6   private
7
8   public :: PairwiseSum
9
10  type, sealed, implements(ISum) :: PairwiseSum{ISum :: U}
11    private
12    type(U), allocatable :: other
13  contains
14    initial :: init
15    procedure, pass :: sum
16  end type PairwiseSum
17
18 contains
19
20  function init(other) result(res)
21    type(U), intent(in) :: other
22    type(PairwiseSum{U}) :: res
23    res%other = other
24  end function init
25
26  function sum{INumeric :: T}(self,x) result(s)
27    type(PairwiseSum{U}), intent(in) :: self
28    type(T), intent(in) :: x(:)
29    type(T) :: s
30    if (size(x) <= 2) then
31      s = self%other%sum(x)
32    else
33      m := size(x) / 2
34      s = self%sum(x(:m)) + self%sum(x(m+1:))
35    end if
36  end function sum
37
38 end module pairwise_library
39
40 module averager_library
41
42   use interfaces, only: IAverager, ISum, INumeric
43
44   implicit none
45   private
46
47   public :: Averager
48
49   type, sealed, implements(IAverager) :: Averager{ISum :: U}
50     private
51     type(U), allocatable :: drv
52  contains
53    initial :: init
54    procedure, pass :: average
55  end type Averager
56
57 contains

```

```

58
59  function init(drv) result(res)
60      type(U), intent(in) :: drv
61      type(Averager{U})    :: res
62      res%drv = drv
63  end function init
64
65  function average{INumeric :: T}(self,x) result(a)
66      type(Averager{U}), intent(in) :: self
67      type(T),           intent(in) :: x(:)
68      type(T)            :: a
69      a = self%drv%sum(x) / T(size(x))
70  end function average
71
72 end module averager_library

```

| object declaration | dynamic dispatch           | static dispatch                           |
|--------------------|----------------------------|-------------------------------------------|
| class(Trait)       | always                     | never                                     |
| class(DerivedType) | if DerivedType is extended | if DerivedType is unextended <sup>2</sup> |
| type(DerivedType)  | never                      | always                                    |
| type(Trait)        | —                          | —                                         |

Table 6.1: Correspondence between object declarations and method dispatch strategies that would be typically employed by an optimizing Fortran compiler according to the present document. A dash indicates that the case in question is presently undefined, but could be used for future extensions as mentioned in Sect. 5.1.5, and further discussed in Sect. 7.3.

The complete source code of this (statically dispatched) version can be found in the Code subdirectory that accompanies this document<sup>3</sup>. Notice, from this complete listing (or from the original reproduction of the main program in Listing 6.3), that the object `av` (of `IAverager` type) still needs to make use of run-time (as opposed to compile-time) polymorphism. Because instances of types that *differ* from `av`'s declared own type must remain assignable to this variable.

As in corresponding (static method dispatch) Swift and Rust implementations of the standard test problem, that are not reproduced here (see the Code directory for this), the instantiation (in the main program) of the objects `other` and `drv` through constructor calls, has the benefit that the compiler can infer the correct type arguments, that are required to automatically instantiate all the involved generic derived types (cf. Sect 5.4.4). Hence (and similar to both the functional and dynamically dispatched OO code versions of Listings 6.1 and 6.3, respectively), not a single manual instantiation of generics is necessary, anywhere, also in the statically dispatched version.

### 6.2.3 Strict conformance to the Open/Closed Principle

In the examples that were presented so far in this chapter, we made predominant use of generics constraints that, in the form of the `INumeric` trait, were expressed as (Go-like) type sets of intrinsic types. As it was already discussed in Sect. 3.1.1, the conciseness of formulation, that type sets permit, comes at the price of having to give up strict conformance to the Open/Closed Principle (OCP) of object-oriented programming.

<sup>2</sup>Or the method is declared as `non_overridable`.

<sup>3</sup>See the file `Code/Fortran/static.f90`.

However, the present design offers the programmer the possibility to formulate generics constraints also in the alternative manner of always specifying both the signatures of procedures in abstract interfaces, and their implementations, explicitly. An example of this was already given in the form of the (universal) `ISum` trait, that was used as a generics constraint within the derived types `PairwiseSum` and `Averager` of Listing 6.4. The explicit approach has the advantage of allowing generic object-oriented code to strictly conform to the OCP. This can be accomplished independently of the employed (dynamic or static) method dispatch strategy. Even though we will confine ourselves here to illustrating this using only Listing 6.3, i.e. the dynamically dispatched code version of the standard test problem, as a baseline.

The following Listing 6.5 gives the changes that are required to make the code of Listing 6.3 strictly conform to the OCP. The `INumeric` abstract interface is now formulated as a trait with associated types, through explicit specification of the signatures of the type-bound initializer (i.e. cast function), and the addition and division operators, that the different summation or averaging methods of Listing 6.3 require.

Notice, that Fortran's integer and `real(real64)` types already come with implementations of all this functionality that is predefined for them by the language. Hence, it suffices to acknowledge their conformance to the `INumeric` trait through empty `implements` statements (that are provided here in the separate new module `intrinsic`). This is fully equivalent to the Swift implementation that was presented in Listing 3.4. It is also equivalent to the Rust example of Sect. 3.2.1, in that now no pre-existing code needs to be modified, to make this Fortran version work with additional, new, types. Simply adding a new module that implements the `INumeric` trait for any such new type is now sufficient.

Listing 6.5: Required changes to enforce strict conformance of Listing 6.3 to the OCP.

```

1  module interfaces
2
3      implicit none
4      private
5
6      public :: INumeric, ISum, IAverager
7
8      abstract interface :: INumeric
9          function init(n)
10             integer, intent(in) :: n
11         end function init
12         function operator(+)(lhs,rhs) result(res)
13             type(self), intent(in) :: lhs, rhs
14             type(self) :: res
15         end function operator(+)
16         function operator(/)(lhs,rhs) result(res)
17             type(self), intent(in) :: lhs, rhs
18             type(self) :: res
19         end function operator(/)
20     end interface INumeric
21
22     abstract interface :: ISum
23         function sum{INumeric :: T}(x) result(s)
24             type(T), intent(in) :: x(:)
25             type(T) :: s
26         end function sum
27     end interface ISum
28
29     abstract interface :: IAverager
30         function average{INumeric :: T}(x) result(a)

```

```

31      type(T), intent(in) :: x(:)
32      type(T)              :: a
33      end function average
34      end interface IAvrager
35
36 end module interfaces
37
38 module intrinsics
39
40     use, intrinsic :: iso_fortran_env, only: real64
41     use interfaces, only: INumeric
42
43     implements INumeric :: integer
44     end implements integer
45
46     implements INumeric :: real(real64)
47     end implements real(real64)
48
49 end module intrinsics

```

When comparing Listing 6.5 to the more compact (type set based) code of Listing 6.3, the need for explicit provision of empty implements statements, for all the types that conform to trait INumeric, might appear long-winded. Especially in cases where one would have to deal with a relatively large number of such types. Such implements statements could be completely avoided, though, by adding also structural (i.e. implicit) subtyping to the design, as it is supported in Go. See Sect. 7.2 for a further discussion.

## 6.3 Use of overloaded operators with generic procedures

The following Listing 6.6 provides a worked out example of how, in the present design, traits with overloaded operator signatures can be used with generic routines. This is demonstrated here in terms of two types, namely a derived type, MyType, and the real intrinsic type. Both of these conform to the traits IReducible (with its overloaded multiplication operator signatures), and IPrintable, as they are defined in the module basic\_interfaces.

Since MyType and real conform to the same traits, a set of common operations can be performed on them. This includes (elemental) multiplication of a type with itself, and with integer instances, along with casting to, as well as customized printing of, this type. The generic routines products and prod take advantage of this conformance to carry out a series of array multiplication and reduction operations, and to print out their results in an appropriate format, regardless of whether their argument arrays are of type real or MyType.

Listing 6.6: Use of an overloaded multiplication operator with different generic procedures.

```

1 module basic_interfaces
2
3     implicit none
4     private
5
6     public :: IReducible, IPrintable
7
8     abstract interface :: IReducible
9         function init(n)
10             integer, intent(in) :: n
11         end function init
12         elemental function operator(*) (lhs, rhs) result(res)

```

```

13     type(itself), intent(in) :: lhs, rhs
14     type(itself)           :: res
15     end function operator(*)
16     elemental function operator(*) (lhs, rhs) result(res)
17         type(itself), intent(in) :: lhs
18         integer,          intent(in) :: rhs
19         type(itself)           :: res
20     end function operator(*)
21 end interface IReducible
22
23 abstract interface :: IPrintable
24     subroutine output()
25     end subroutine output
26 end interface IPrintable
27
28 end module basic_interfaces
29
30 module my_type
31
32     use basic_interfaces, only: IReducible, IPrintable
33
34     implicit none
35     private
36
37     type, public, sealed :: MyType
38     private
39     integer :: n
40     end type MyType
41
42     implements (IReducible + IPrintable) :: MyType
43     initial :: init
44     procedure, static :: operator(*) => multiply, multiply_by_int
45     procedure, pass   :: output
46     end implements MyType
47
48 contains
49
50     function init(n) result(res)
51         integer, intent(in) :: n
52         type(MyType)        :: res
53         res%n = n
54     end function init
55
56     elemental function multiply(lhs, rhs) result(res)
57         type(MyType), intent(in) :: lhs, rhs
58         type(MyType)           :: res
59         res%n = lhs%n * rhs%n
60     end function multiply
61
62     elemental function multiply_by_int(lhs, rhs) result(res)
63         type(MyType), intent(in) :: lhs
64         integer,          intent(in) :: rhs
65         type(MyType)           :: res
66         res%n = lhs%n * rhs
67     end function multiply_by_int
68
69     subroutine output(self)
70         type(MyType), intent(in) :: self
71         write(*, '(a,i8)') "I_am:_", self%n

```

```

72     end subroutine output
73
74 end module my_type
75
76 module real_type
77
78     use basic_interfaces, only: IReducible, IPrintable
79
80     implicit none
81     private
82
83     implements (IReducible + IPrintable) :: real
84     procedure, pass :: output
85     end implements real
86
87 contains
88
89     subroutine output(self)
90         real, intent(in) :: self
91         write(*,'(a,f8.1)') "I_am:_", self
92     end subroutine output
93
94 end module real_type
95
96 program main
97
98     use basic_interfaces, only: IReducible, IPrintable
99     use my_type,              only: MyType
100
101     implicit none
102
103     ! declare and initialize arrays with four elements
104     at := [(MyType(i),i:=1,4)] ! Use user-defined initializer for MyType
105     ar := [(real(i),i:=1,4)]   ! Use built-in initializer/cast for real type
106
107     call products(at,ar)
108
109 contains
110
111     subroutine products{IReducible + IPrintable :: T, &
112                     IReducible + IPrintable :: R}(at,ar)
113         type(T), intent(in) :: at(4)
114         type(R), intent(in) :: ar(4)
115         st := prod( (at * at) * [4,3,2,1] ) ! Reduce arrays to scalars after
116         sr := prod( (ar * ar) * [4,3,2,1] ) ! elementwise multiplications
117         call st%output() ! Print the results
118         call sr%output() ! (13824 in both cases)
119     end subroutine products
120
121     function prod{IReducible :: T}(arr) result(res)
122         type(T), intent(in) :: arr(:)
123         type(T)              :: res
124         res = T(1) ! Use initializer for cast
125         do i := 1, size(arr)
126             res = res * arr(i) ! Use operator(*) for reduction
127         end do
128     end function prod
129
130 end program main

```

Of note, here, is the provision of a user-defined initializer for `MyType`, and the use of the `static` attribute for binding actual implementations of `operator(*)` to this type. Derived-type methods that are static, are bound to the type itself rather than its instances. Hence, they cannot take any passed-object arguments, and they must always be dispatched statically.

Notice, also, how none of this is actually required for the `real` type. Because it already comes with suitable, language-defined, implementations of both an intrinsic initializer for casts to this type, and an elemental multiplication operator – as they are required by the `IReducible` trait. We therefore had to merely acknowledge `real`'s conformance to this trait. Whereas we had to provide an own implementation of procedure output, to make `real` conform to the `IPrintable` trait.

Furthermore, and as it was discussed in Sect. 4.2.4, the present Fortran language *rightly prohibits* the programmer (for reasons of predictability) from overloading the intrinsic numeric operators for language-intrinsic types! Any attempt by the programmer to provide some alternative implementation(s) of `operator(*)` for the `real` type should thus be treated as an error. The present generics design makes it easy for a compiler to detect, and ultimately reject, such an error.

## 6.4 Use of associated types with generic implementing types

Lastly, we will demonstrate, now, how associated types in traits (as they were introduced in Sect. 5.1.3) can be used in conjunction with implementing types that are parameterized over some generic parameter. The following example makes use of the `IAppendable` trait, that was defined in Sect. 5.1.3, in order to require an implementing type to provide functionality for appending items to itself. This trait contains an associated type, `Element`, acting as a placeholder for the type of the appendable items, that is to be inferred from an actual implementation of that trait.

```

1  module traits
2
3      implicit none
4
5      abstract interface :: IAppendable
6          typedef, deferred :: Element
7          subroutine append(item)
8              type(Element), intent(in) :: item
9          end subroutine append
10     end interface IAppendable
11
12 end module traits
13
14 module vector_library
15
16     use traits, only: IAppendable
17
18     implicit none
19     private
20
21     type, public, sealed, implements(IAppendable) :: Vector{U}
22         private
23         type(U), allocatable :: elements(:)
24     contains
25         initial :: init
26         procedure, pass :: append
27         procedure, pass :: printout
28     end type Vector
29

```

```

30 contains
31
32     function init(item) result(res)
33         type(U), intent(in) :: item
34         type(Vector{U})      :: res
35         res%elements = [item]
36     end function init
37
38     subroutine append(self,item)
39         type(Vector{U}), intent(inout) :: self
40         type(U),          intent(in)    :: item
41         self%elements = [self%elements,item]
42     end subroutine append
43
44     subroutine printout(self)
45         type(Vector{U}), intent(in) :: self
46         print *, self%elements
47     end subroutine printout
48
49 end module vector_library
50
51 program test_vector
52
53     use vector_library, only: Vector
54
55     implicit none
56
57     doubles := Vector(0.d0)
58     call doubles%append(1.5d0)
59     call doubles%append(2.2d0)
60     call doubles%printout()           ! prints "0. 1.5 2.2"
61
62     bools := Vector{logical}(.true.)
63     call bools%append(.false.)
64     call bools%append(.true.)
65     call bools%printout()             ! prints "T F T"
66
67     strings := Vector("John_")
68     call strings%append("Mary_")
69     call strings%append("Anne_")
70     call strings%printout()           ! prints "John Mary Anne "
71
72 end program test_vector

```

An example of such an implementation of trait `IAppendable` is provided here by the `Vector` derived type. Notice, how the definition of this type is expressed in terms of the declaration `Vector{U}`, that (for the sake of brevity) does *not* contain an explicitly provided constraint for the type parameter `U`. The latter is therefore, according to Sect. 5.1.1, implicitly inferred by the compiler to conform to the `IAnyType` trait. Thus enabling `Vector` to store items of any existing type, `U`, within the generic elements array that it contains.

Notice, also, that for maintaining consistency between the type, `U`, of the elements array, and any new item that we may wish to append to this array, we must force the `item` argument of method `append` to be of type `U` as well, as it is shown in this method's actual implementation. To accomplish this enforcement without contradicting trait `IAppendable`'s definition (that doesn't know anything about type `U`), we made use of the placeholder (i.e. associated) type `Element` in this trait. Given method `append`'s implementation, the compiler will then infer `Element` to be of the

actual type `U`.

Associated types thus allow one to make traits/abstract interfaces generic, without having to parameterize either them or their method signatures by actual generic type parameters, and to thereby entail the instantiation needs of the latter, and the potential duplication of client code that this can result in (cf. Sect. 3.1.1). Associated types are therefore an essential element of the present generics design, and its philosophy not to burden the user with superfluous manual instantiations of generics.

This philosophy, of avoiding the manual provision of boilerplate code, is also evident in the main `test_vector` program of the above listing. This program demonstrates how to use Vectors of different language-intrinsic types, by making use of automatic type inference – both in order to determine a Vector’s appropriate generic type argument from its initializer’s regular argument, and to declare actual Vector instances. Even so, it is important to provide a consistent user-experience, in cases where the programmer does decide to provide generic type arguments manually. This is illustrated in program `test_vector` through the instantiation of object `bools`. There, Vector’s initializer is called with a generic type argument, `logical`, in curly braces; even though the user-provided implementation of this initializer does *not* contain an actual generic parameters list. This list is, instead, understood to be provided implicitly by the compiler, as it was explained in Sect. 5.4.4.

## Chapter 7

# Outlook and final conclusions

This final chapter lists a number of foreseeable extensions, that could be added to the present design once a prototype of it is already successfully implemented.

### 7.1 Rank genericity

Fortran’s special role, as a language that caters to numerical programming, demands that any generics design for the language must allow for the possibility to also handle genericity of array rank. The present design offers a lot of room in this respect, but since this is a largely orthogonal issue, and since we consider a discussion of such functionality to be non-essential for the purpose of a very first prototype implementation of the generics features that were described here, we defer it to a separate document.

### 7.2 Structural subtyping

A pretty natural extension of the present design would be to *supplement* the currently used nominal subtyping with structural subtyping. The latter could be easily incorporated, by making lists of interface/trait names optional (rather than mandatory) when providing trait implementations. The `implements` statement, for instance, would still serve to bind methods (retroactively) to types, and to provide optional nominal subtyping, whose use would remain highly recommended, for documentation purposes [15].

But in all cases where such nominal subtyping would be intentionally skipped, e.g. because it might appear more cumbersome, the compiler would take over. It would infer any traits that a type conforms to, given the functionality that the type supports. This would, for instance, allow one to delete the entire module `intrinsic`s from the code example of Listing 6.5, and to thereby achieve greater conciseness of formulation, while still maintaining strict conformance to the Open/Closed Principle. We ranked structural subtyping as an optional feature, here, mainly in order to ease prototype implementations of the present design.

### 7.3 Union types

Another natural extension would be to allow both type sets, and traits that are formulated in terms of them (cf. Sect 5.1.2), to be used directly within the type specifier of variable declarations. This would offer additional compile-time polymorphism support, via union types (also called sum

types). The example of a generic sum function, that was given in Sect. 5.4.1, could for instance be re-written in the following, schematic, manner:

```
1  function sum(x) result(s)
2      type(integer(*) | real(*)), intent(in) :: x(:)
3      typeof(x)                                :: s
4      s = typeof(x)(0)
5      do i := 1, size(x)
6          s = s + x(i)
7      end do
8  end function sum
```

This no longer employs a generic type parameters list, and instead makes use of both, a union type, and Fortran 2023's `typeof` declaration specifier, in order to allow this code to be used with any of the integer or real types that the language supports. An extension of the present design by union types would take it to its logical conclusion, offer an alternative to generic type parameters for certain use cases, such as the one shown above, and fill the gap that is present in Table 6.1.

## 7.4 Conclusions

The Fortran extensions, that are described in this document for both run-time and compile-time polymorphism, resulted from the consistent application of orthogonal language design. That is, significant new capabilities are provided through the systematic extension of already existing Fortran features, and their mutual interaction, rather than the piecemeal addition of mutually incompatible (and superfluous) new constructs. Indeed, only in one single case (the implements statement of Sect. 4.2) did we find it necessary to introduce a major new language feature. But even then, it was ensured that the new feature would be of *equal* utility in supporting both compile-time and run-time polymorphism, so that orthogonality, again, prevailed.

The end result is a language that is both far less verbose, much easier to use, and incomparably more powerful than that of competing non-holistic approaches, which prefer to abandon the proven philosophy of orthogonal design. The presented extensions are fully backwards compatible with the present Fortran language, and (in stark contrast to competing approaches) allow for a consistent use of fully type-checked generics not just in procedural and functional, but *also* in OO programming settings, *including* the support of static method dispatch (which can improve the performance of polymorphic methods by facilitating inlining via the compiler).

Taken together, these capabilities will allow for the uniform management of source code dependencies on both language-intrinsic and user-defined types in Fortran. They will enable the Fortran programmer to write unprecedented modular code, that is on par with the most modern languages in terms of reusability, and moreover does not sacrifice any computational performance. The present design achieves all of this largely *without* requiring manual instantiations of generics, and it furthermore provides a solid foundation for a number of possible future extensions, like support for array-rank genericity, structural subtyping, and compile time polymorphism via union (sum) types.

# Bibliography

- [1] Luca Cardelli and Peter Wegner. “On understanding types, data abstraction, and polymorphism”. In: *ACM Computing Surveys* 17 (4 Dec. 1985), pages 471–523. URL: <https://dl.acm.org/doi/pdf/10.1145/6041.6042>.
- [2] Brad J. Cox, Steve Naroff, and Hansen Hsu. “The Origins of Objective-C at PPI/Stepstone and Its Evolution at NeXT”. In: *Proceedings of the ACM on Programming Languages* 4 (June 2020), pages 1–74. URL: <https://dl.acm.org/doi/10.1145/3386332>.
- [3] Ole-Johan Dahl. “The Birth of Object Orientation: the Simula Languages”. In: *From Object-Orientation to Formal Methods*. Edited by O. Owe, S. Krogdahl, and T. Lyche. Berlin: Springer, 2004. URL: <https://www.mn.uio.no/ifi/english/about/ole-johan-dahl/bibliography/the-birth-of-object-orientation-the-simula-languages.pdf>.
- [4] Erich Gamma et al. *Design Patterns. Elements of Reusable Object-Oriented Software*. Prentice Hall, 1994.
- [5] *Generics in Swift*. URL: <https://github.com/swiftlang/swift/blob/main/docs/archive/Generics.rst>.
- [6] Magne Haverdaen, Jaakko Järvi, and Damian Rouson. *Reflecting on Generics for Fortran*. July 2019. URL: <https://j3-fortran.org/doc/year/19/19-188.pdf>.
- [7] Allen Holub. *Holub on Patterns, Learning Design Patterns by Looking at Code*. Apress, 2004. ISBN: 1-59059-388-X.
- [8] Steve Klabnik and Carol Nichols. *The Rust Programming Language*. URL: <https://doc.rust-lang.org/stable/book/ch10-02-traits.html#implementing-a-trait-on-a-type>.
- [9] George Lyon. URL: <https://gist.github.com/GeorgeLyon/c7b07923f7a800674bc9745ae45ddc7f>.
- [10] Robert C. Martin. “The Open/Closed Principle”. In: *C++ Report* (Jan. 1996). URL: <https://web.archive.org/web/20060822033314/http://www.objectmentor.com/resources/articles/ocp.pdf>.
- [11] Robert C. Martin. *Agile Software Development: Principles, Patterns, and Practices*. Prentice Hall, 2002. ISBN: 978-0-135-97444-5.
- [12] Robert C. Martin. *OO vs FP*. Nov. 2014. URL: <https://blog.cleancoder.com/uncle-bob/2014/11/24/FPvsOO.html>.
- [13] Robert C. Martin. *Clean Architecture: A Craftsman’s Guide to Software Structure and Design*. Boston, MA: Prentice Hall, 2017. ISBN: 978-0-13-449416-6.
- [14] Nicholas D Matsakis and Felix S Klock II. “The Rust language”. In: *ACM SIGAda Ada Letters*. Volume 34. ACM. 2014, pages 103–104.
- [15] *Mojo Manual: Traits*. URL: <https://docs.modular.com/mojo/manual/traits/>.

- [16] Slava Pestov. *Compiling Swift Generics*. URL: <https://github.com/swiftlang/swift/tree/main/docs/Generics>.
- [17] Benjamin C. Pierce. *Programming with intersection types, union types, and polymorphism*. URL: <https://api.semanticscholar.org/CorpusID:118907019>.
- [18] PL22.3 and WG5. *J3/18-007r1 (F2018 interpretation document)*. Oct. 2018. URL: <https://j3-fortran.org/doc/year/18/18-007r1.pdf>.
- [19] Ian Lance Taylor. URL: <https://github.com/golang/go/issues/45346>.
- [20] *The Java Tutorials. Abstract Methods and Classes*. URL: <https://docs.oracle.com/javase/tutorial/java/IandI/abstract.html>.
- [21] *The Swift Programming Language (6 beta). Extensions*. URL: <https://docs.swift.org/swift-book/documentation/the-swift-programming-language/extensions/>.
- [22] Wolfgang Weck and Clemens Szyperski. *Do We Need Inheritance?* URL: [https://www.researchgate.net/publication/2297653\\_Do\\_We\\_Need\\_Inheritance](https://www.researchgate.net/publication/2297653_Do_We_Need_Inheritance).

## Acknowledgements

The present extensions for Fortran could not have been devised without the guidance we had through the prior work of countless developers on the Swift, Go, Rust, Carbon, and Java languages. We wish to thank all these developers, even though we can only mention here a few of them explicitly by name, who have particularly influenced the present effort. Chris Lattner's design of the Swift language provided the foundational ideas upon which we built the majority of our language extensions for Fortran, including much of the modern OO features, and the generics. While Robert Griesemer's, Ian Lance Taylor's, and the entire Go team's inspirational work on type sets accounts largely not only for our remaining generics design, but also for much of its ease of use. Robert also kindly provided the original code version from which Listing 2.2 was derived, and answered the many questions we had regarding Go's generics. We are also grateful to everyone else who gave us their feedback to our work, in particular Chris Lattner, Magne Haverdaen, and Brad Richardson.